



Deep Learning





Course Introduction



Scope of the Course

Course Duration

12-25 Jan

8 Sessions, 2 Labs, **26 Hours in Total**

Structure of Course



Course Projects

2 Mini Projects & Assignments

Ass-1 **ANN** (Churn Prediction For Bank Customer),
Ass-2 **CNN** (Image Classification),

1 Medium Projects

CNN (Image (dog/cat) Classification),



Deep Learning

Session-1



Table of Contents

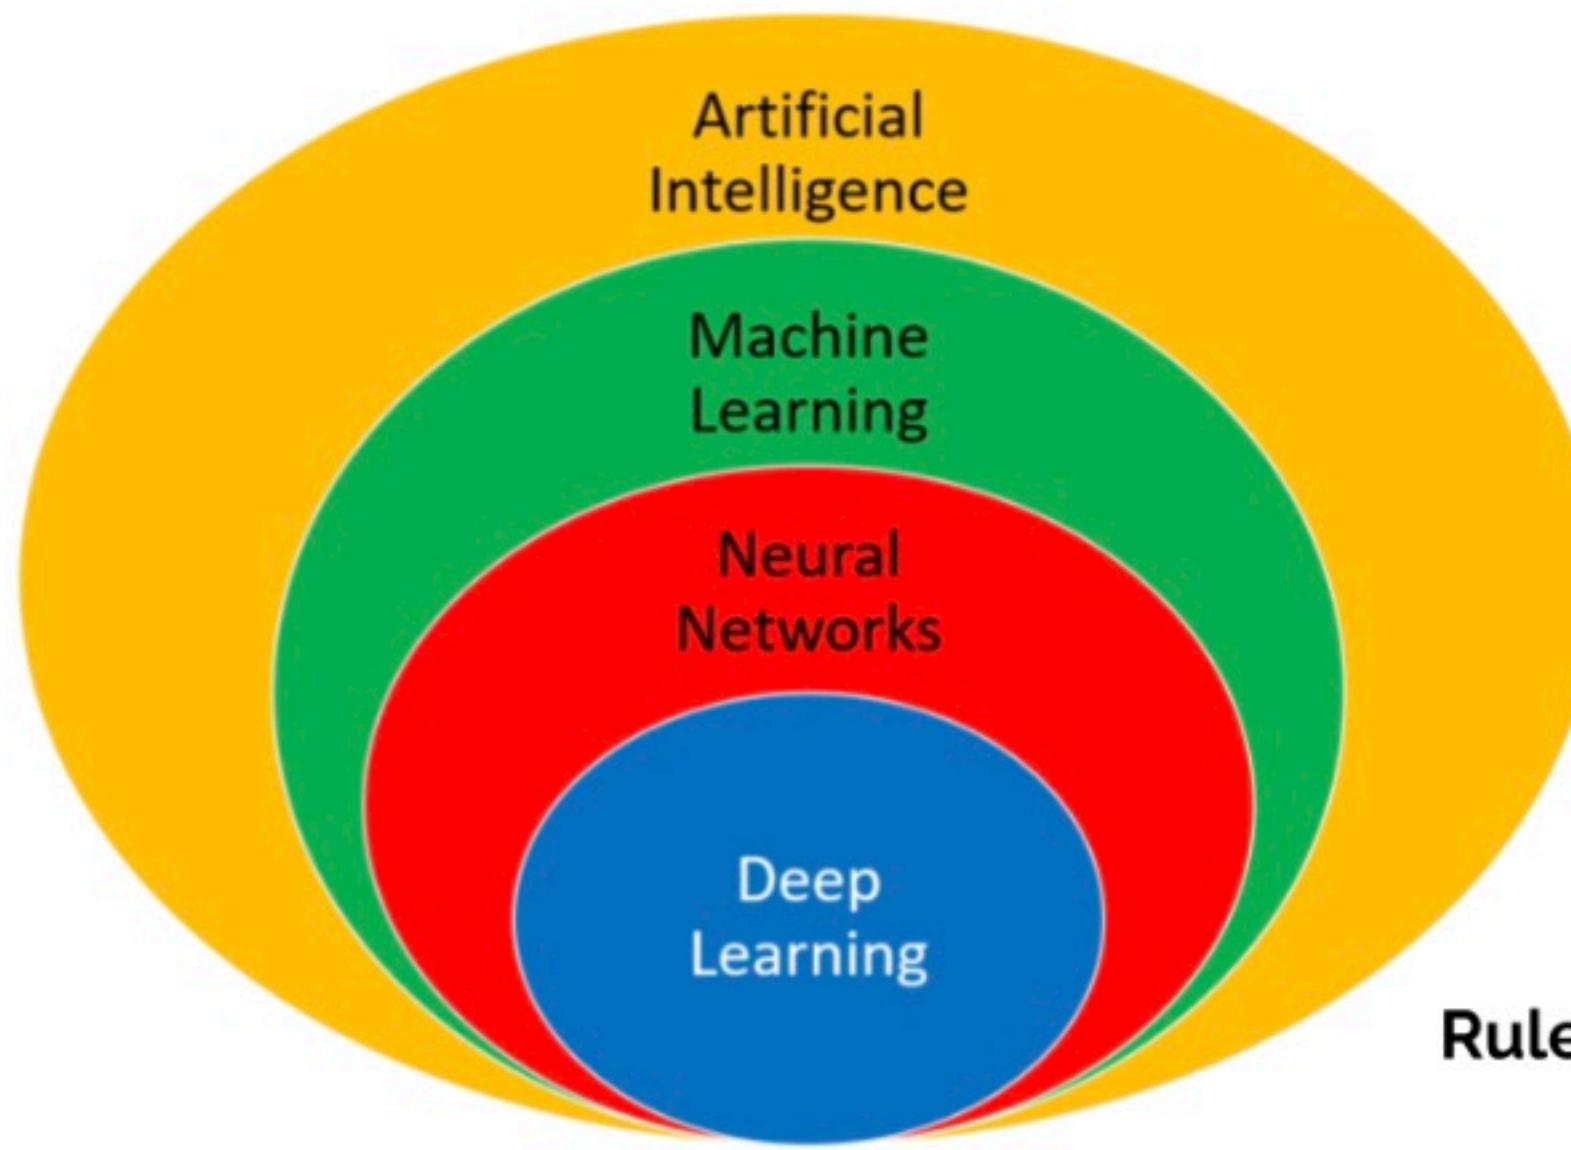
- ▶ What is Deep Learning?
- ▶ Perceptron Model & Neural Networks
- ▶ Activation Functions
- ▶ Multiclass Classification
- ▶ Cost Function



What is Deep Learning?



What is Deep Learning?

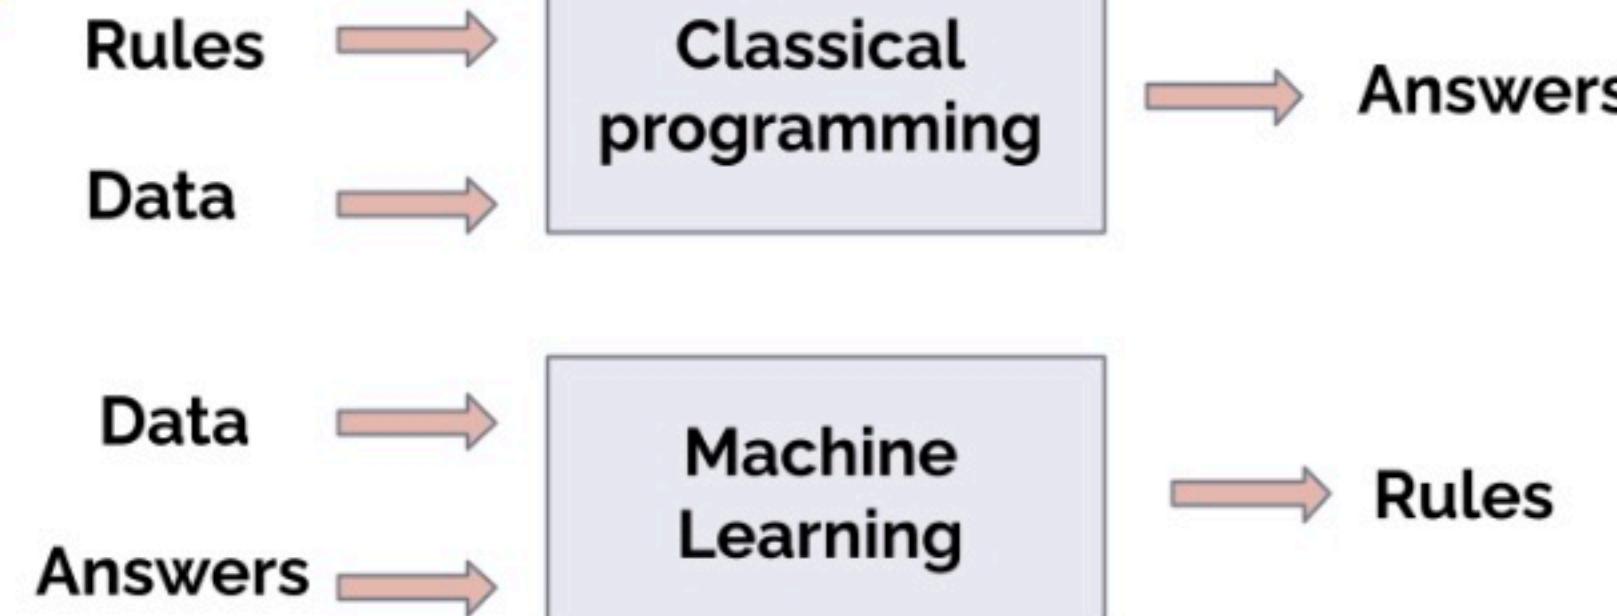


"the field of study that gives computers the ability to learn without being explicitly programmed."

AI

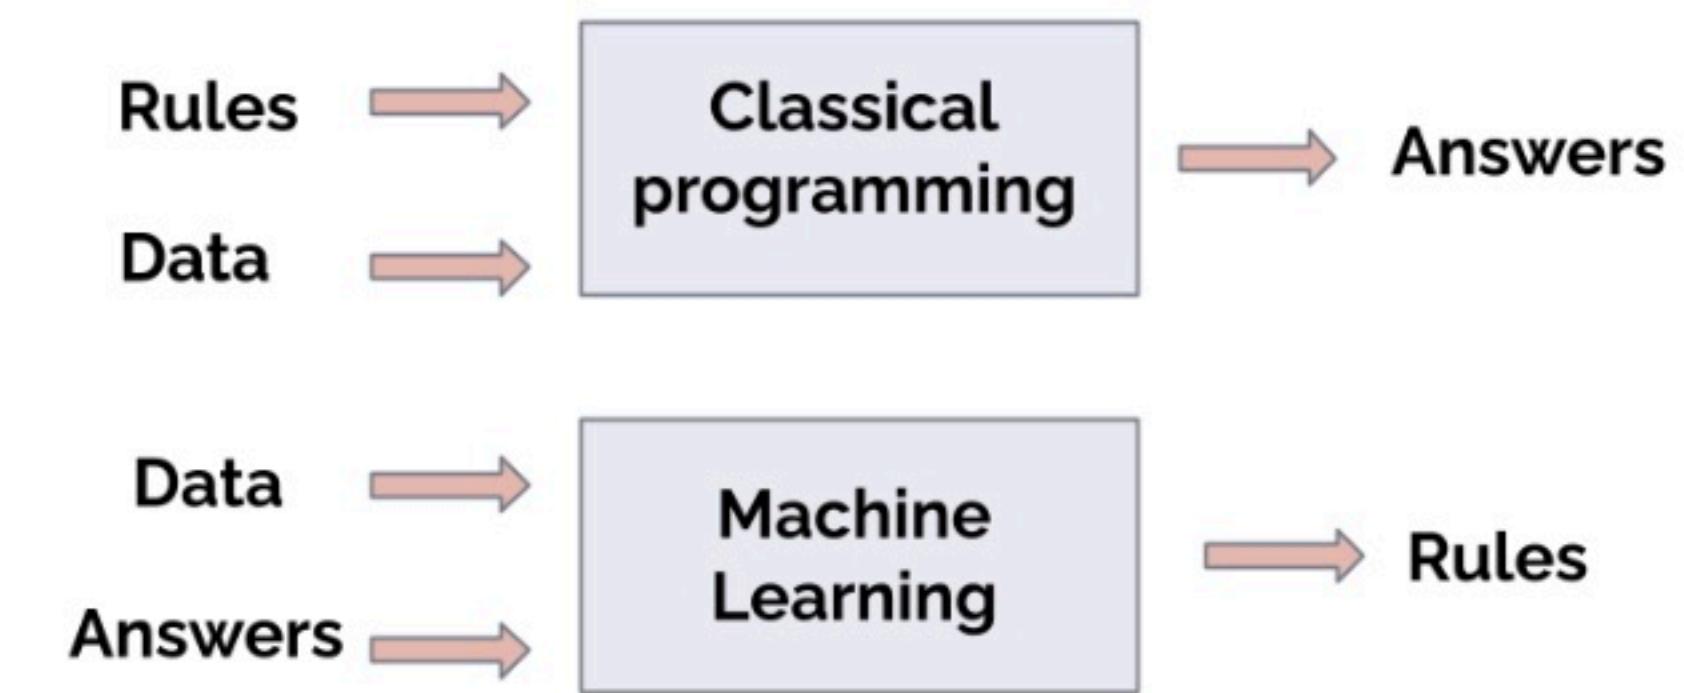
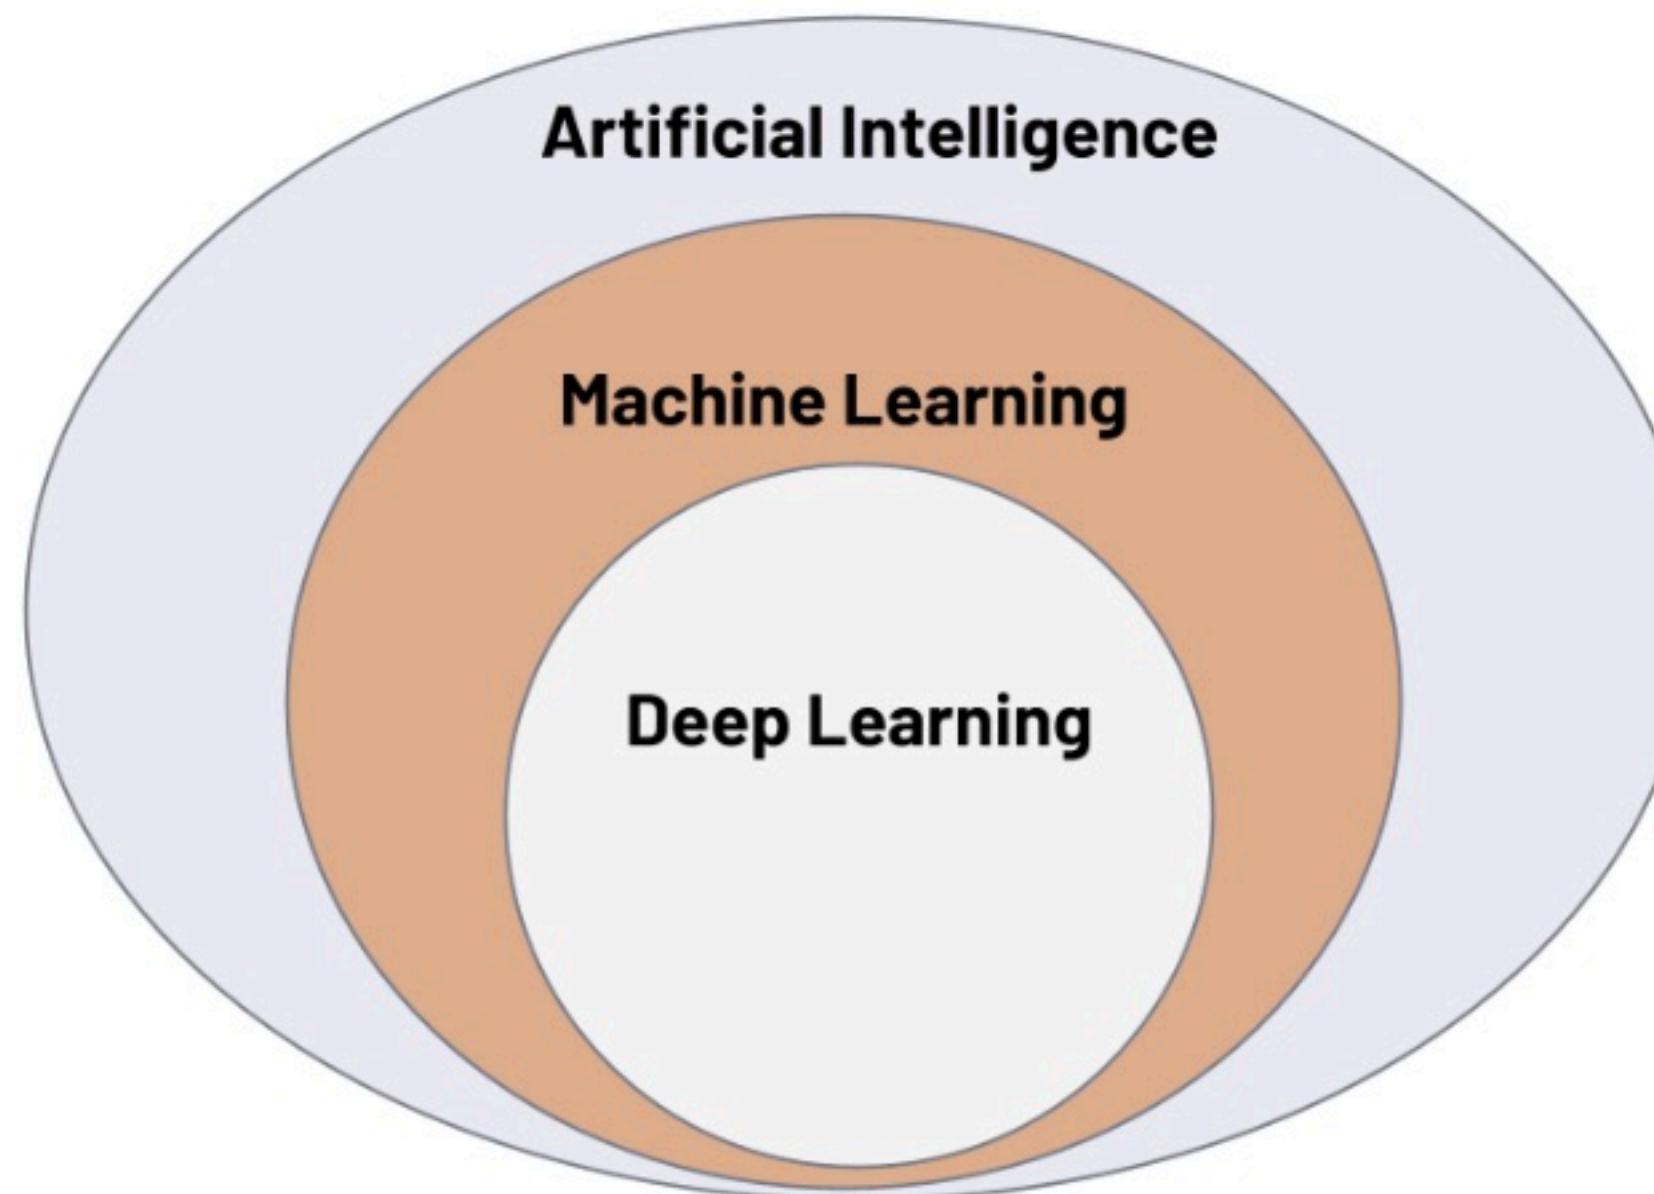
The effort to automate intellectual tasks normally performed by humans. AI is a general field that encompasses machine learning and deep learning.

ML/DL

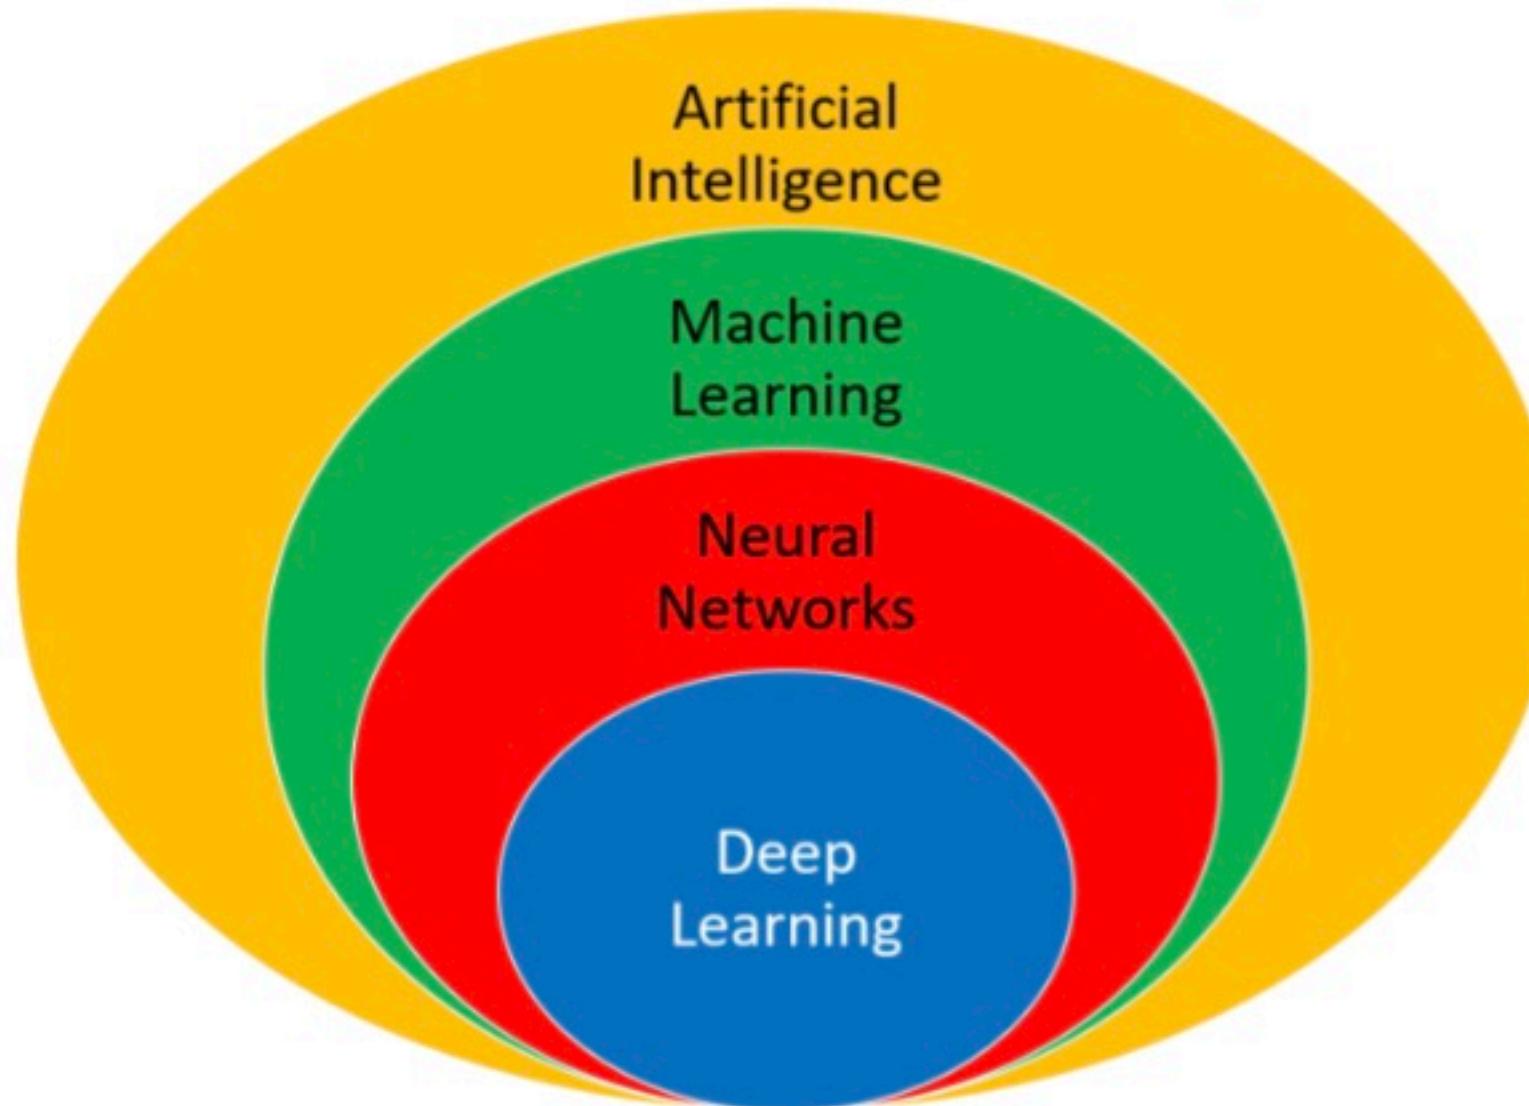


A new
programming
paradigm

What is Deep Learning?



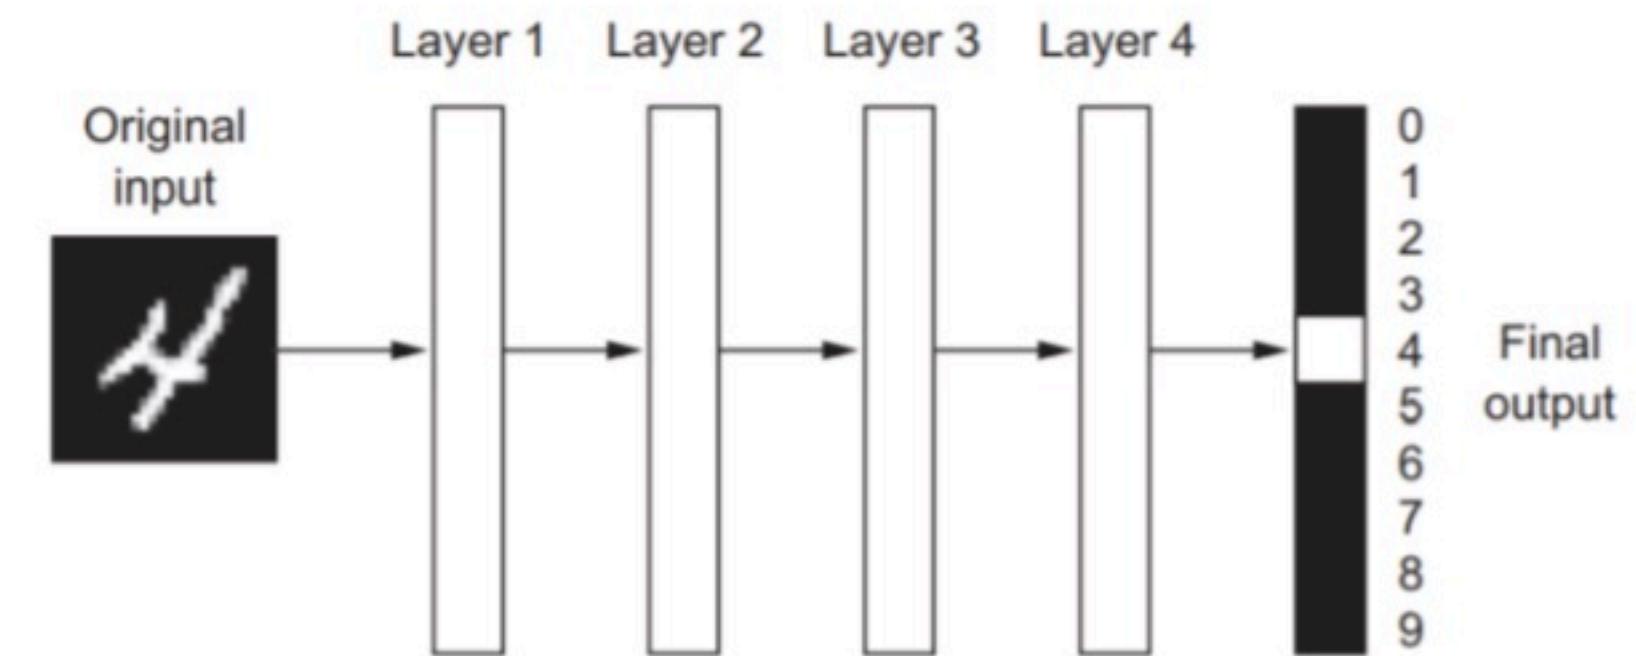
What is Deep Learning?



"The deep in deep learning isn't a reference to any kind of deeper understanding achieved by the approach; rather, it stands for this idea of successive layers of representations"

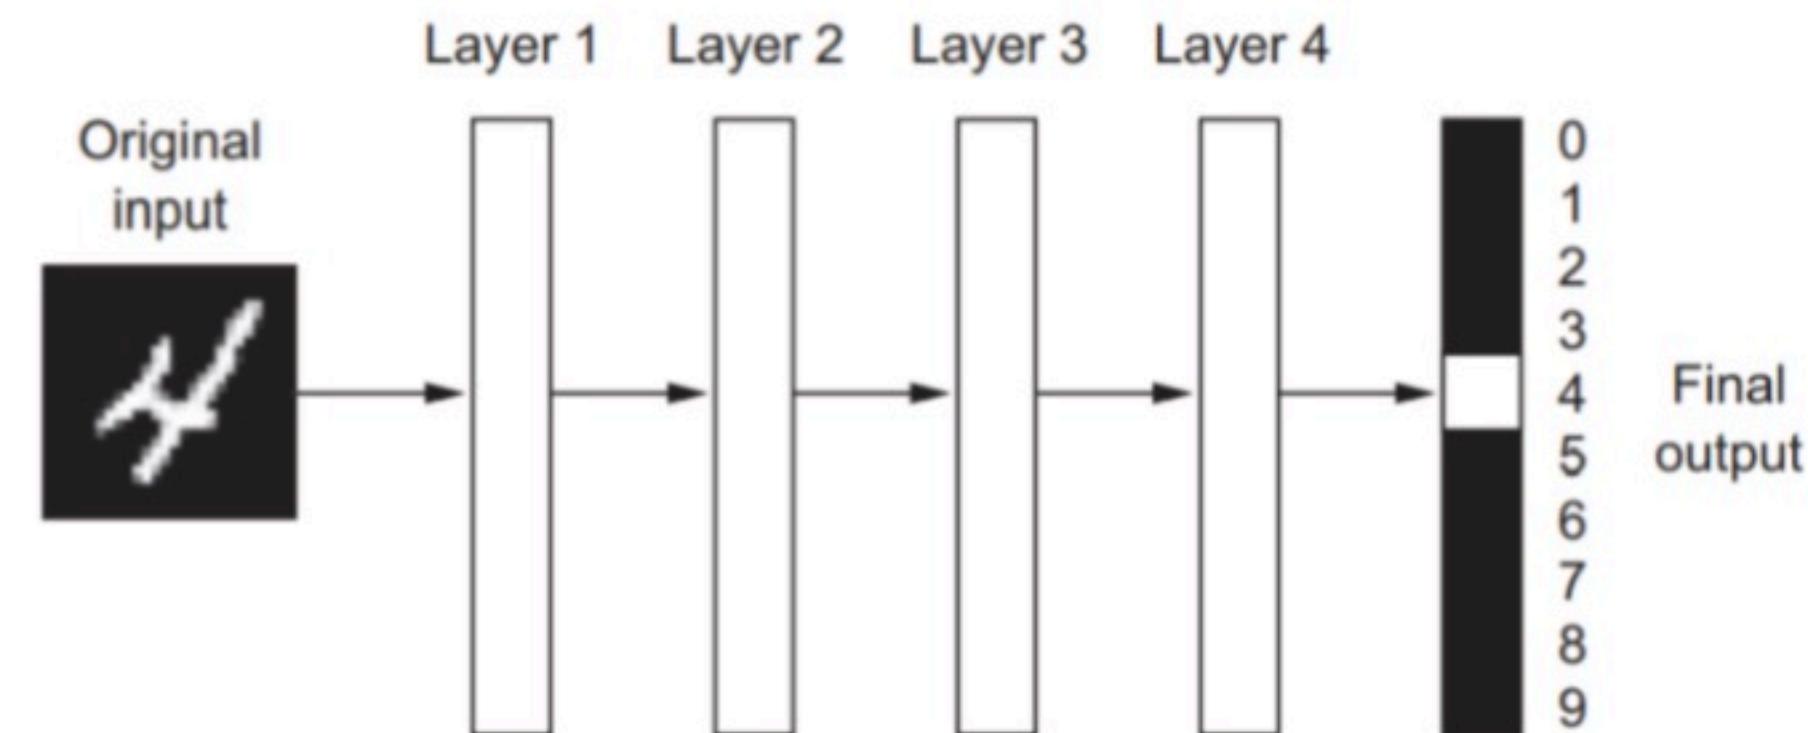
Deep Learning

Subfield of machine learning concerned with algorithms **inspired by the structure and function of the brain** called **artificial neural networks**.



What is Deep Learning?

Subfield of machine learning concerned with algorithms **inspired by the structure and function of the brain** called **artificial neural networks**.



What is Deep Learning?



Andrew Ng

The idea of deep learning as:

Using brain simulations, hope to:

- Make learning algorithms **much better and easier** to use.
- Make **revolutionary advances** in machine learning and AI.

*I believe this is our best shot at progress towards **real AI***

*..almost all the value today of deep learning is through **supervised learning** or learning from labeled data*

*..one reason that deep learning has taken off like crazy is because it is fantastic at **supervised learning***

What is Deep Learning?

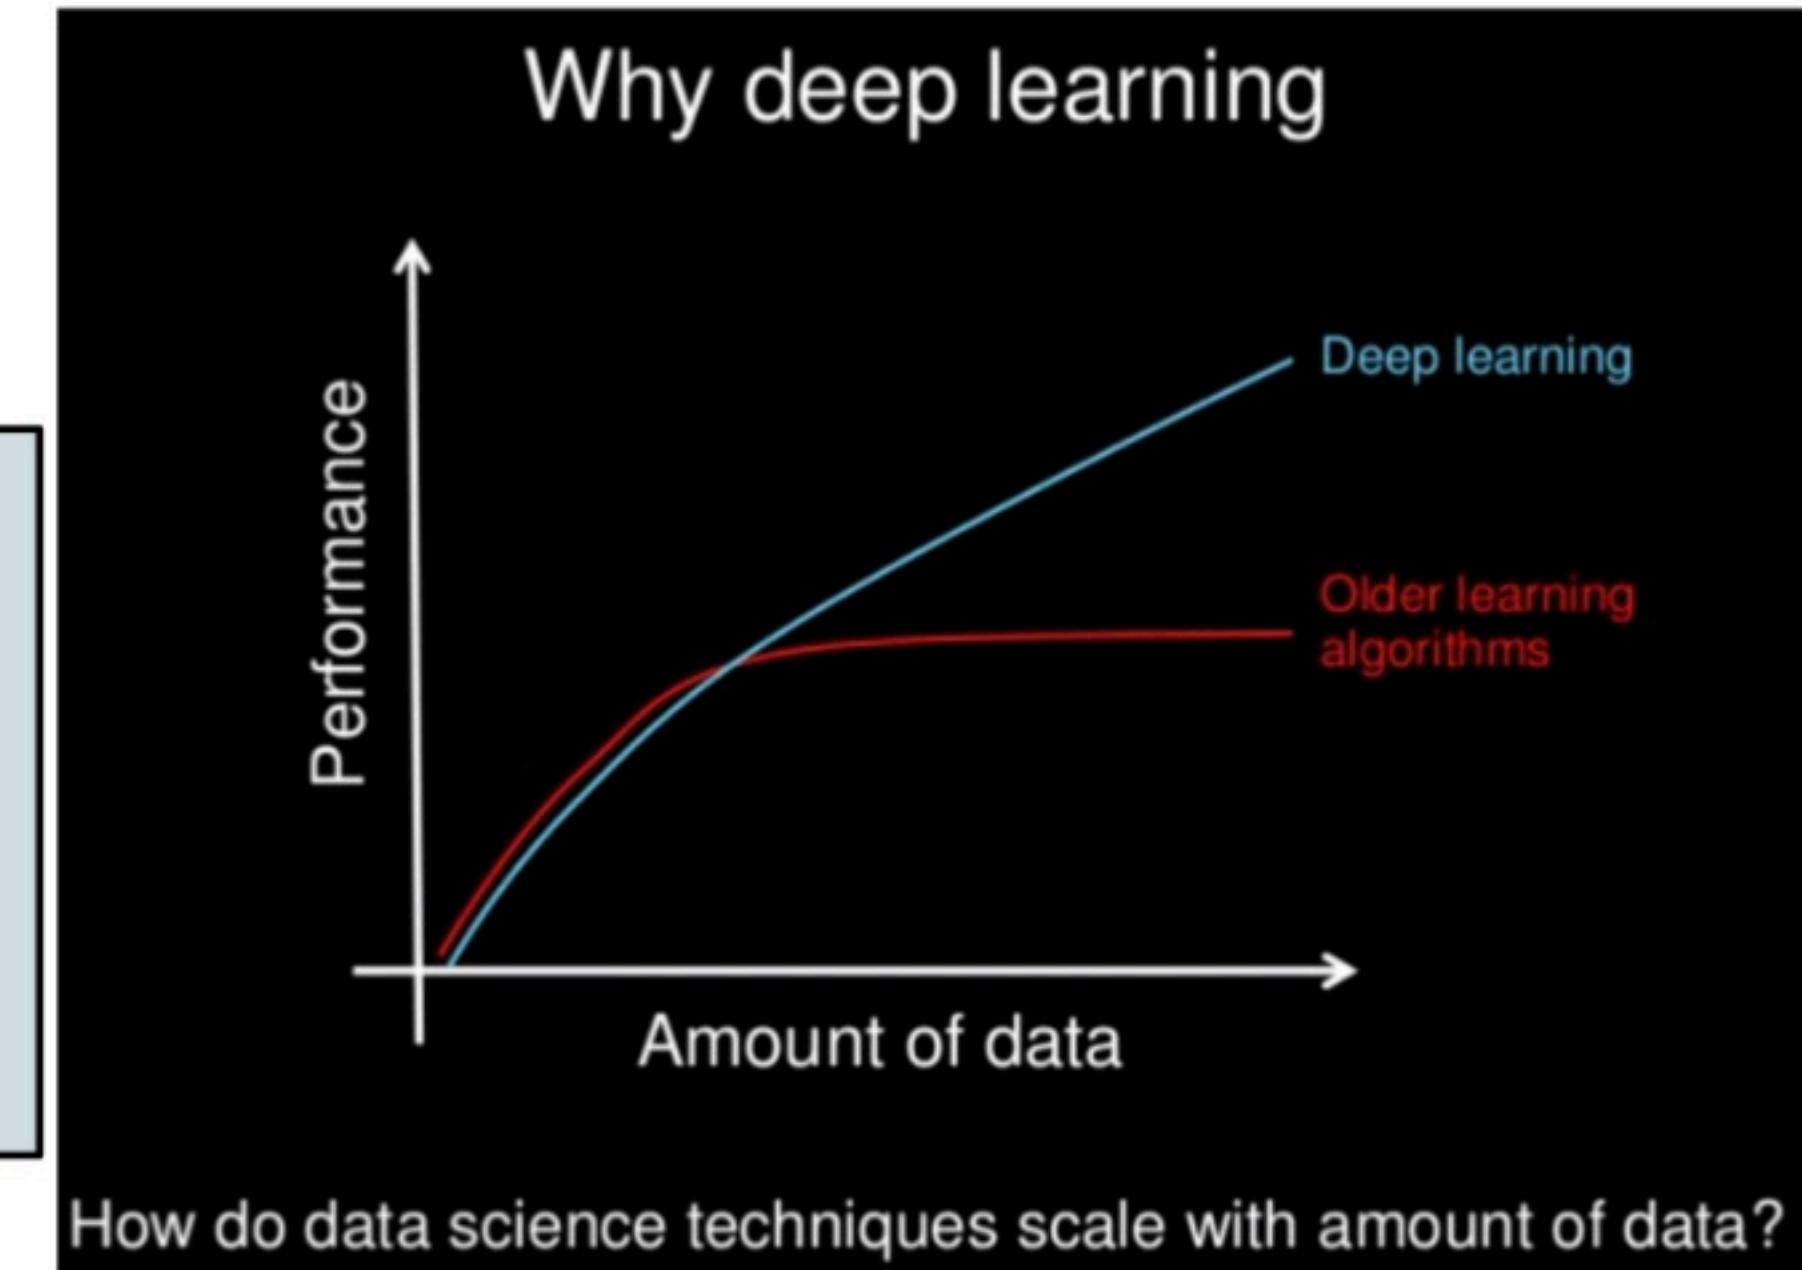
Scalability

very large neural networks we can now have and ... huge amounts of data that we have access to

for most flavors of the old generations of learning algorithms ... performance will plateau.

...

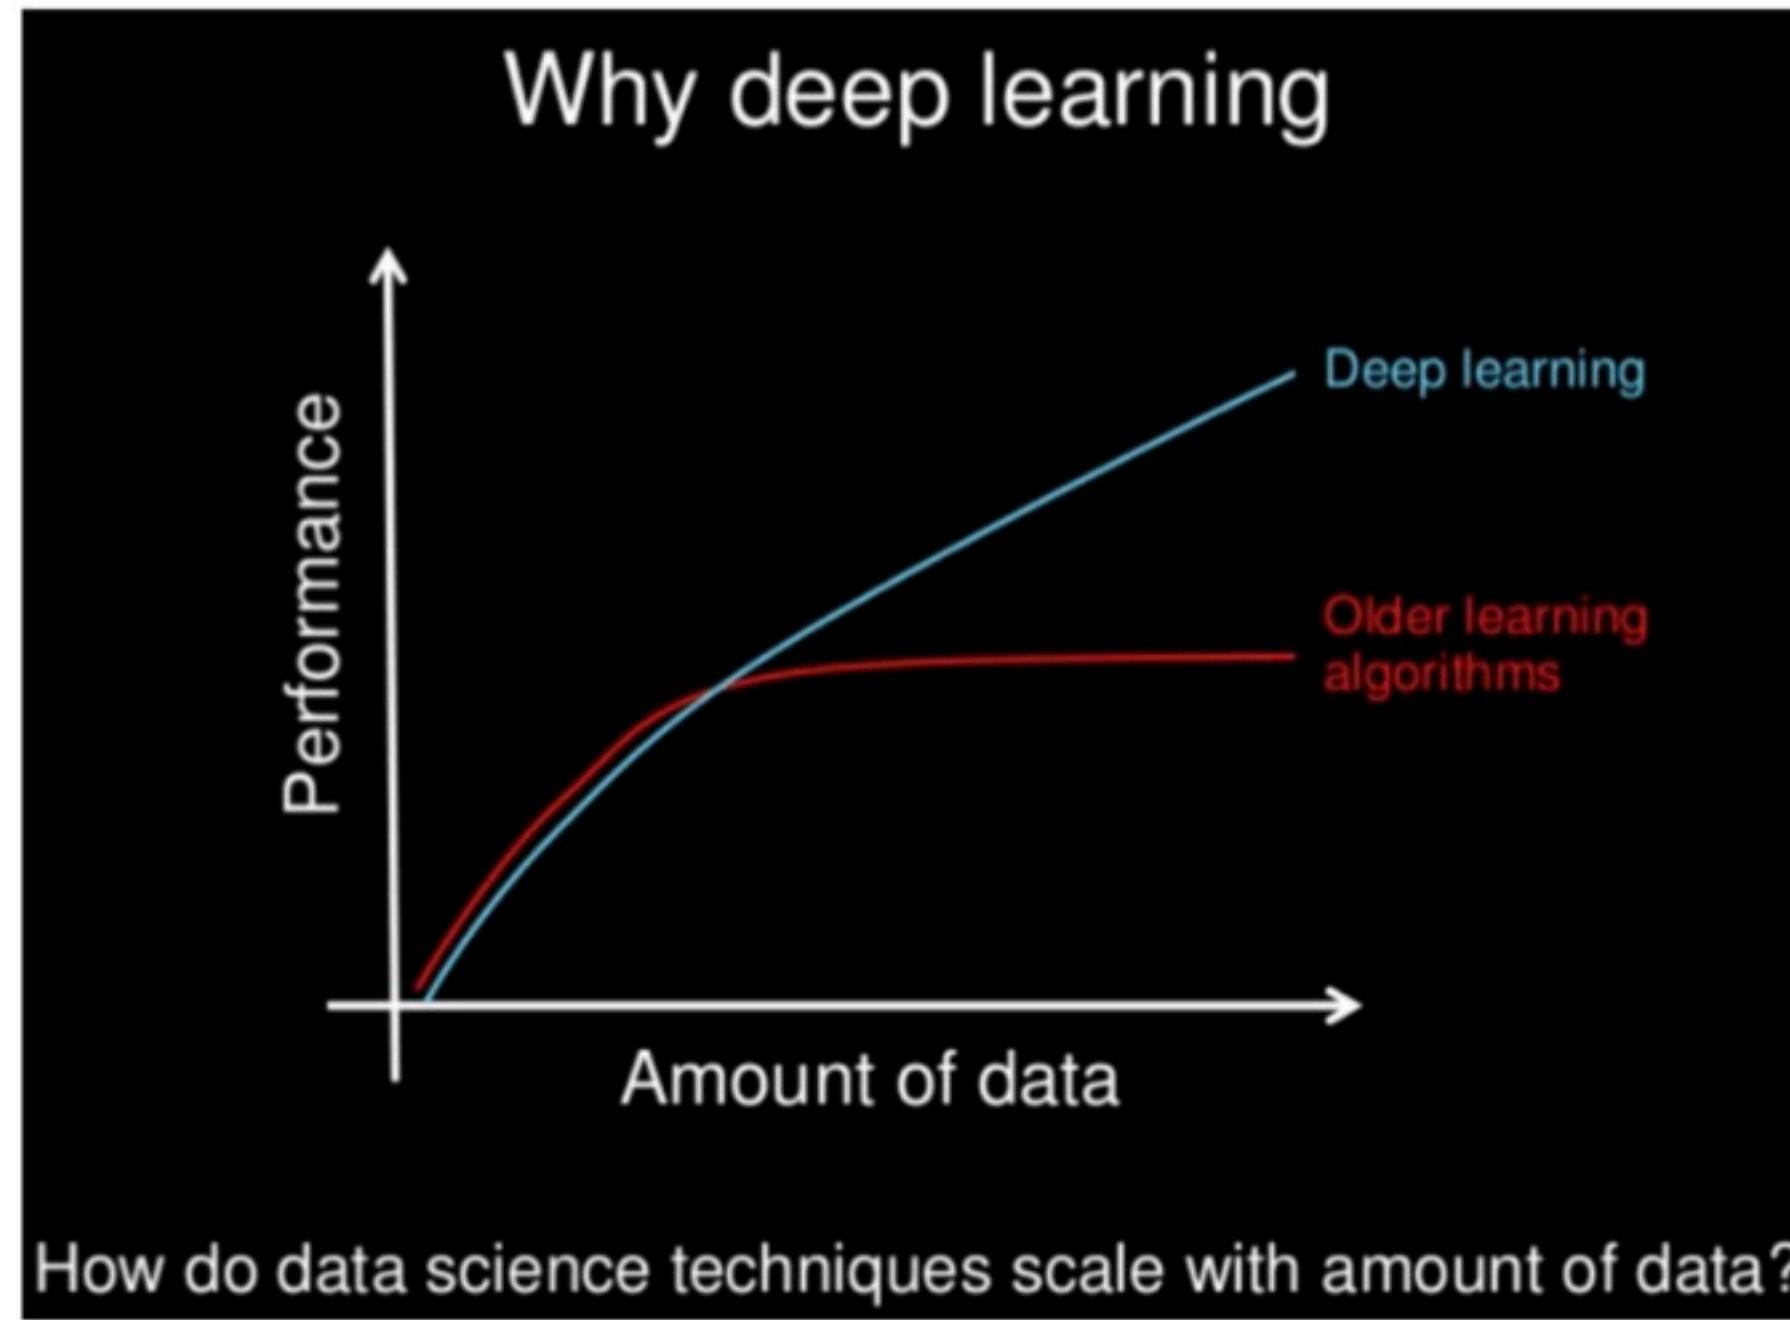
deep learning ... is the first class of algorithms ... that is scalable. ... performance just keeps getting better as you feed them more data



What is Deep Learning?

Scalability

very large neural networks we can now have and ... huge amounts of data that we have access to



What is Deep Learning?

Scalability



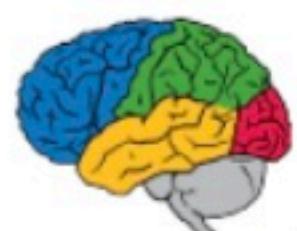
Jeff Dean

Important Property of Neural Networks

Results get better with

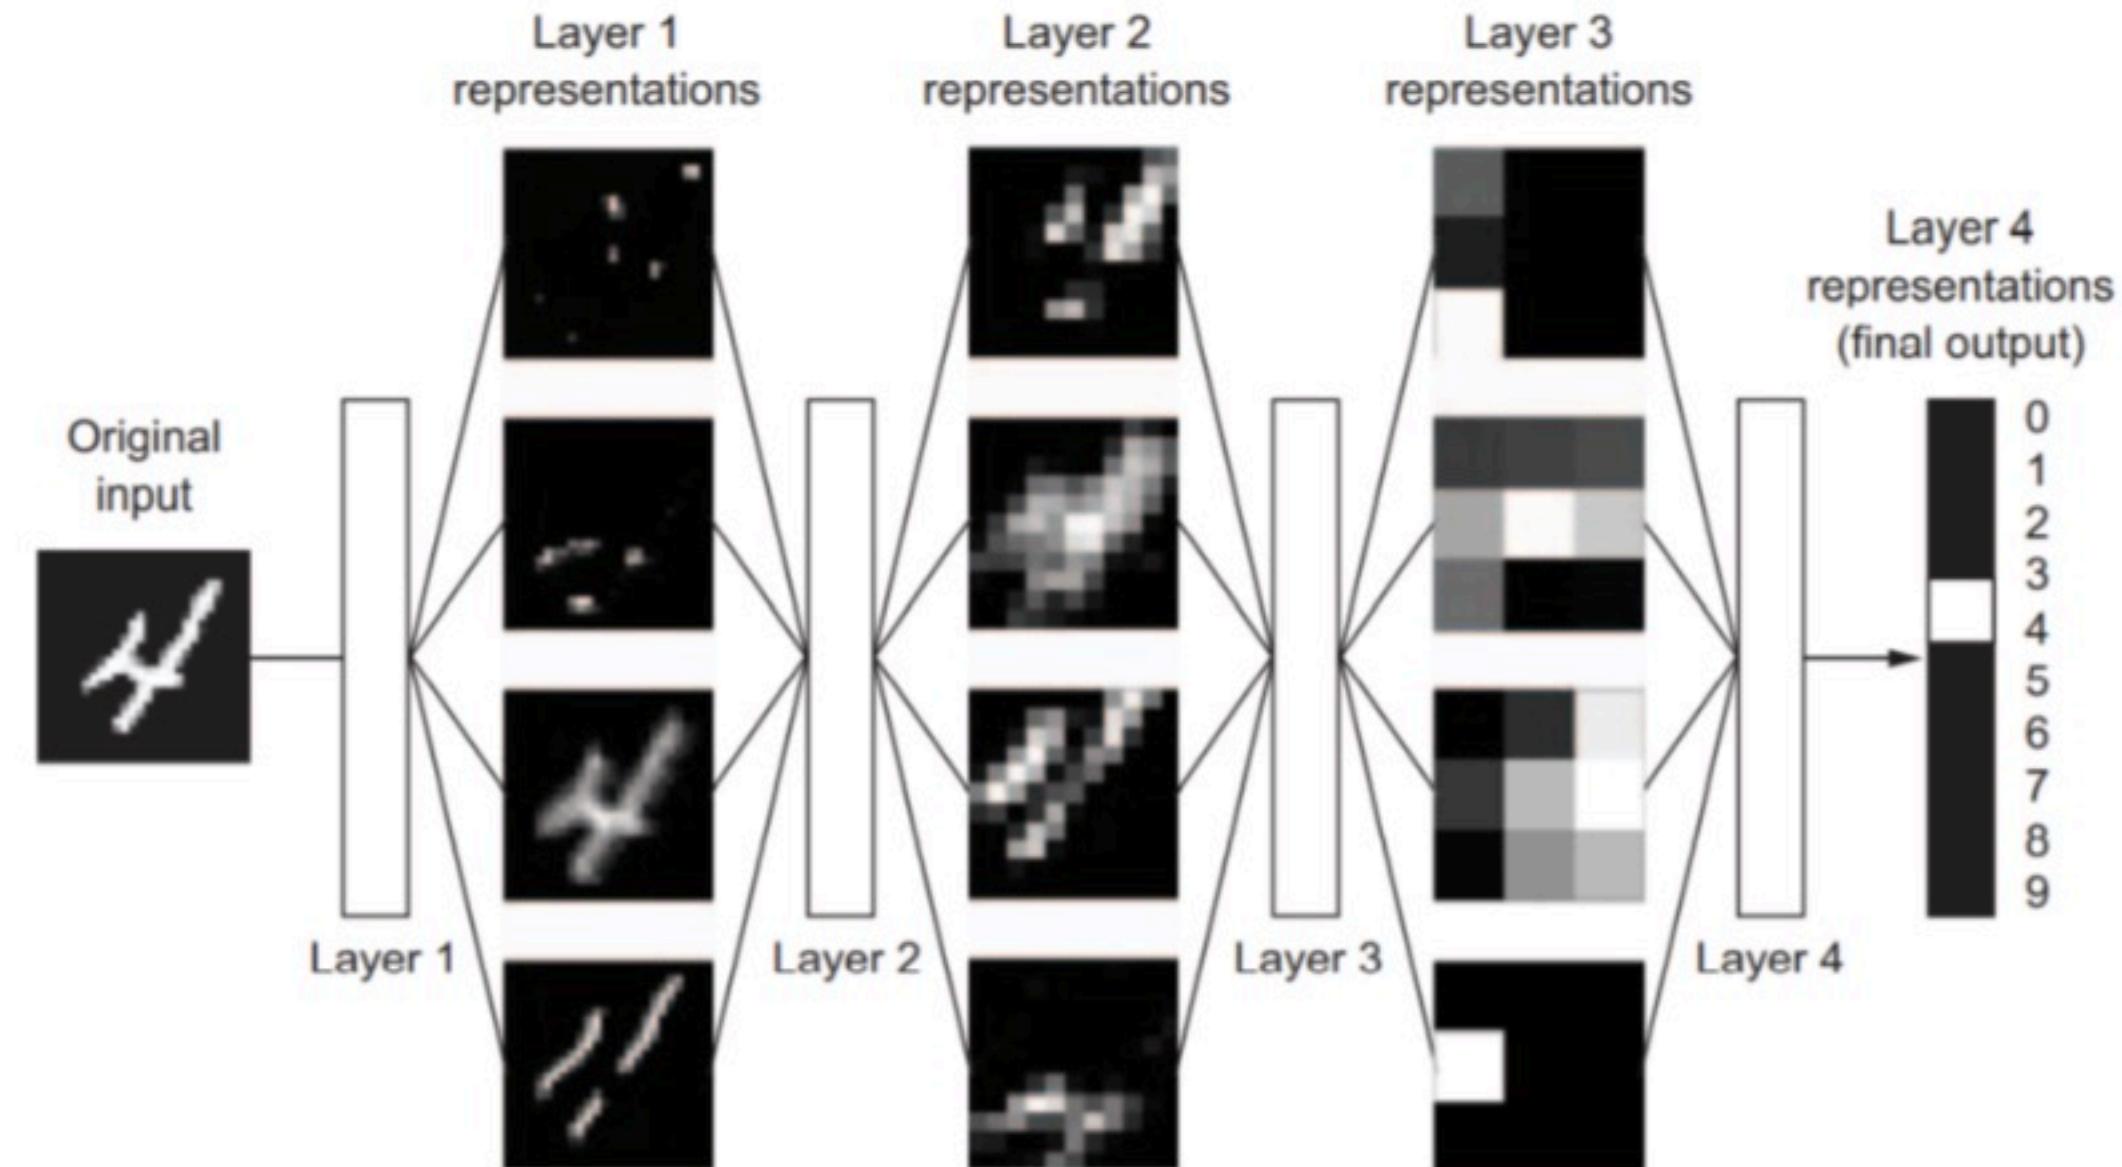
**more data +
bigger models +
more computation**

**(Better algorithms, new insights and improved
techniques always help, too!)**



What is Deep Learning?

Feature Learning



*...a kind of learning where the representation you form have **several levels of abstraction**, rather than a direct input to output*

What is Deep Learning?



The reasons for boost in Deep Learning are:

- ▶ Large amount of training data
- ▶ Faster machines and multicore **CPU/GPUs**
- ▶ New models, algorithms, ideas:
 - Better, more flexible **learning** of intermediate **representations**
 - Effective **learning methods** for using contexts and **transferring** of learned knowledge between tasks
 - Better **regularization** and **optimization** techniques

What is Deep Learning?



*Modern state-of-the-art deep learning is focused on **training deep (many layered) neural network models***

The most popular techniques are:

- ▶ Multilayer Perceptron Networks.
- ▶ Convolutional Neural Networks.
- ▶ Recurrent Neural Networks.

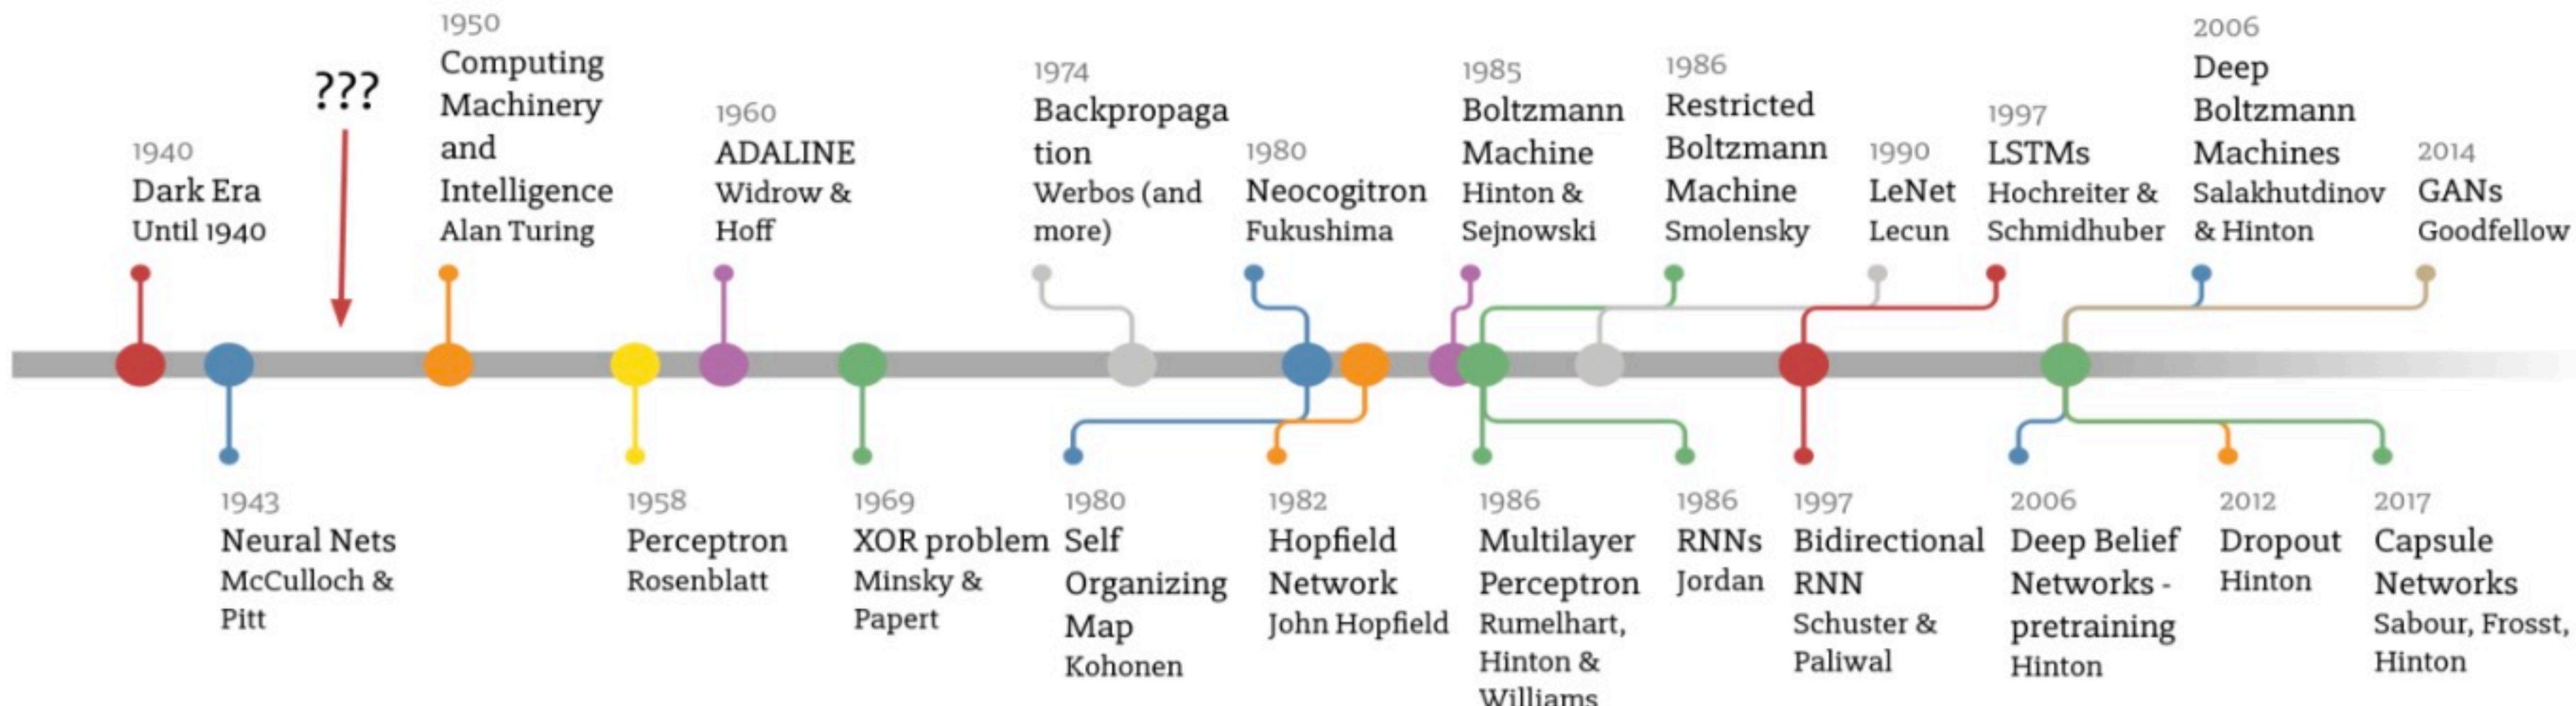
DEEP

Artificial
Neural
Networks

What is Deep Learning?



Deep Learning Timeline





<https://www.nytimes.com/2019/03/27/technology/turing-award-ai.html>

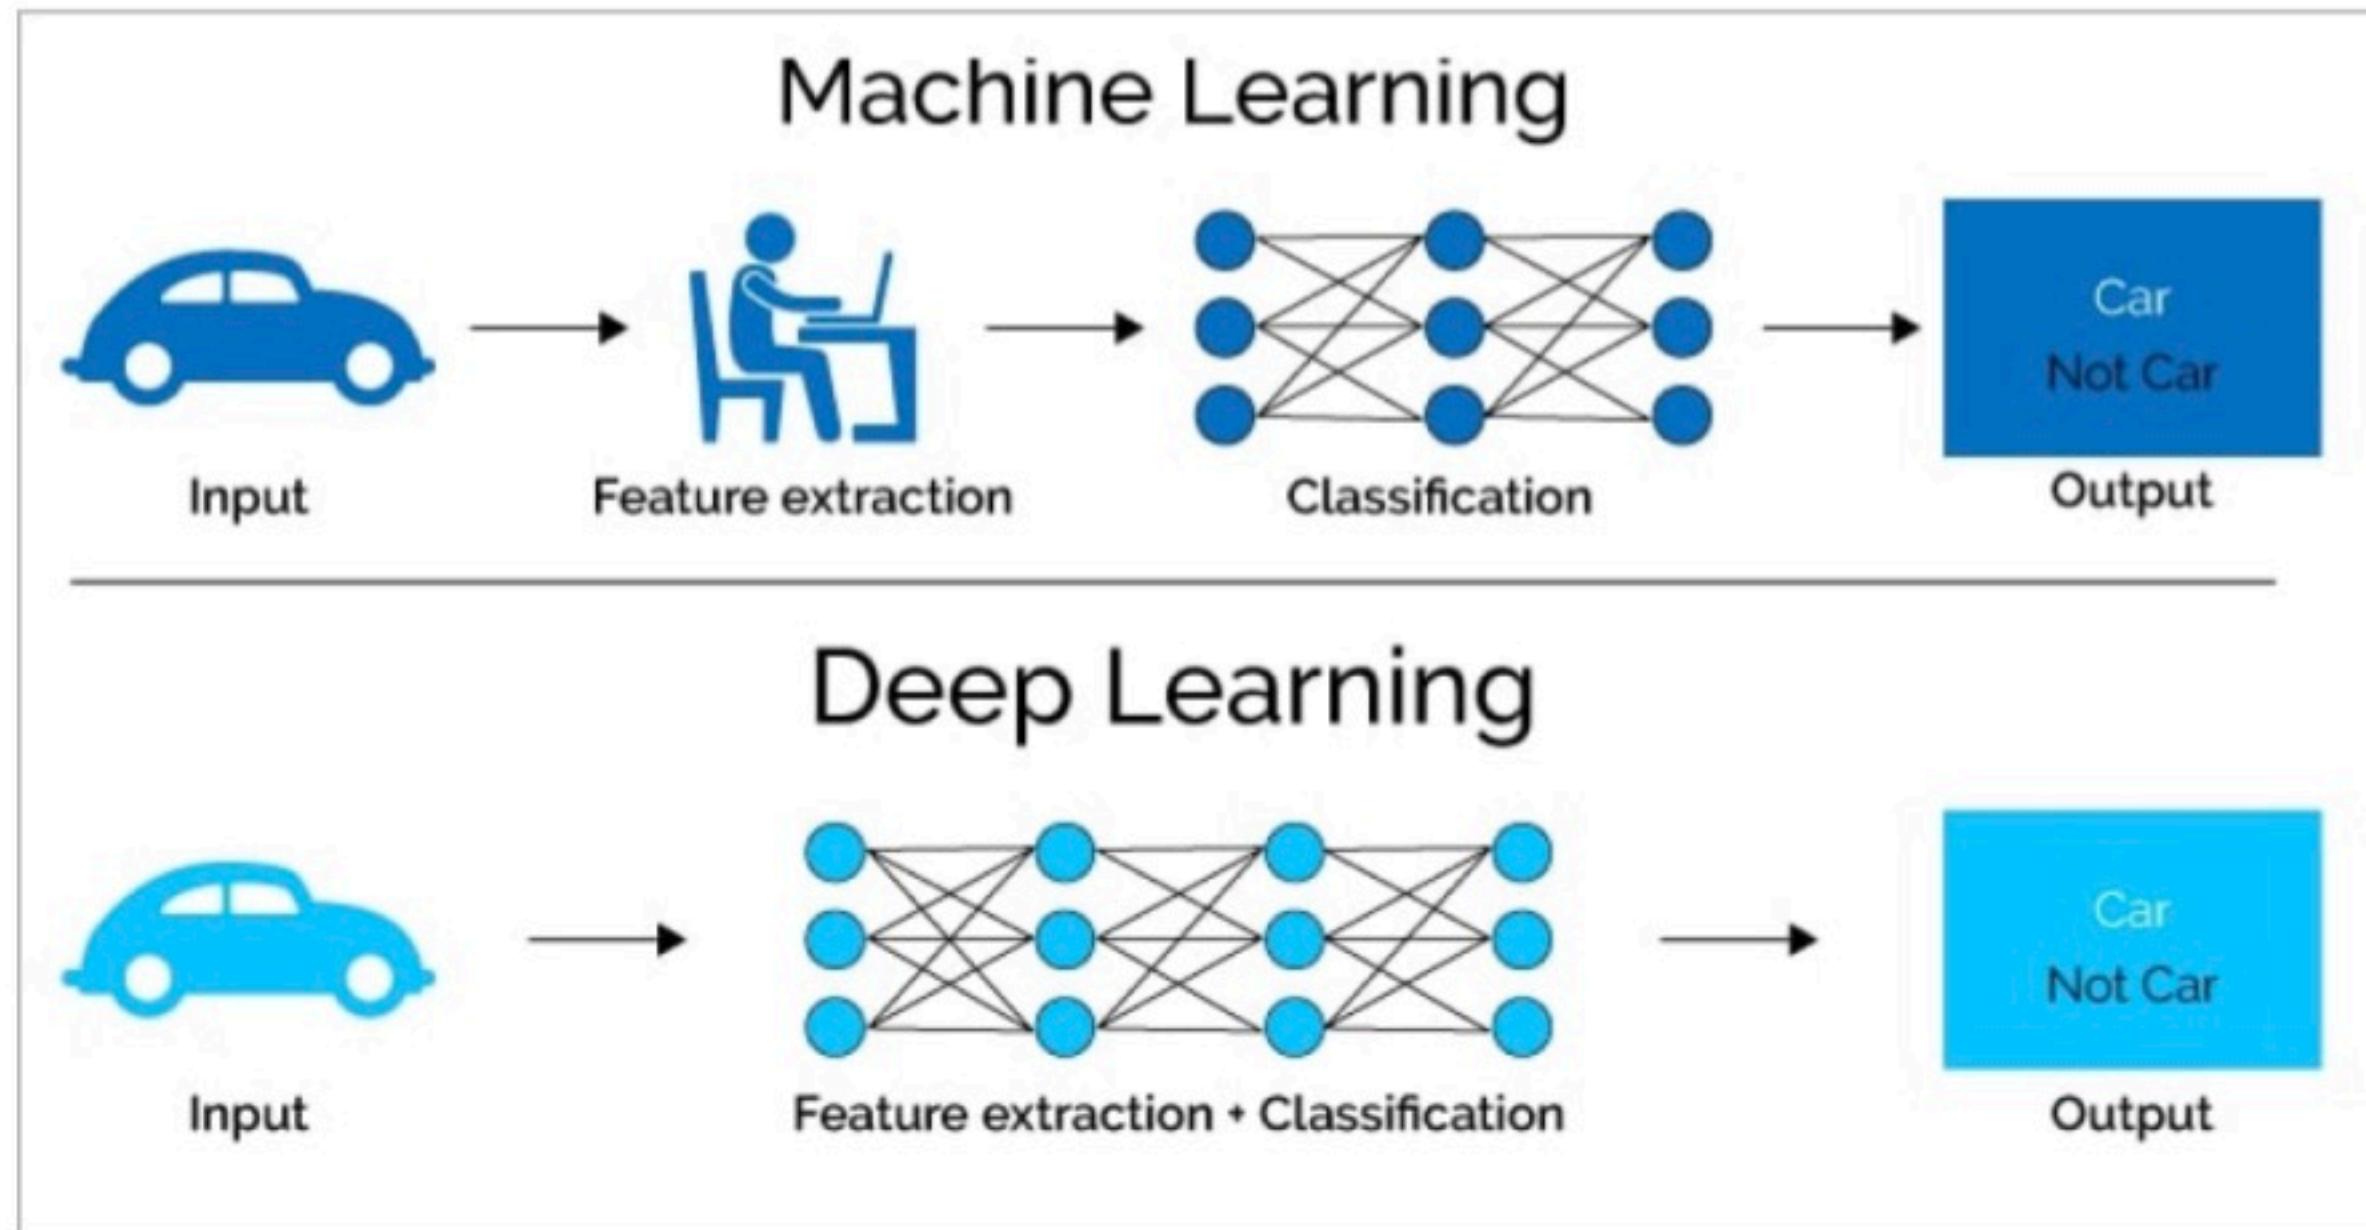
What is Deep Learning?



Briefly;

- ▶ Deep learning is a specialized **subset of machine learning**.
- ▶ Deep learning relies on a **layered structure** of algorithms called an artificial neural network.
- ▶ Deep learning has **huge data needs** but requires little human intervention to function properly.
- ▶ **Transfer learning** is a cure for the needs of large training datasets.

What is Deep Learning?



Deep Learning can process **unstructured data** such as **documents, images, and text.**

The deep learning algorithm doesn't need a software engineer to identify features but is capable of automatic feature engineering through its neural network. (Source: softwaretestinghelp.com)



Perceptron Model & Neural Networks



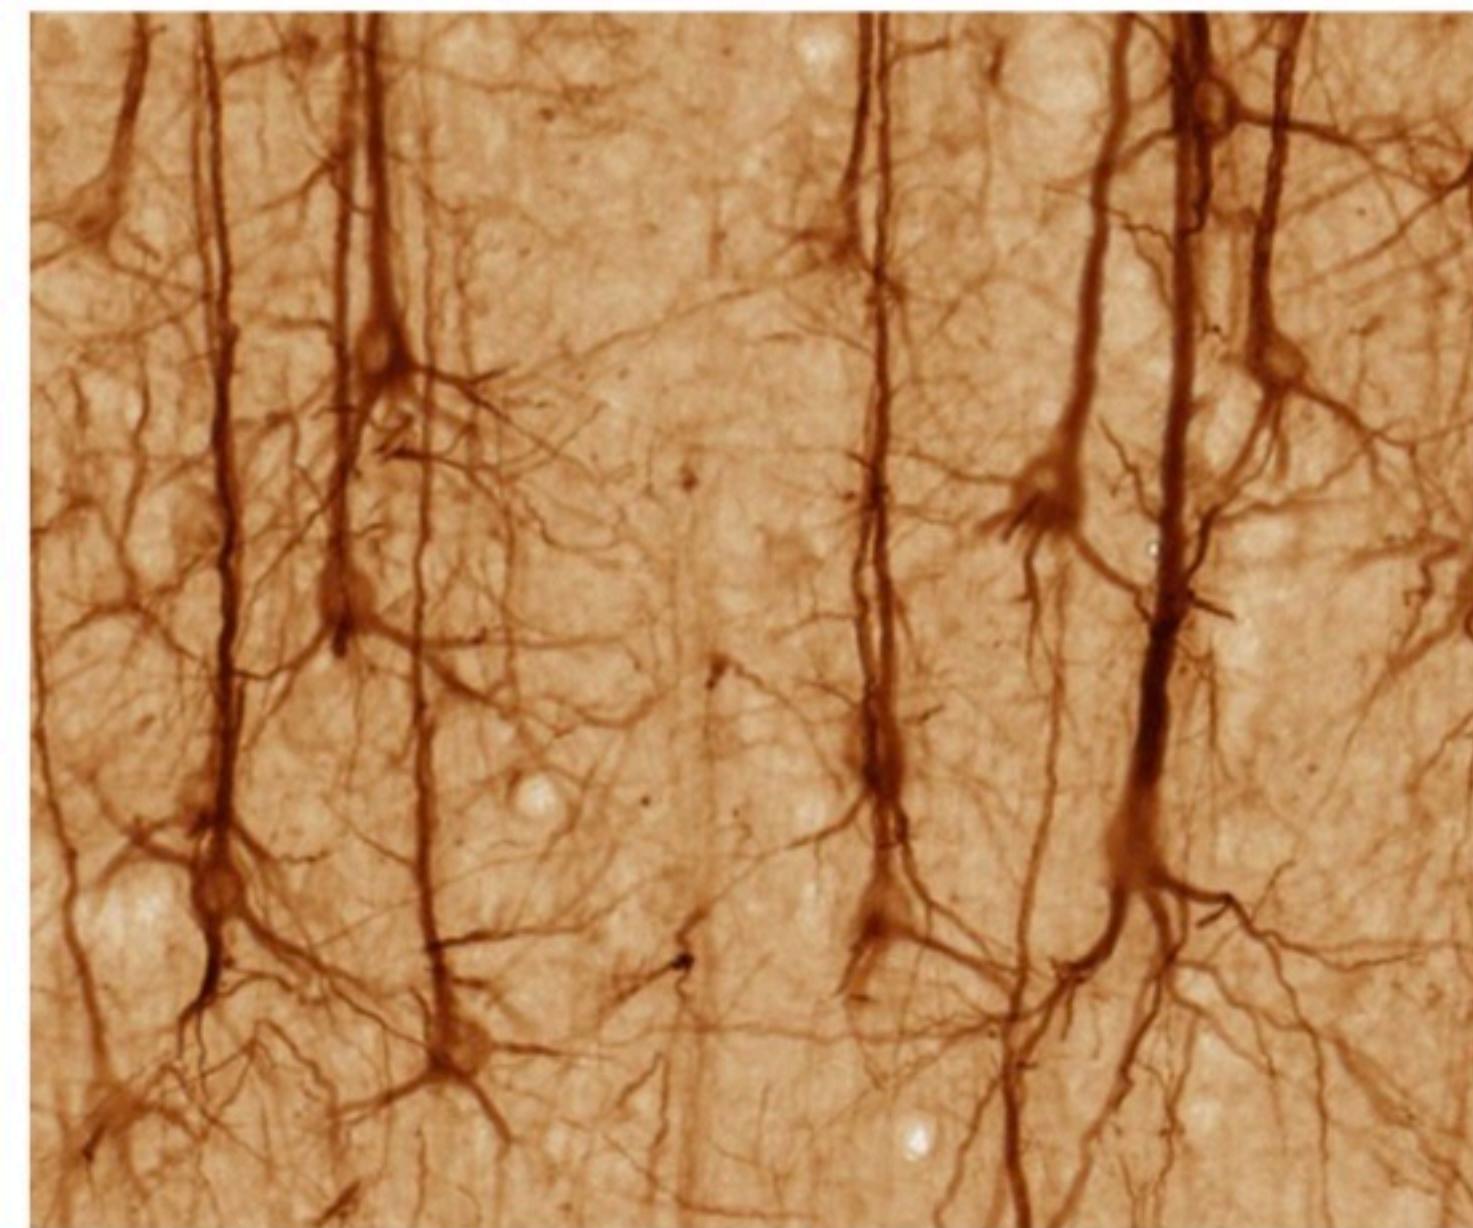


Perceptron Models



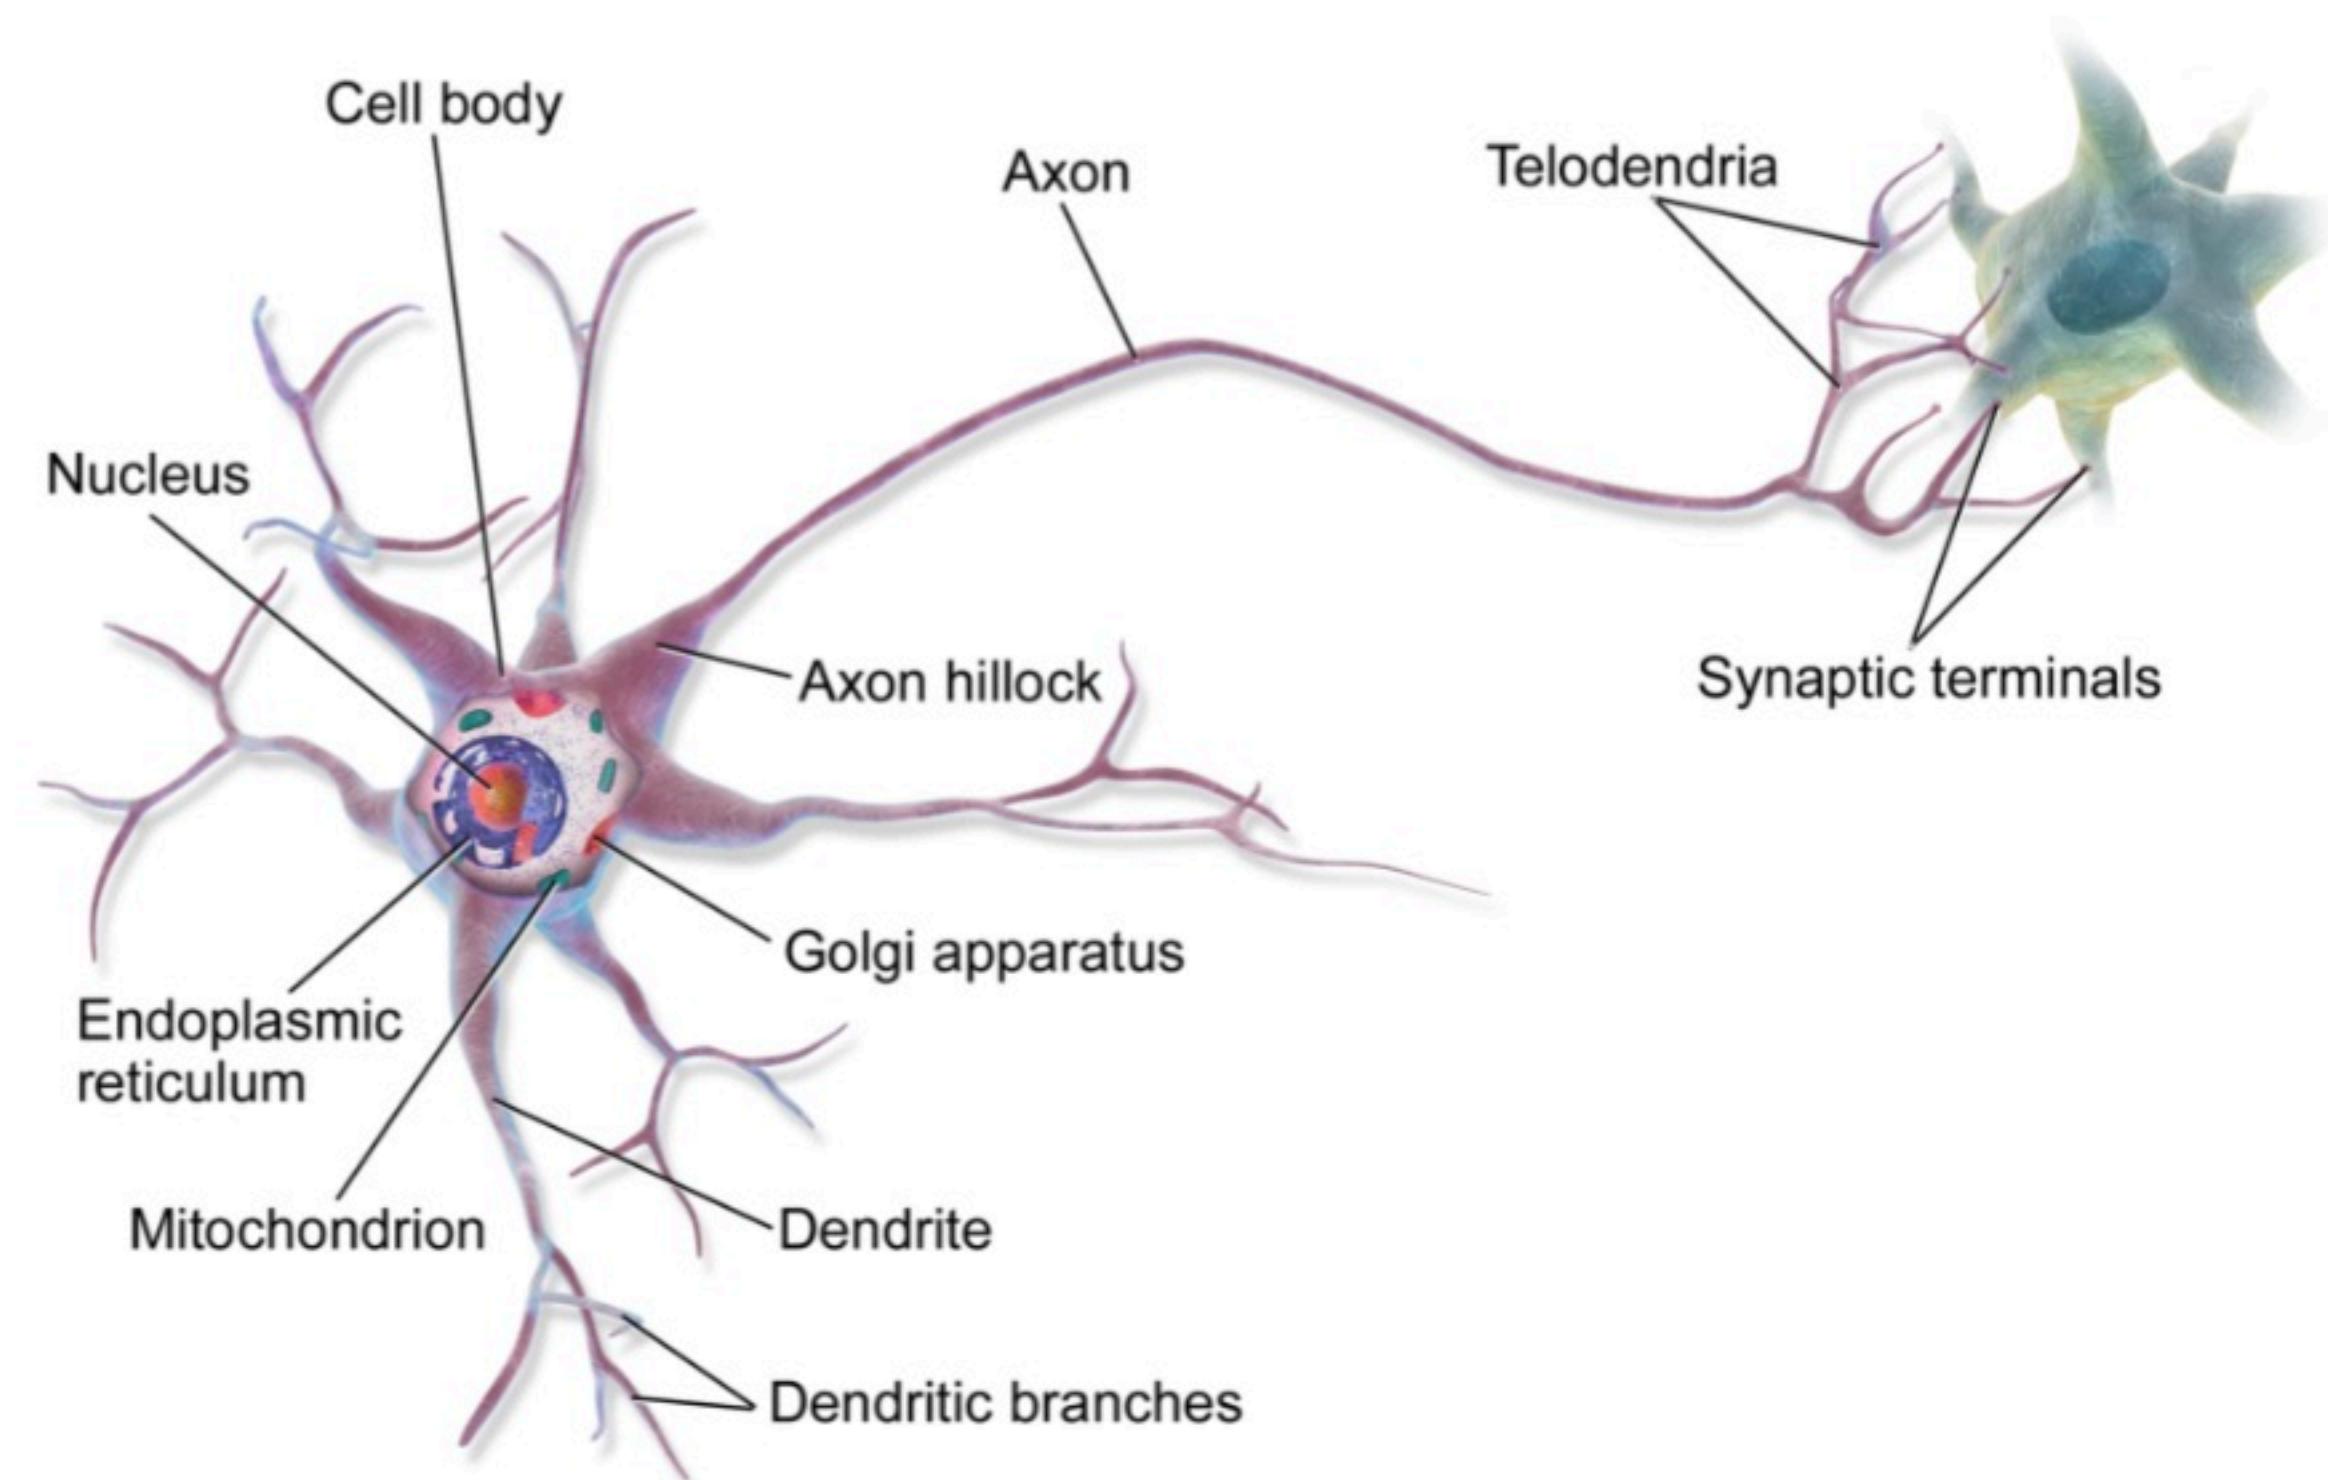
Single Biological Neuron

Neurons in the Cortex



Single Biological Neuron

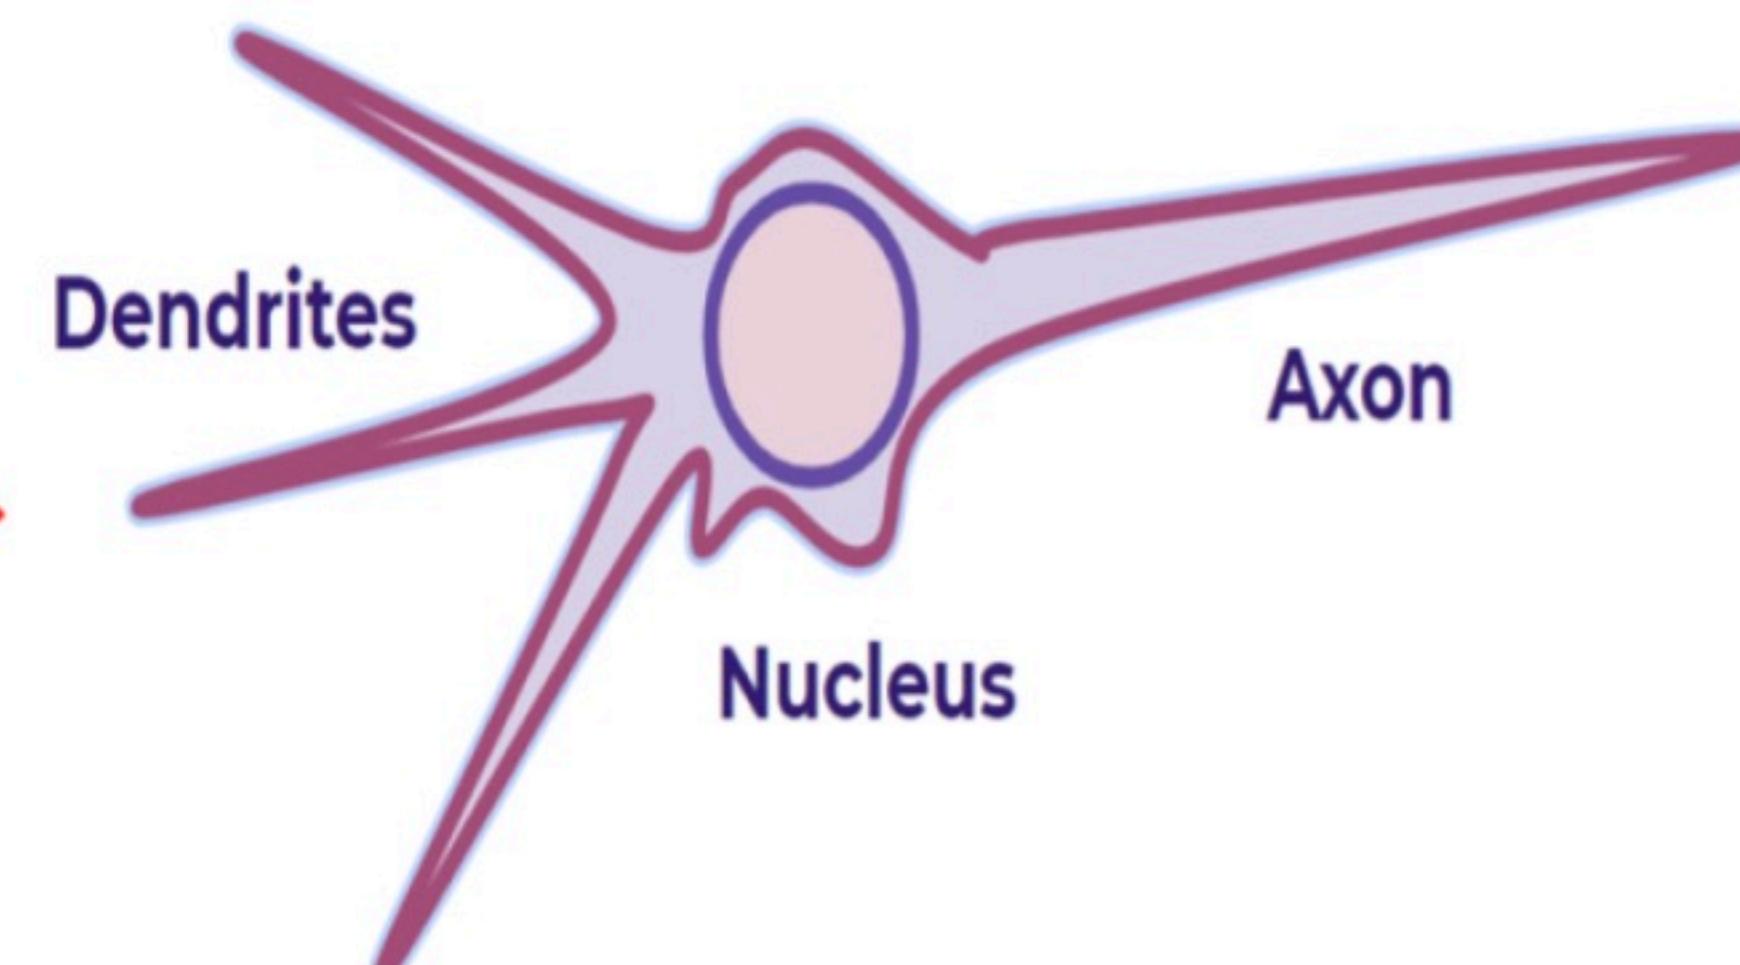
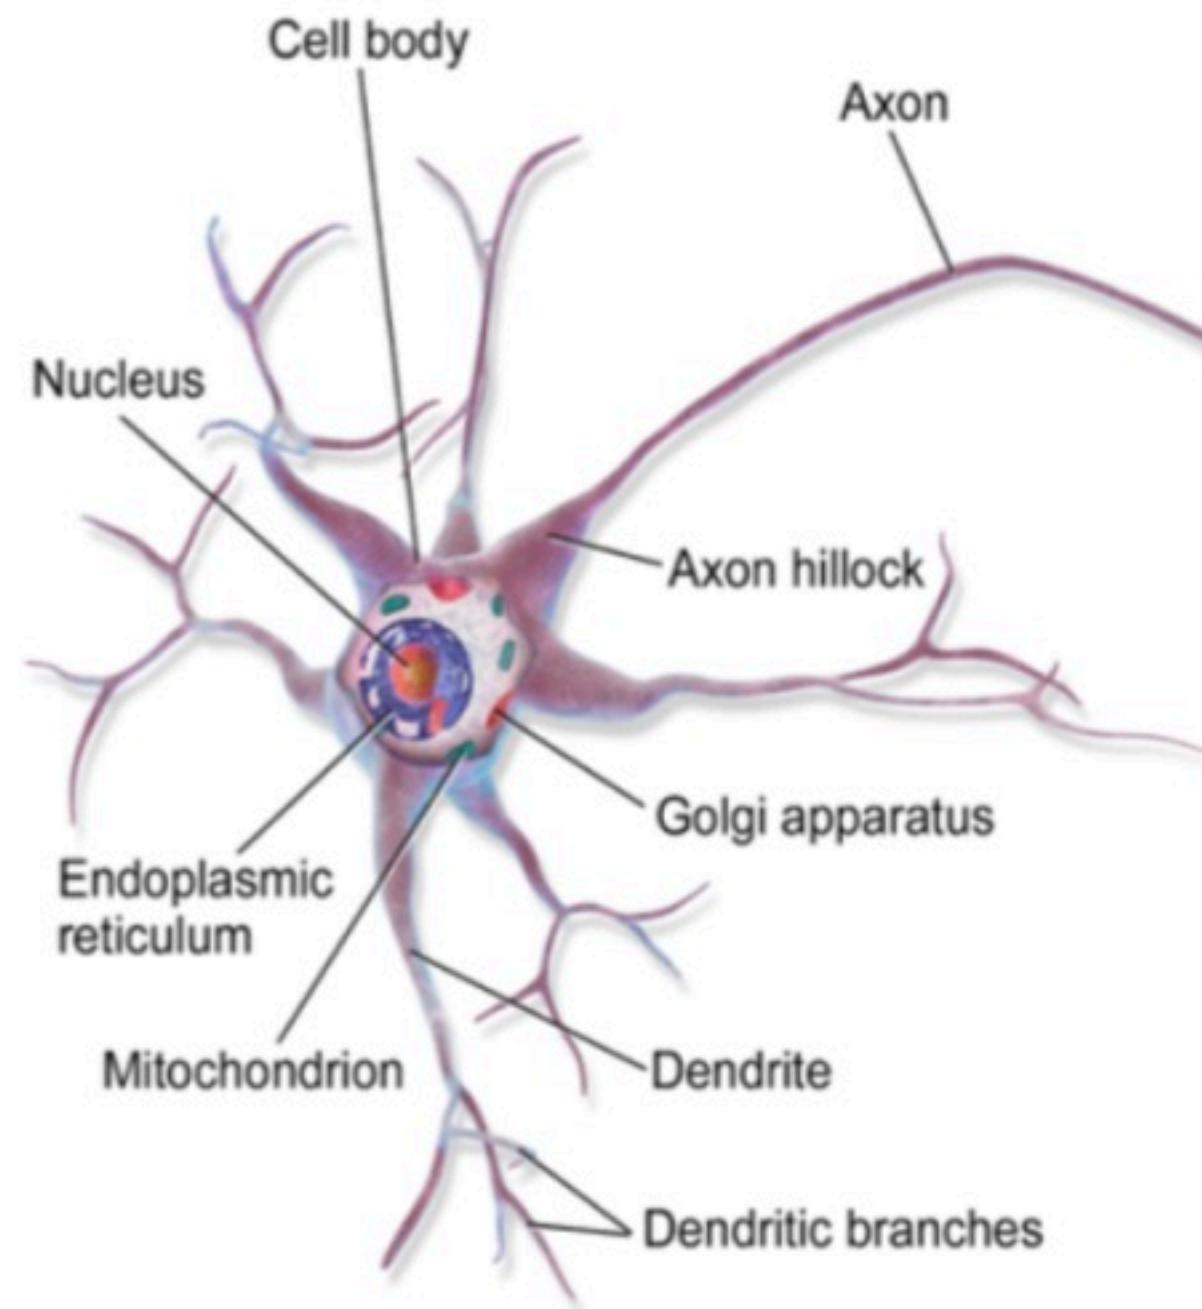
A Single Biological Neuron



Perceptron Models



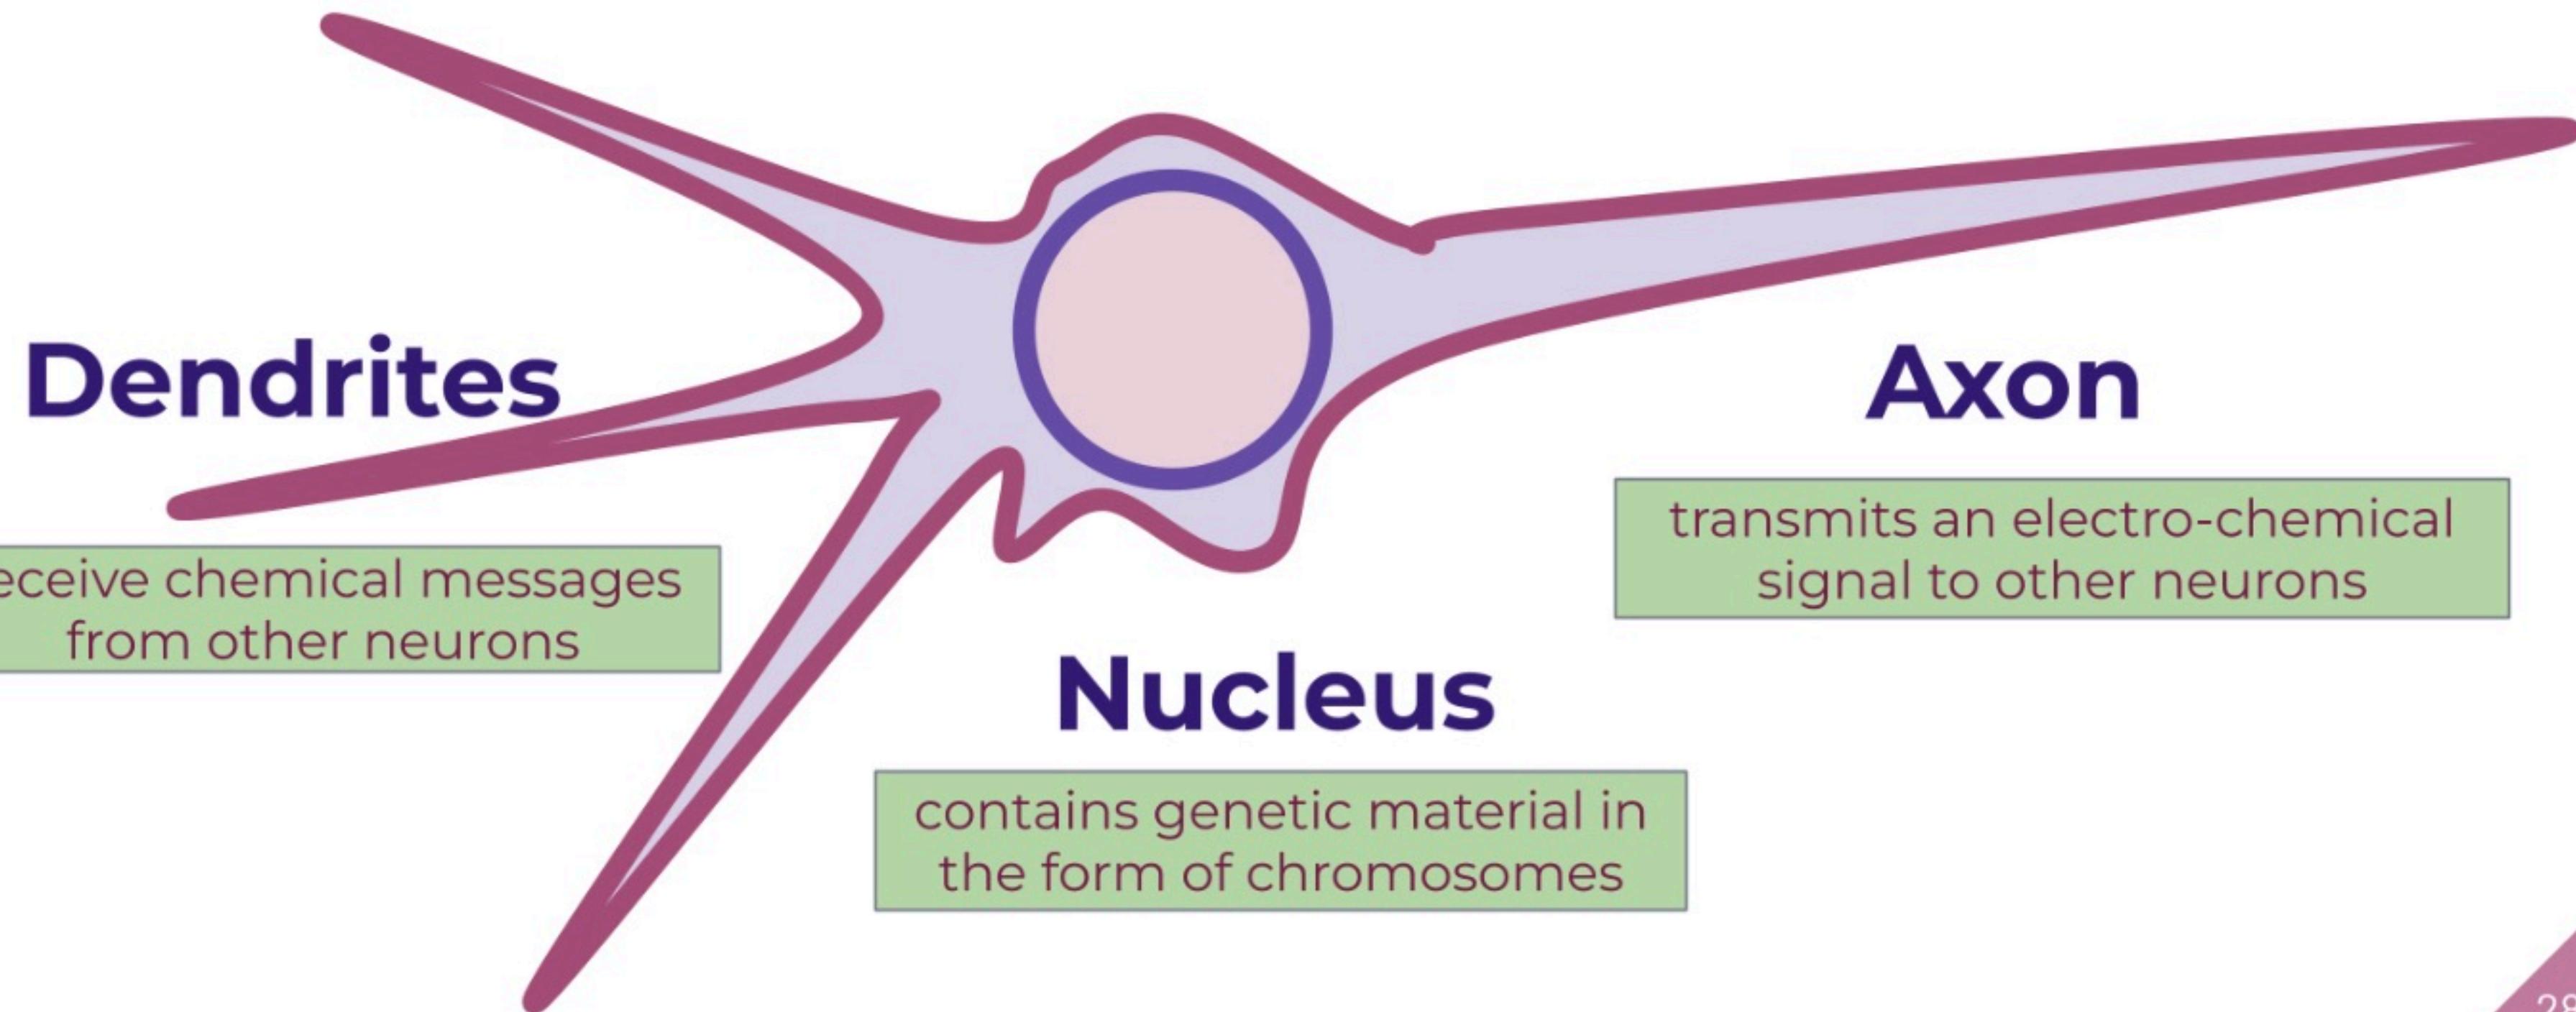
Simplifying the Single Biological Neuron



Perceptron Models



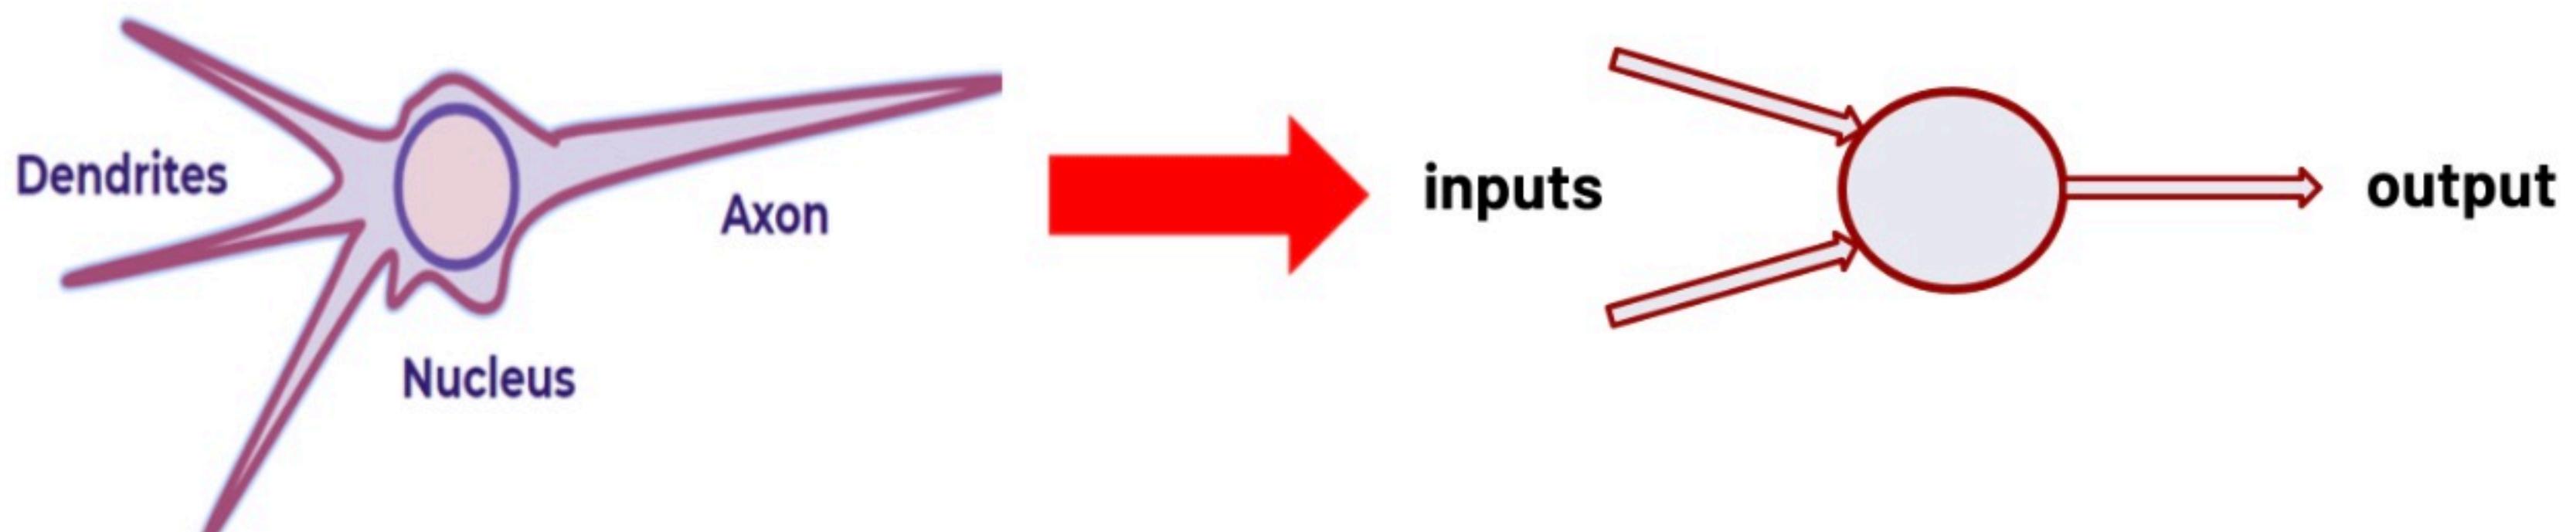
Simplified Biological Neuron Model



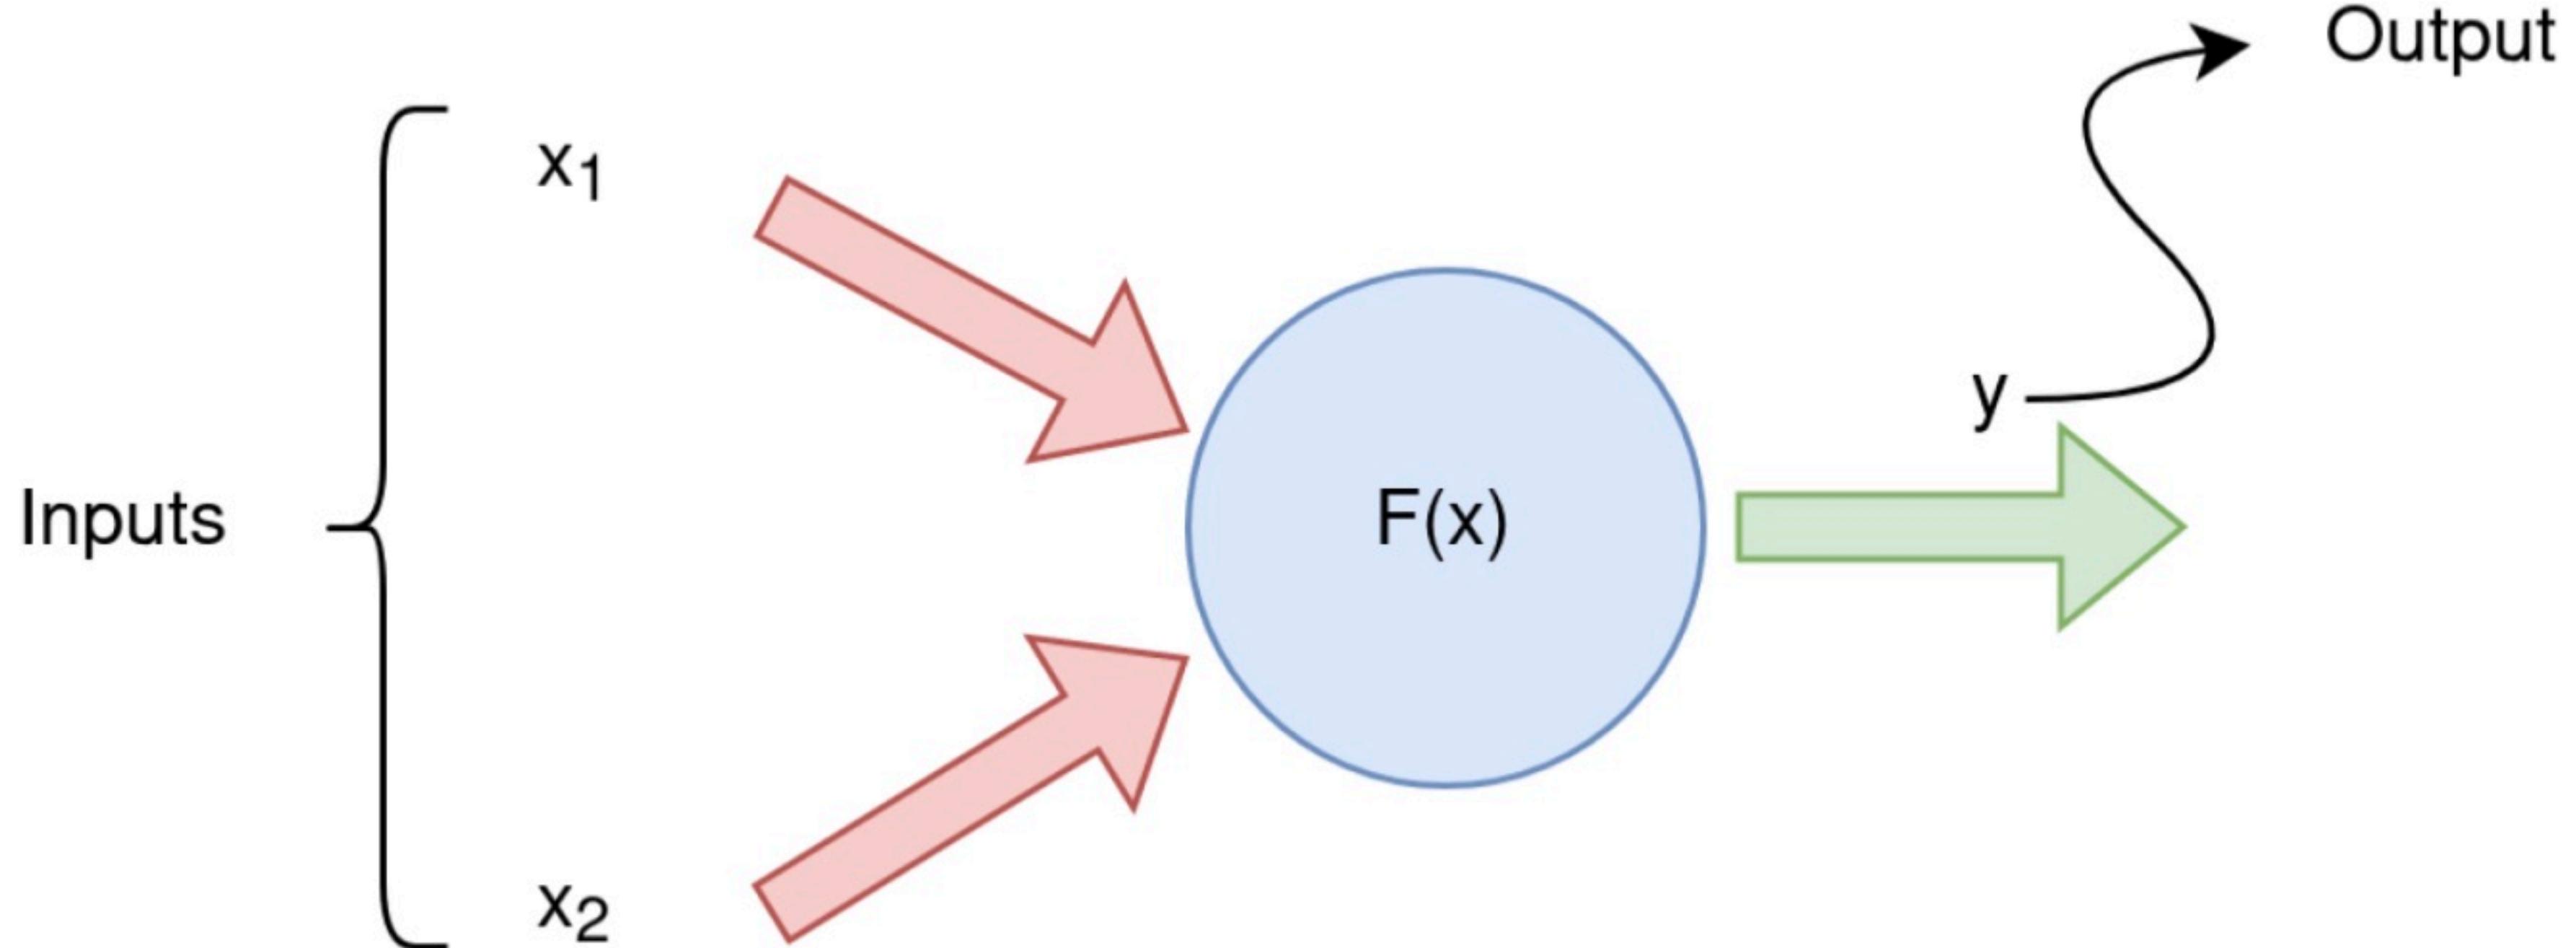
Perceptron Models



Converting the Simplified Single Biological Neuron to Perceptron Model

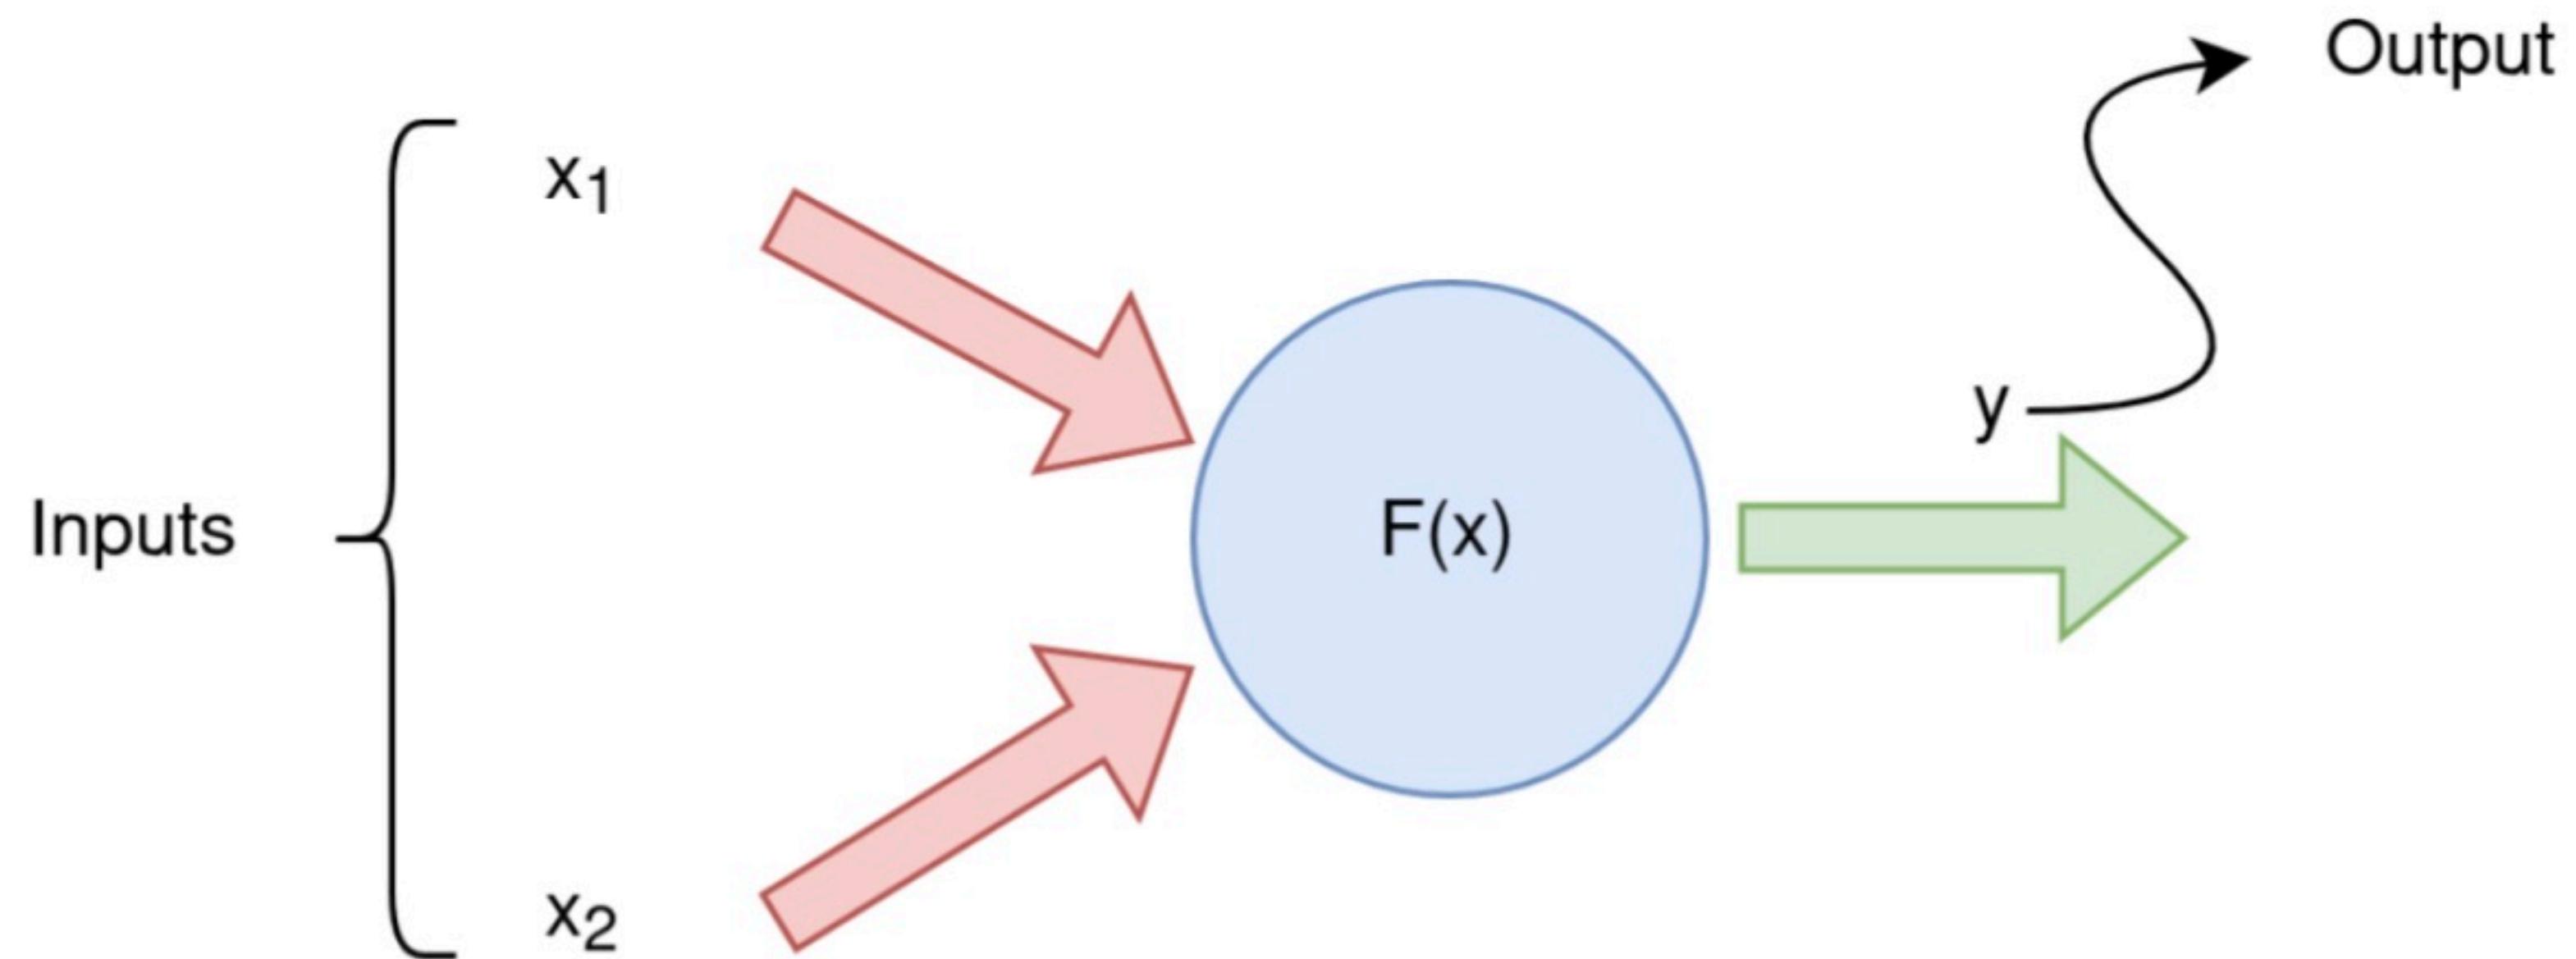


Perceptron Models



Perceptron Models

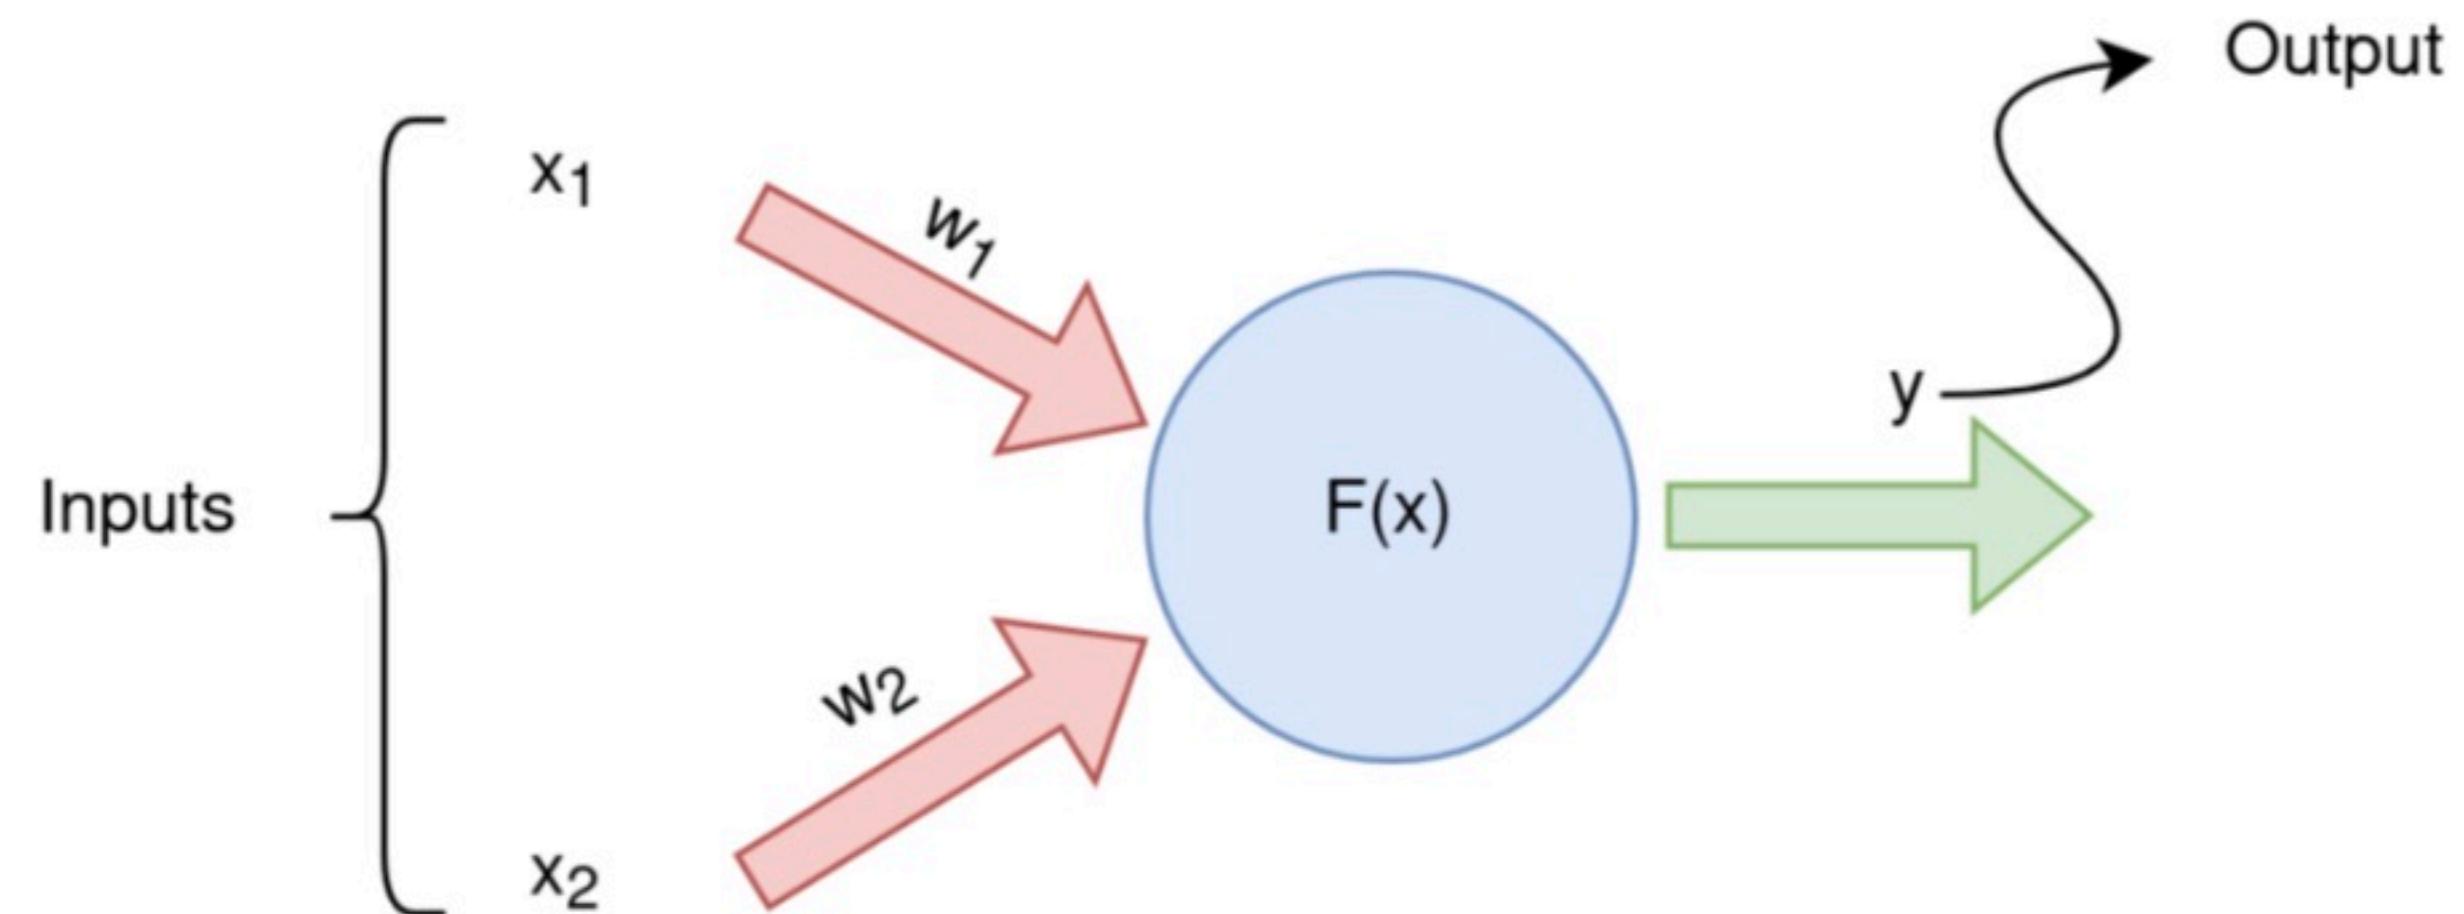
If it is a sum function, then; $y = x_1 + x_2$



Perceptron Models

How to insert another parameter in order to learn: w

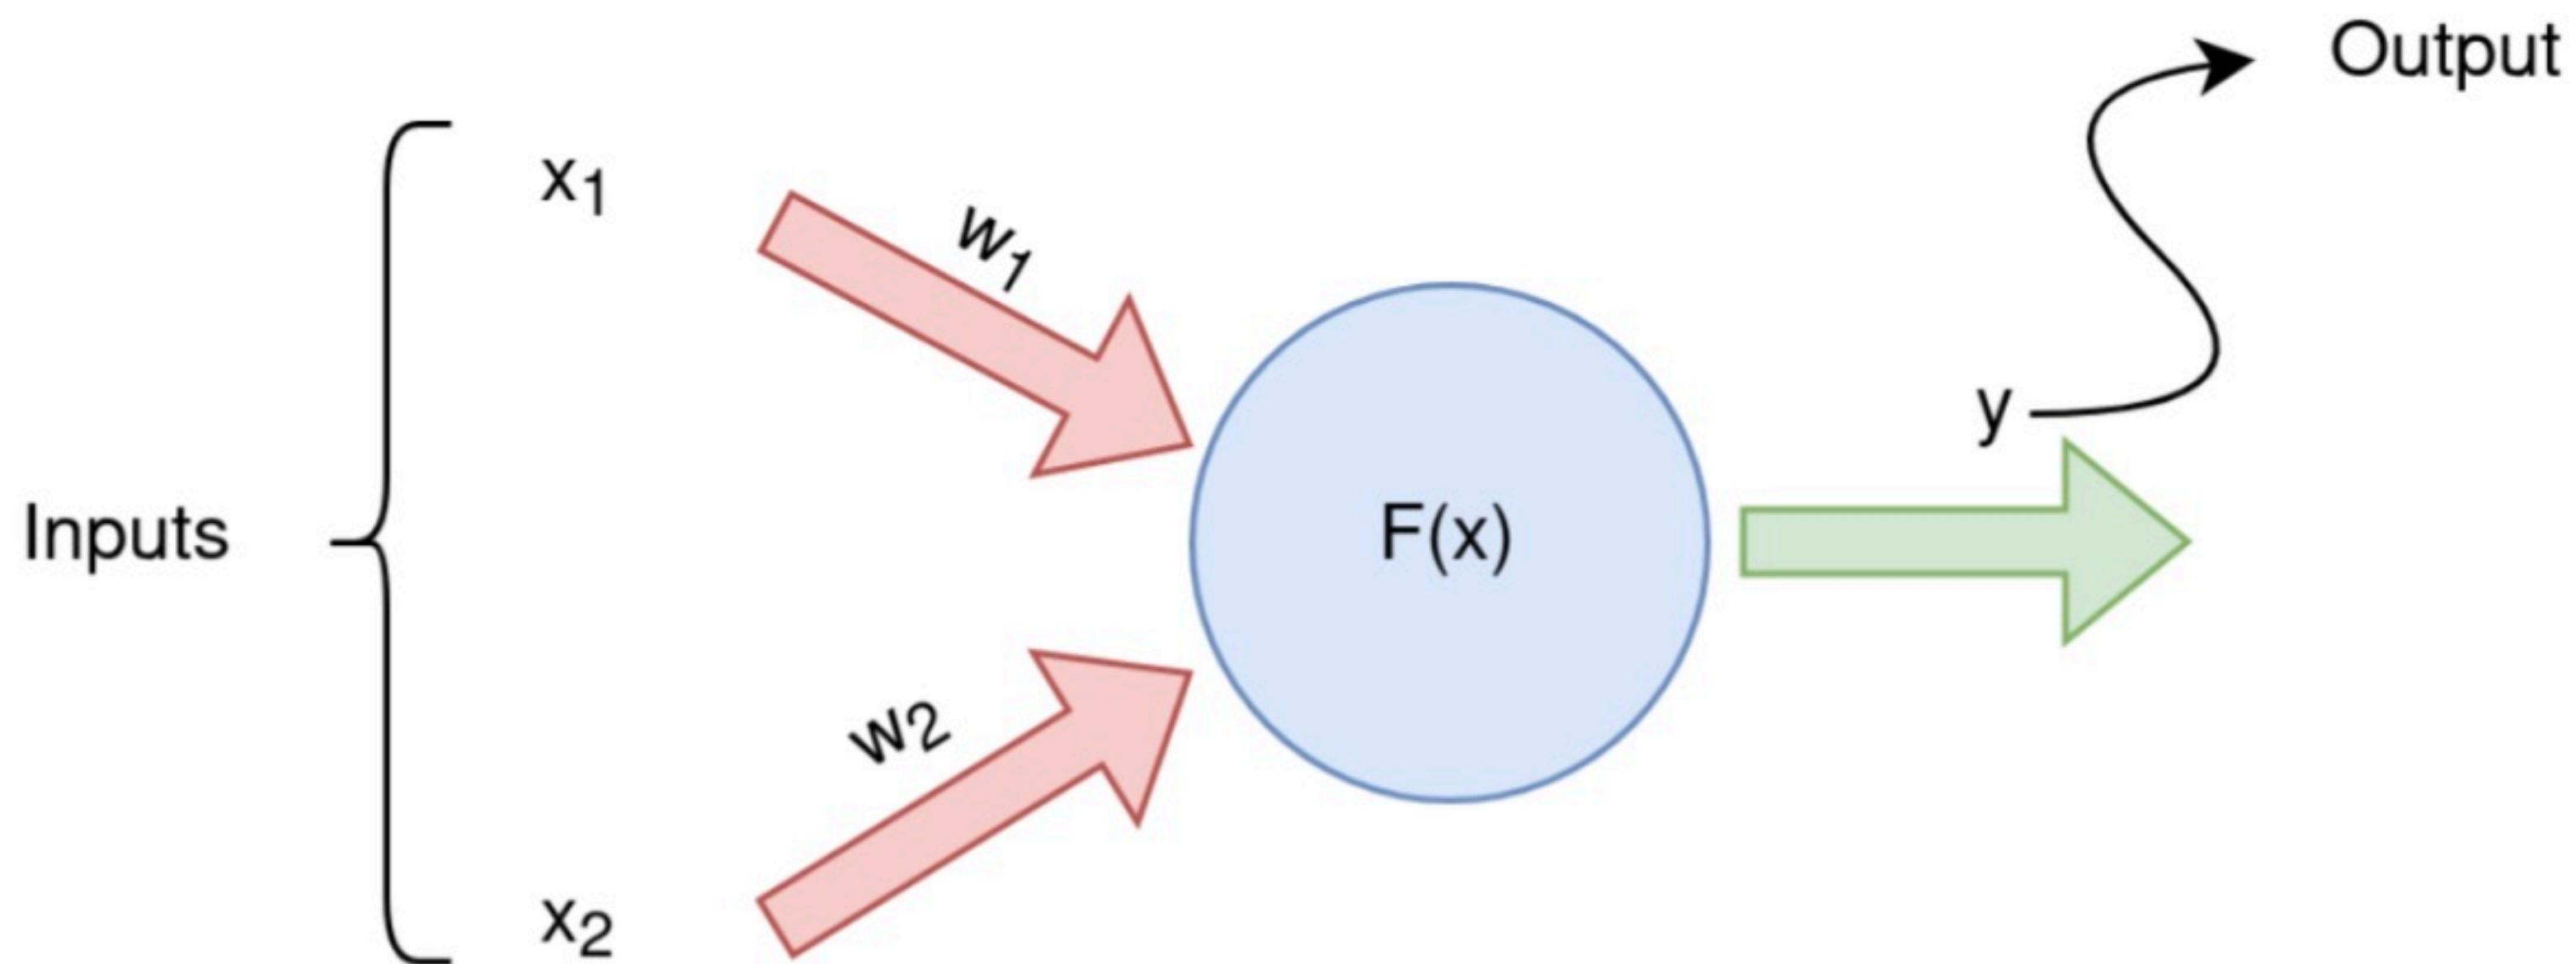
$$y = x_1 \cdot w_1 + x_2 \cdot w_2$$



Now we can update the **weights** to affect **y** (learning)

Perceptron Models

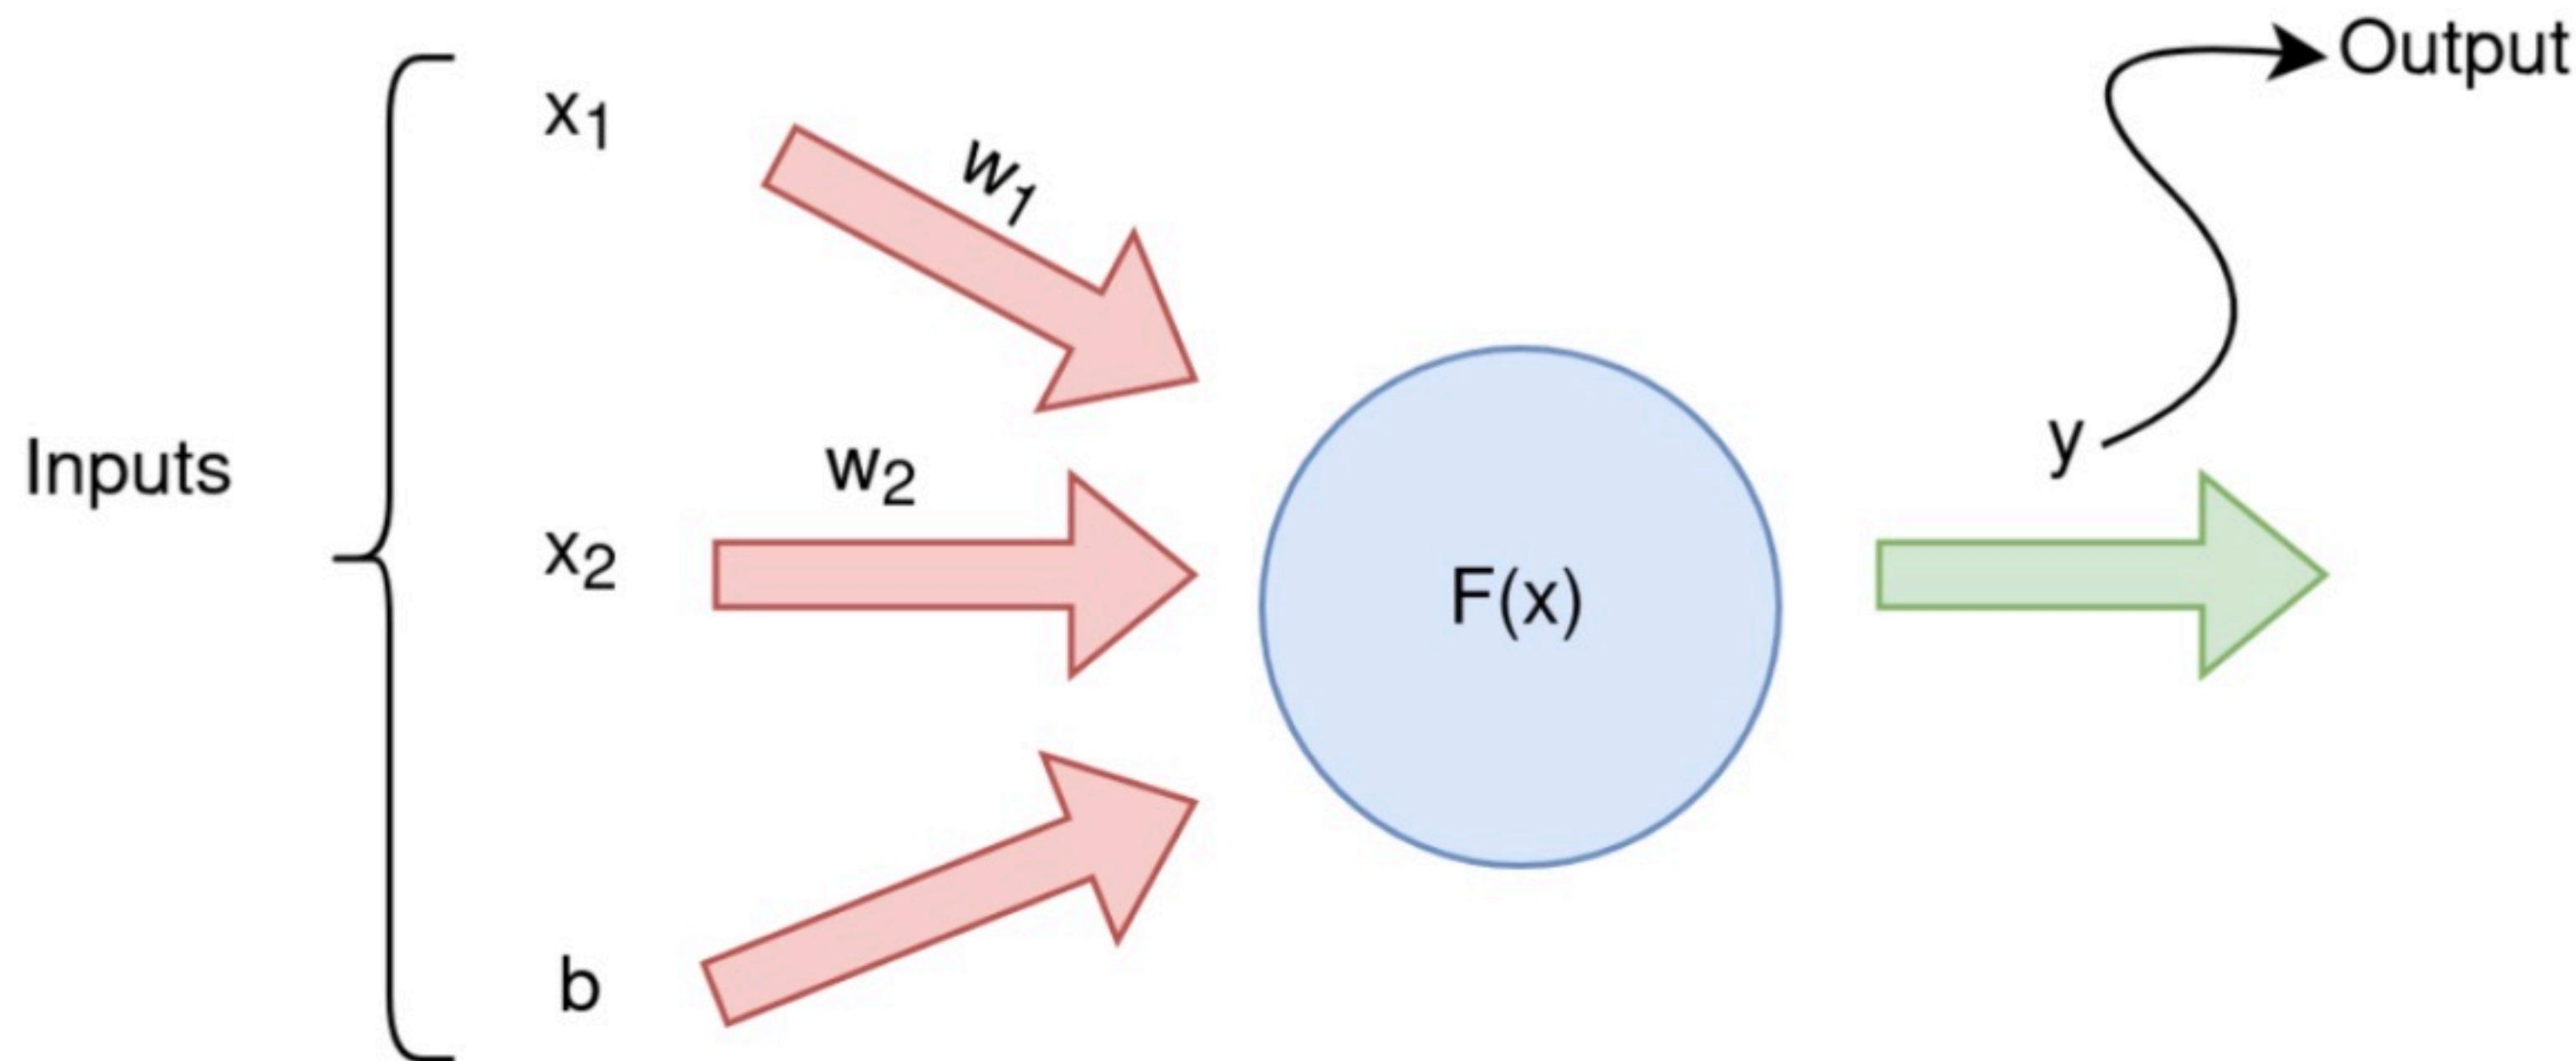
What if an \mathbf{x} is zero? \mathbf{w} can not change anything by itself



Perceptron Models

We should add a **bias** term **b** to the inputs.

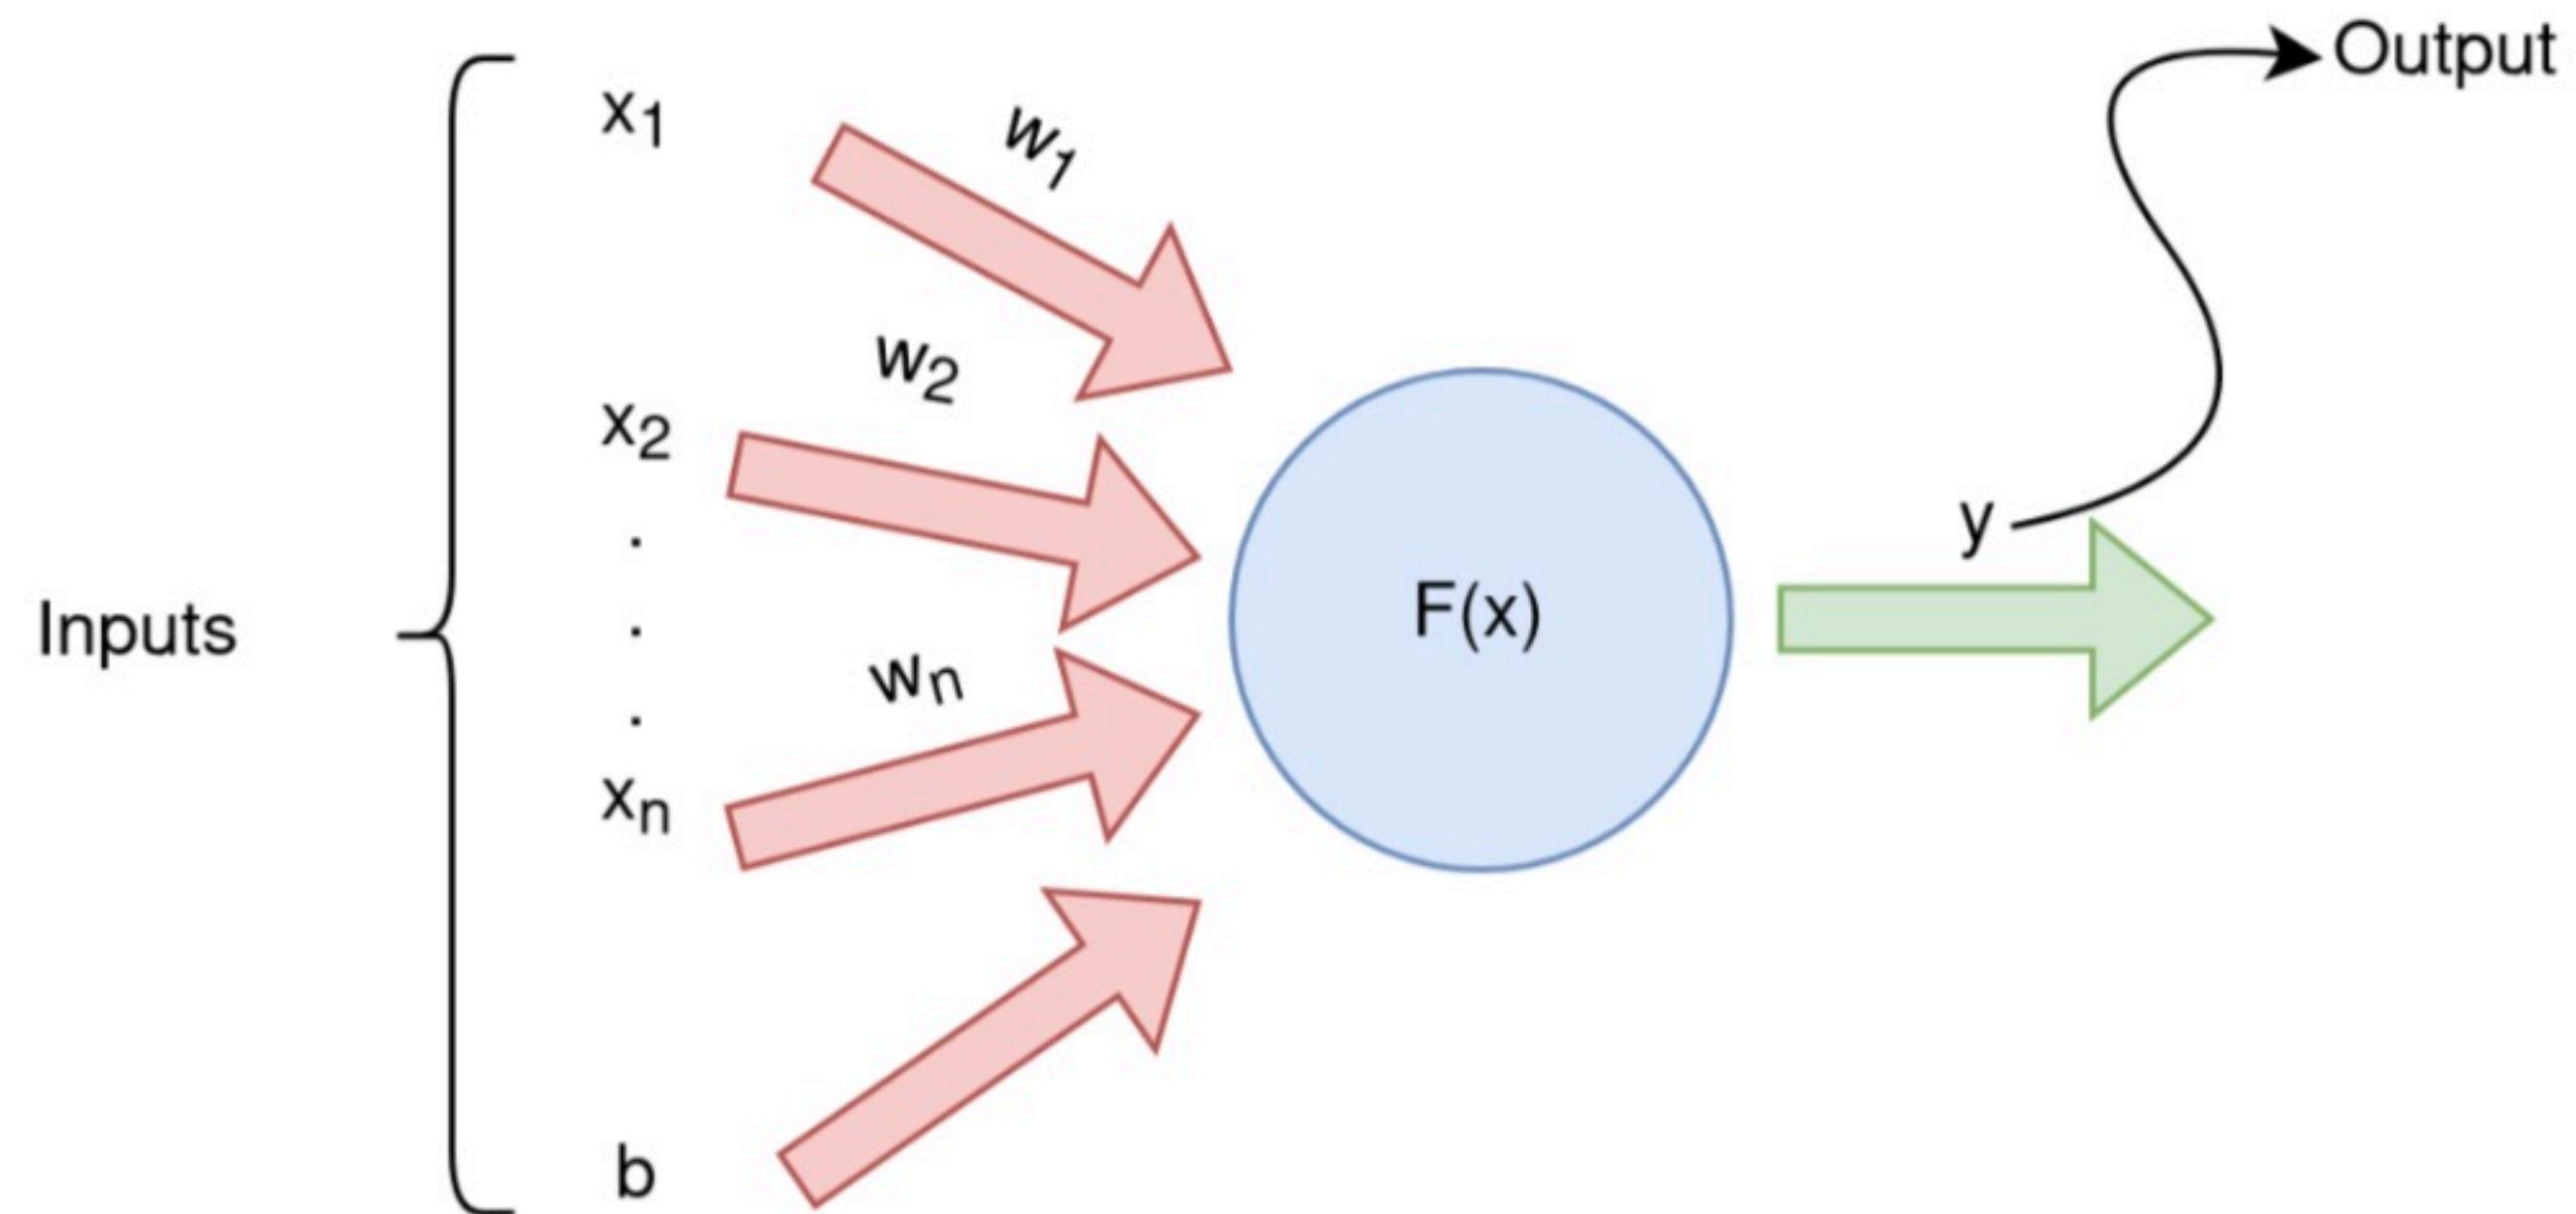
Now; $y = x_1 \cdot w_1 + x_2 \cdot w_2 + b$



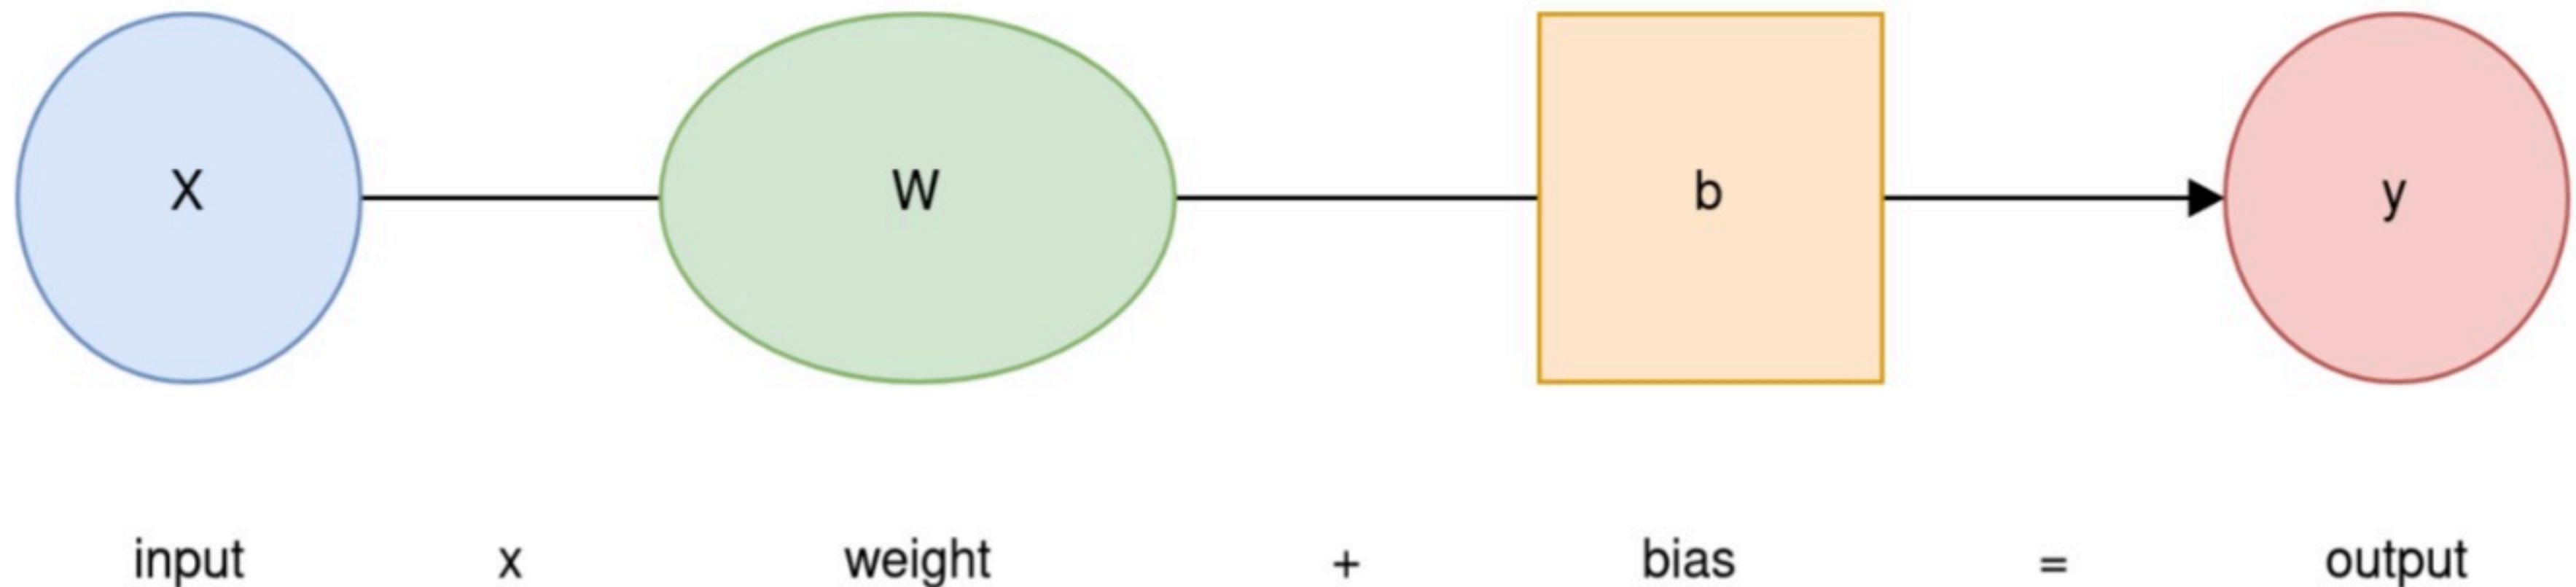
Perceptron Models

Generalization of the formula:

$$y = x_1 \cdot w_1 + x_2 \cdot w_2 + \dots + x_n \cdot w_n + b$$



Perceptron Models



The formula is: $z = \sum x_i \cdot w_i + b$ *b: bias*

$$z = x_1 \cdot w_1 + x_2 \cdot w_2 + x_3 \cdot w_3 + \dots + x_n \cdot w_n + b$$

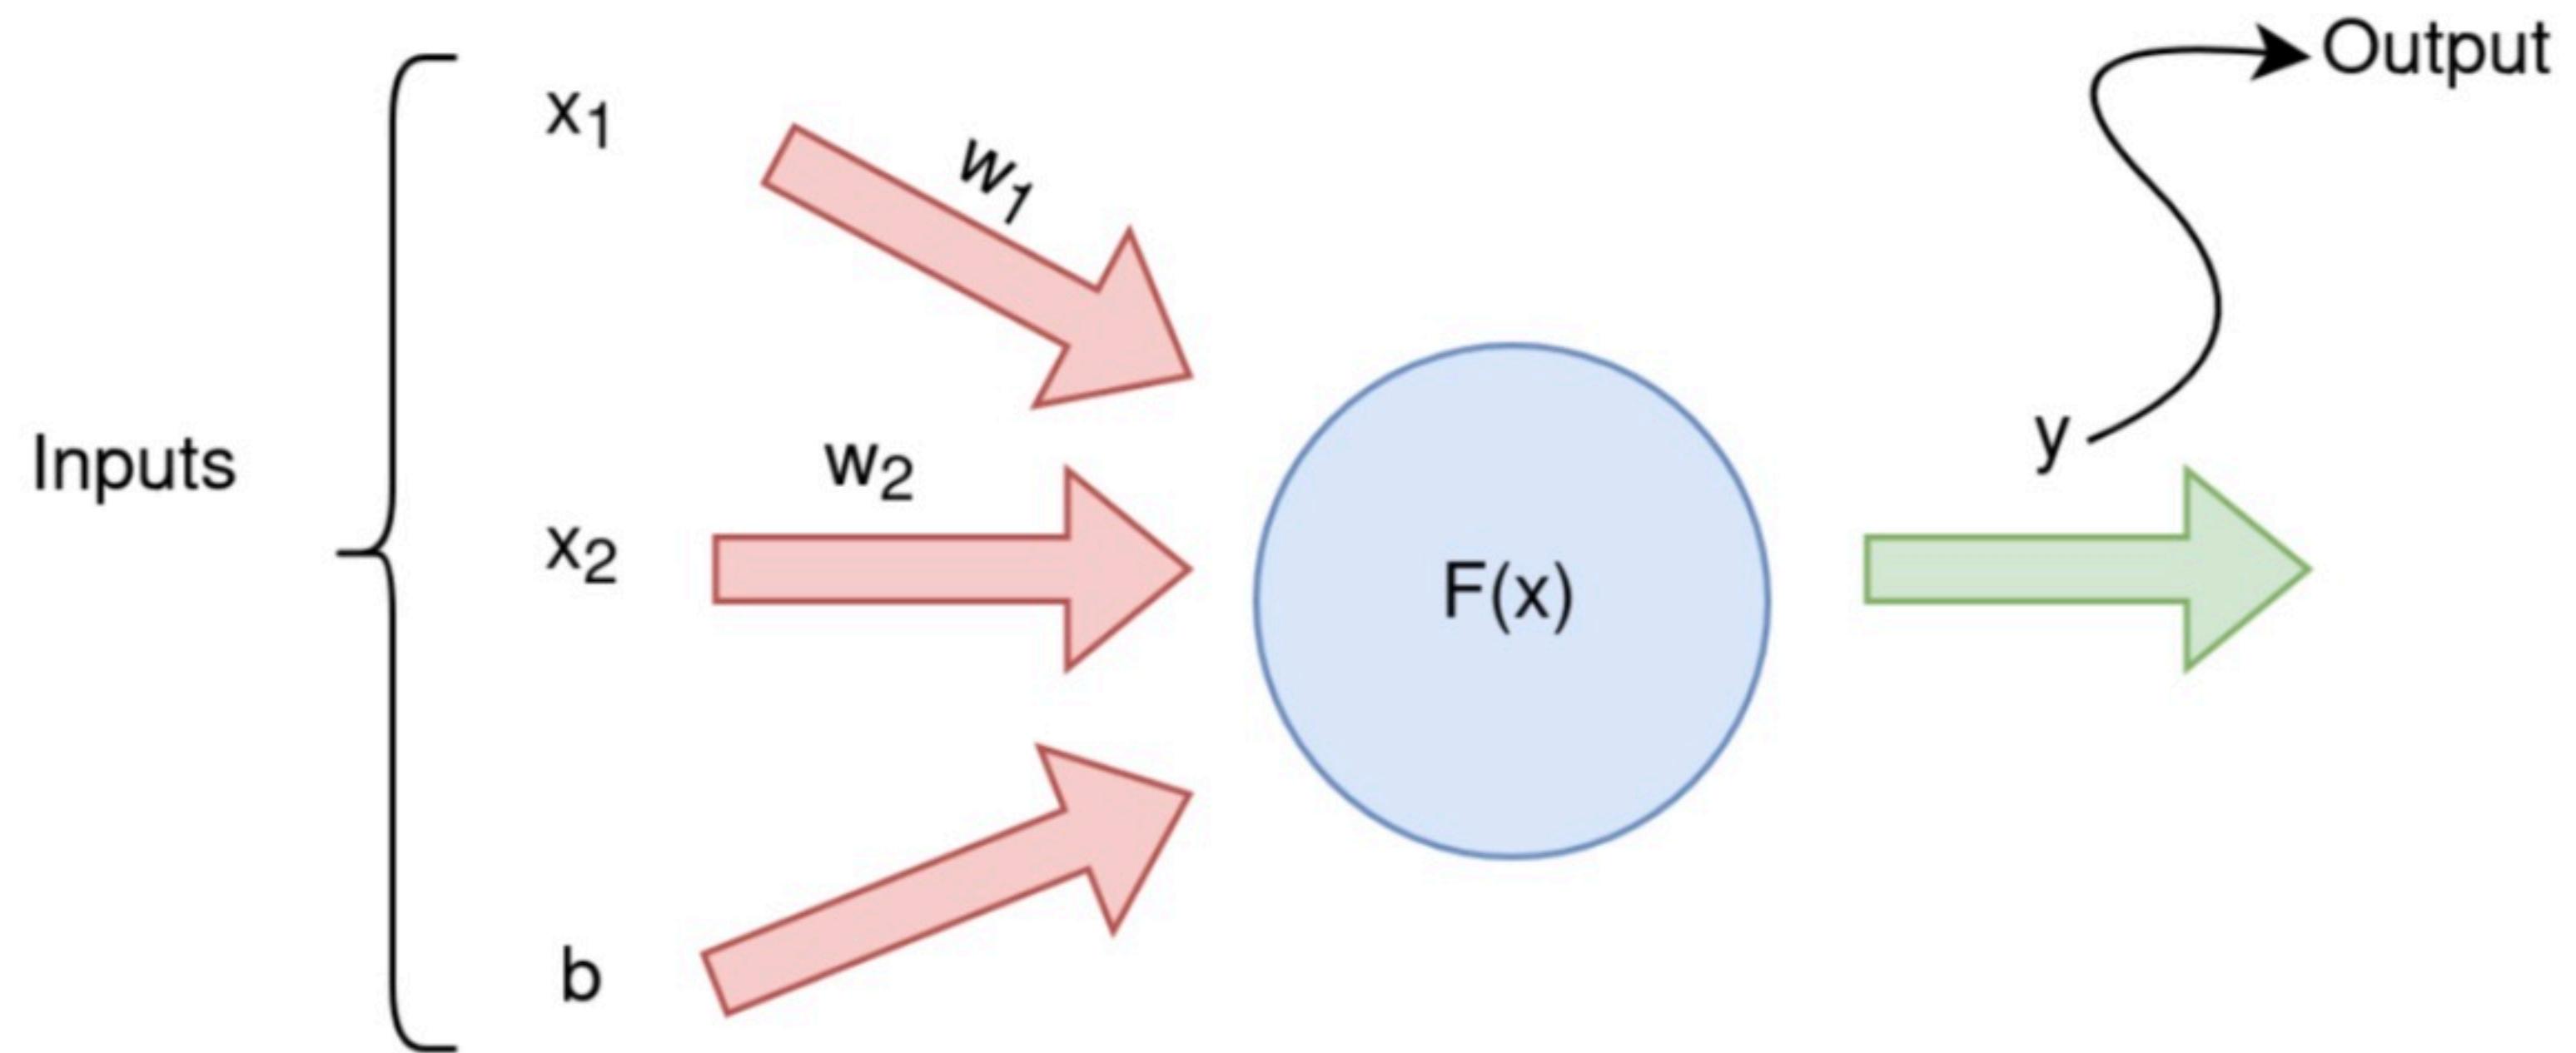


Neural Networks

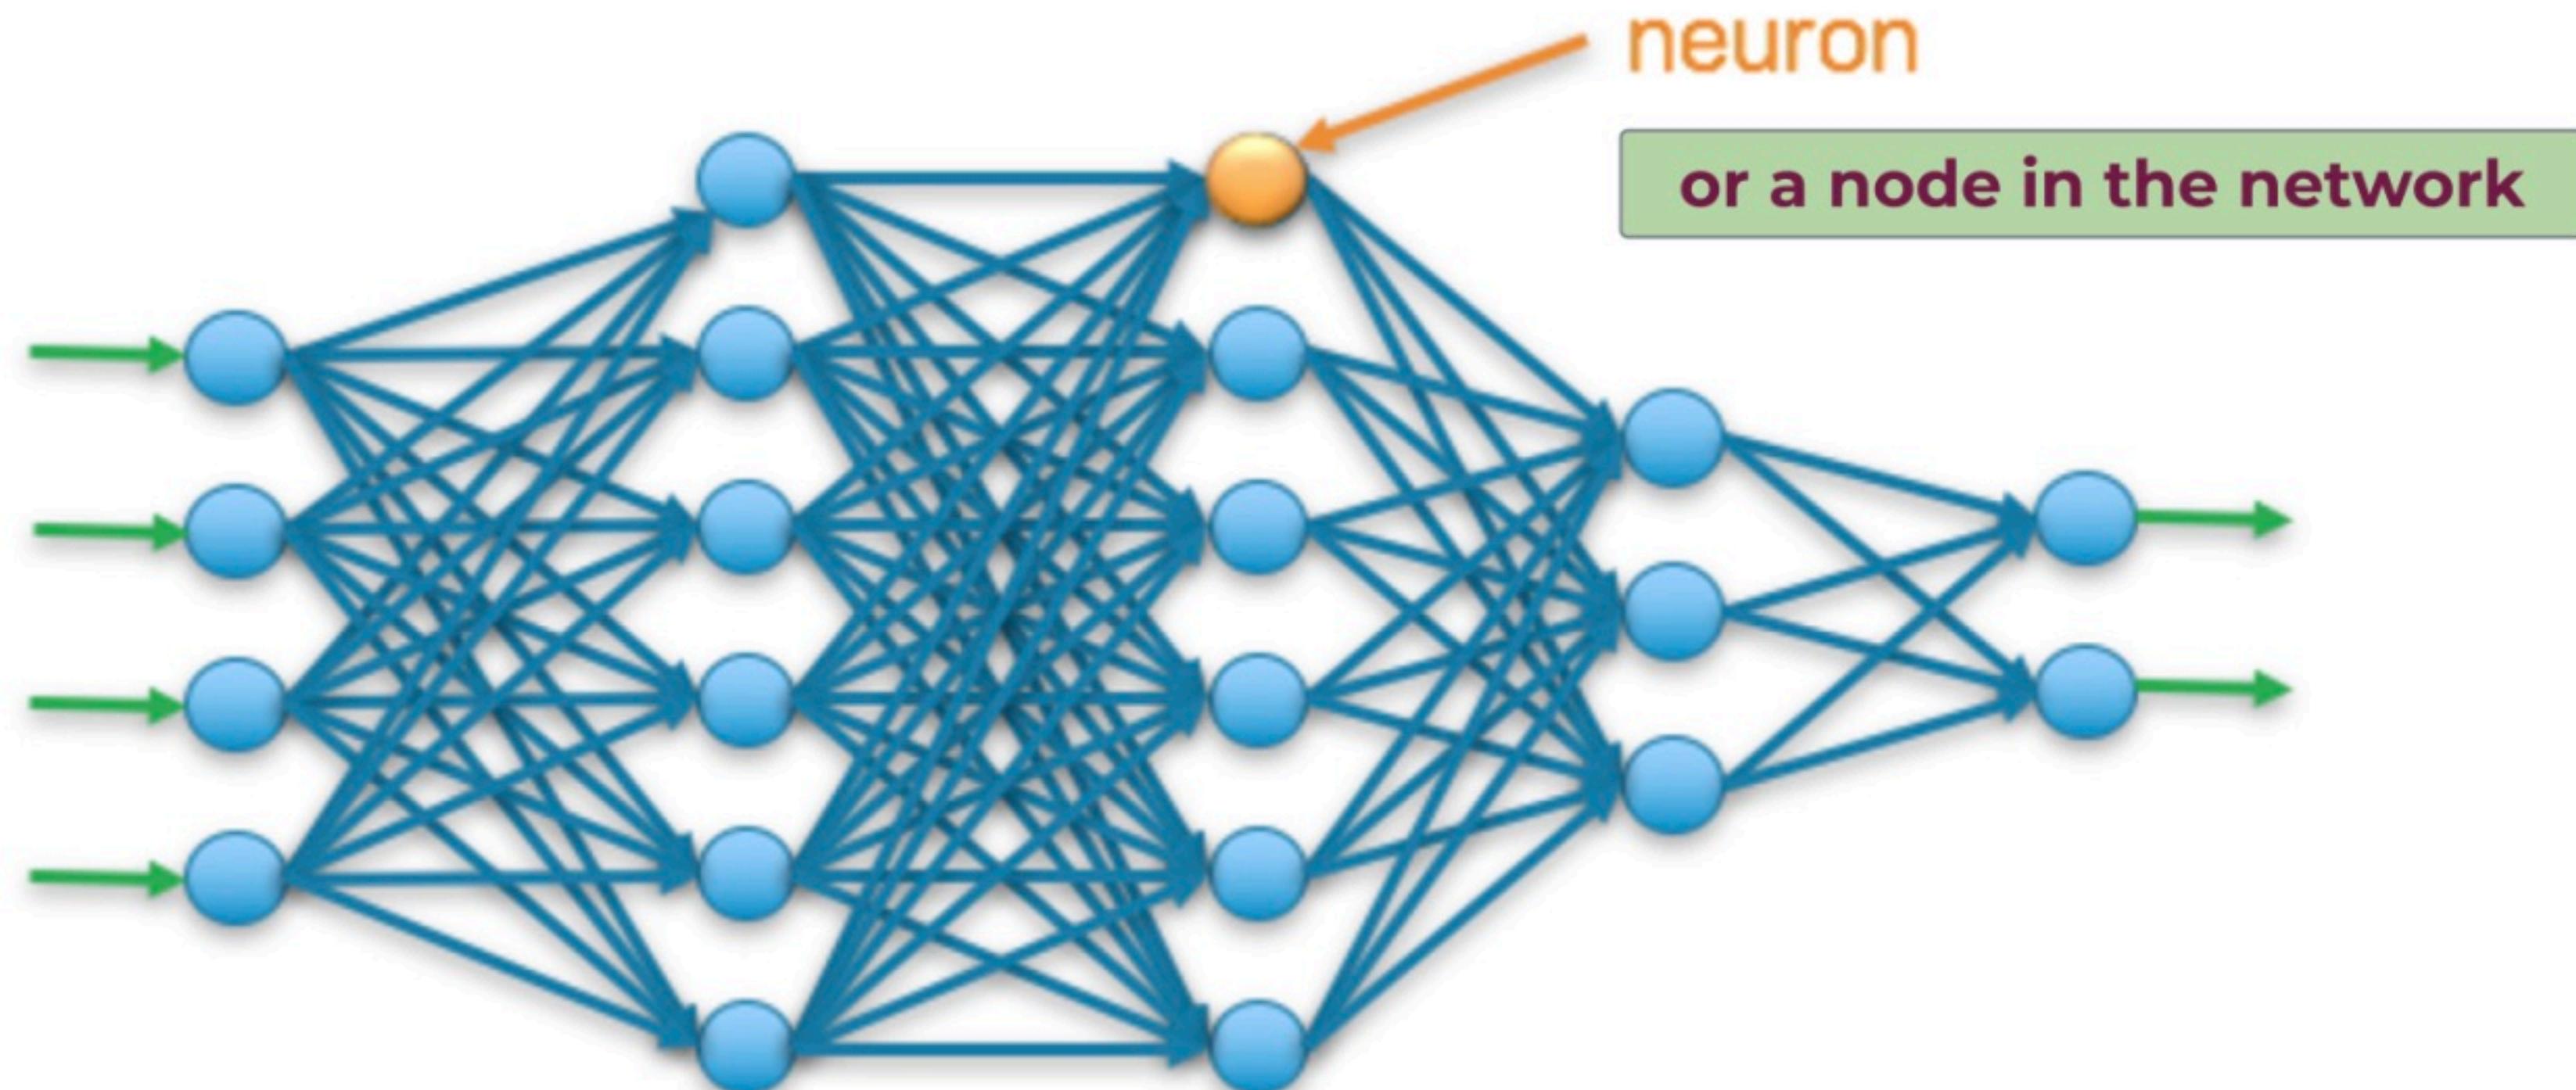


What is a Neuron ?

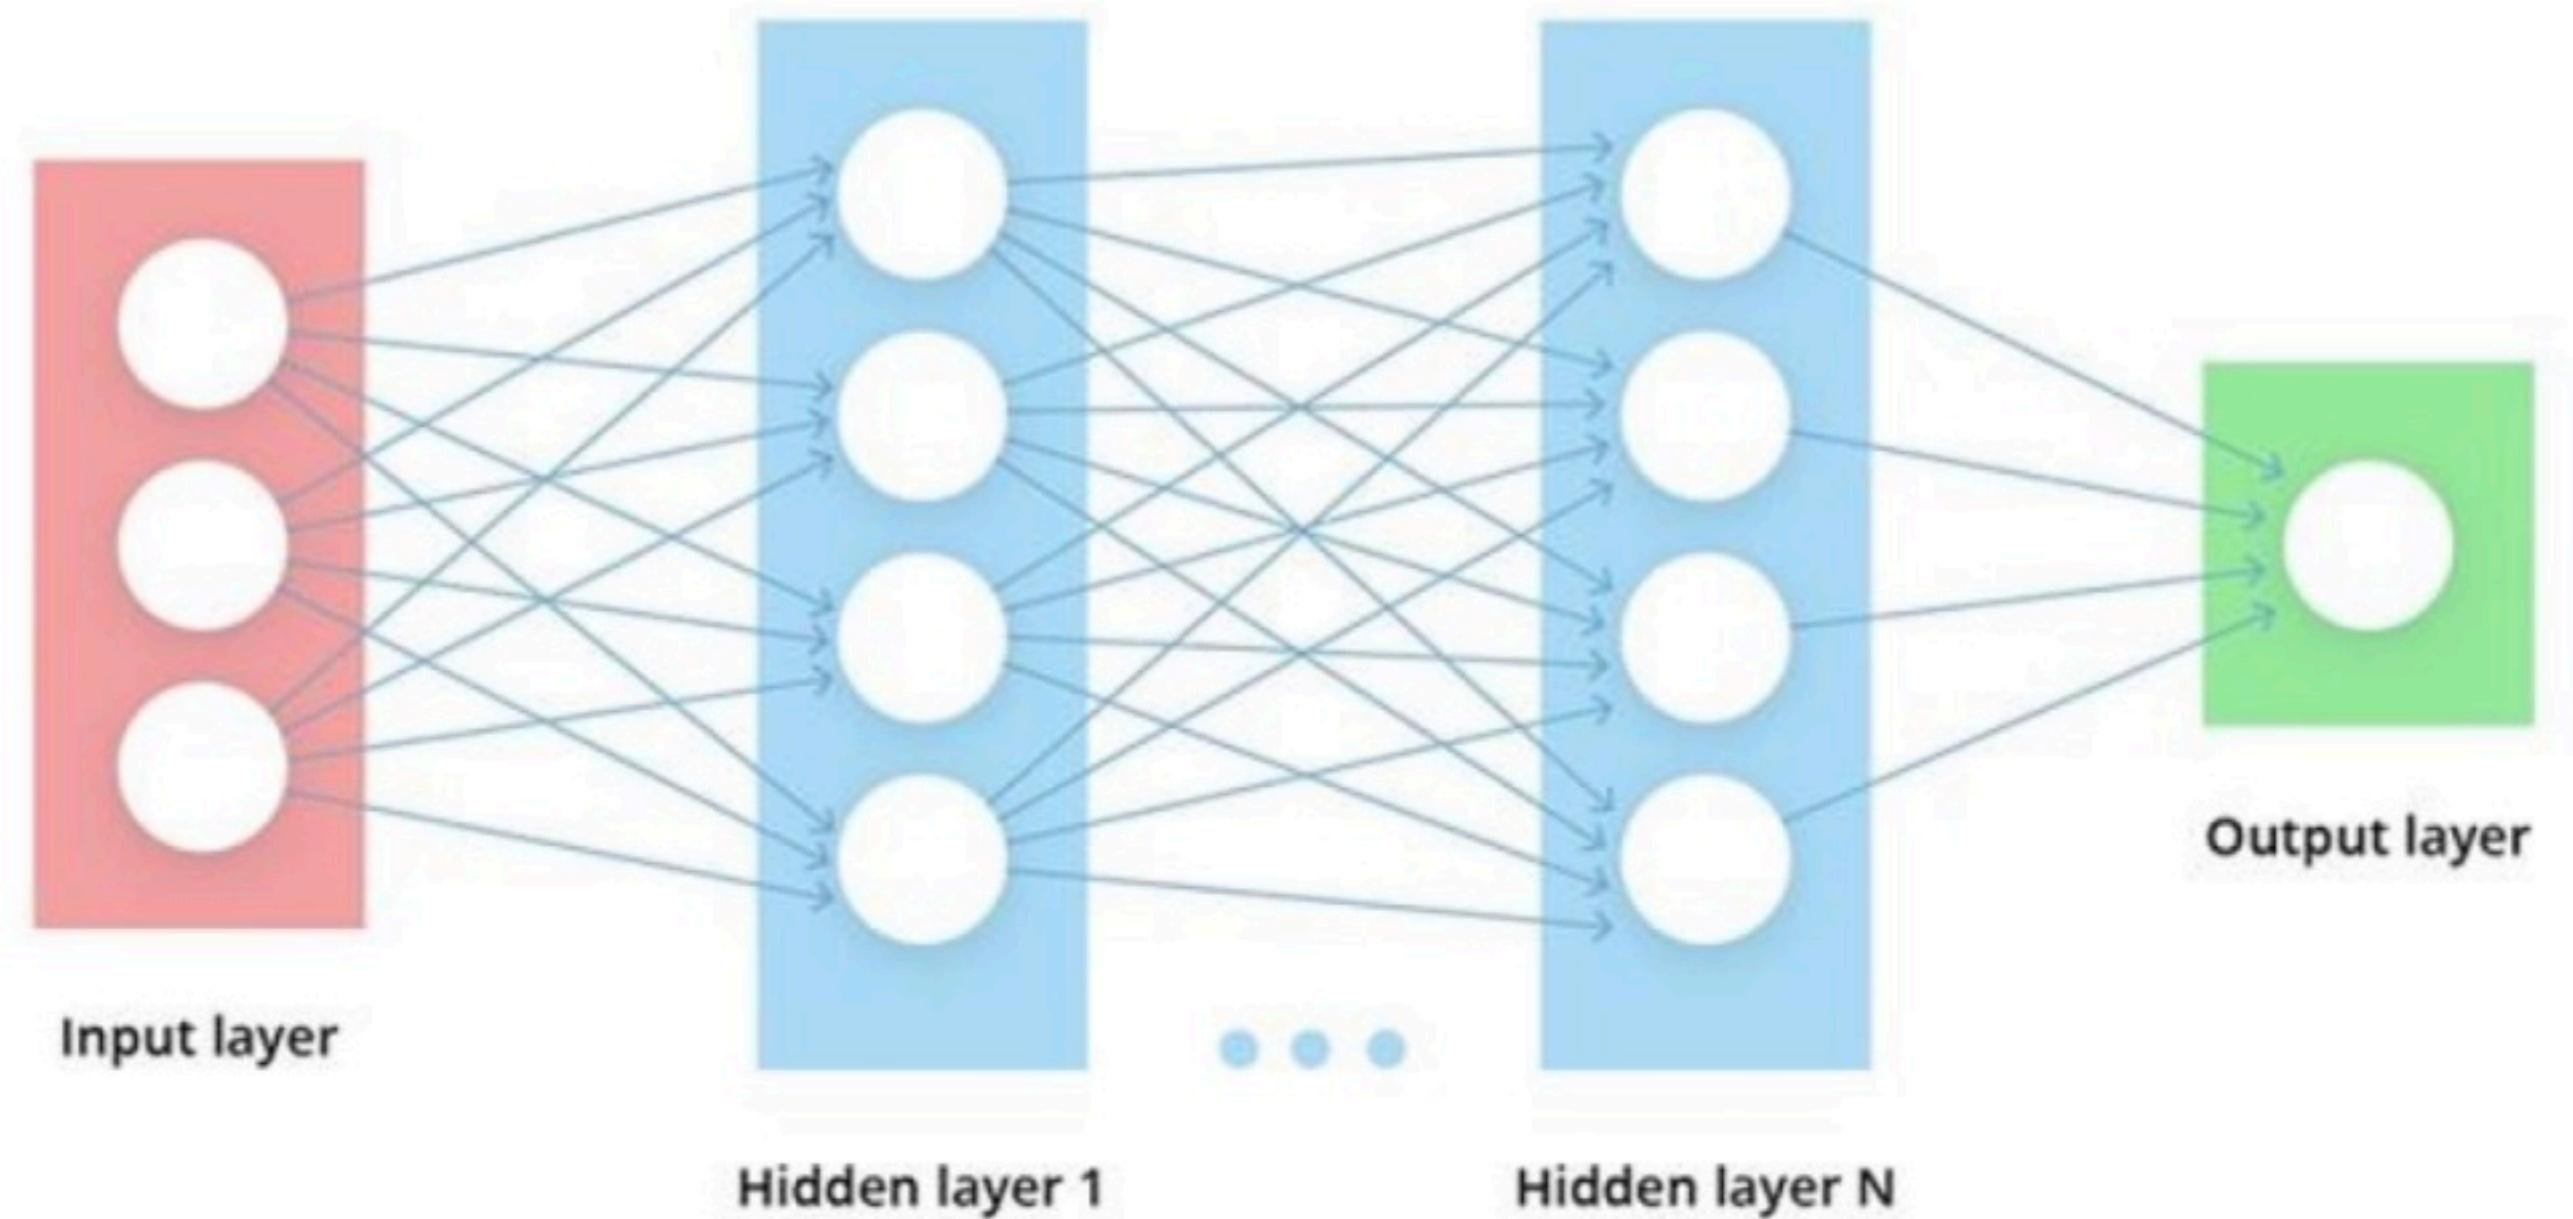
A single perceptron model = A neuron = A node in a network



Neurons in a Network



Layers



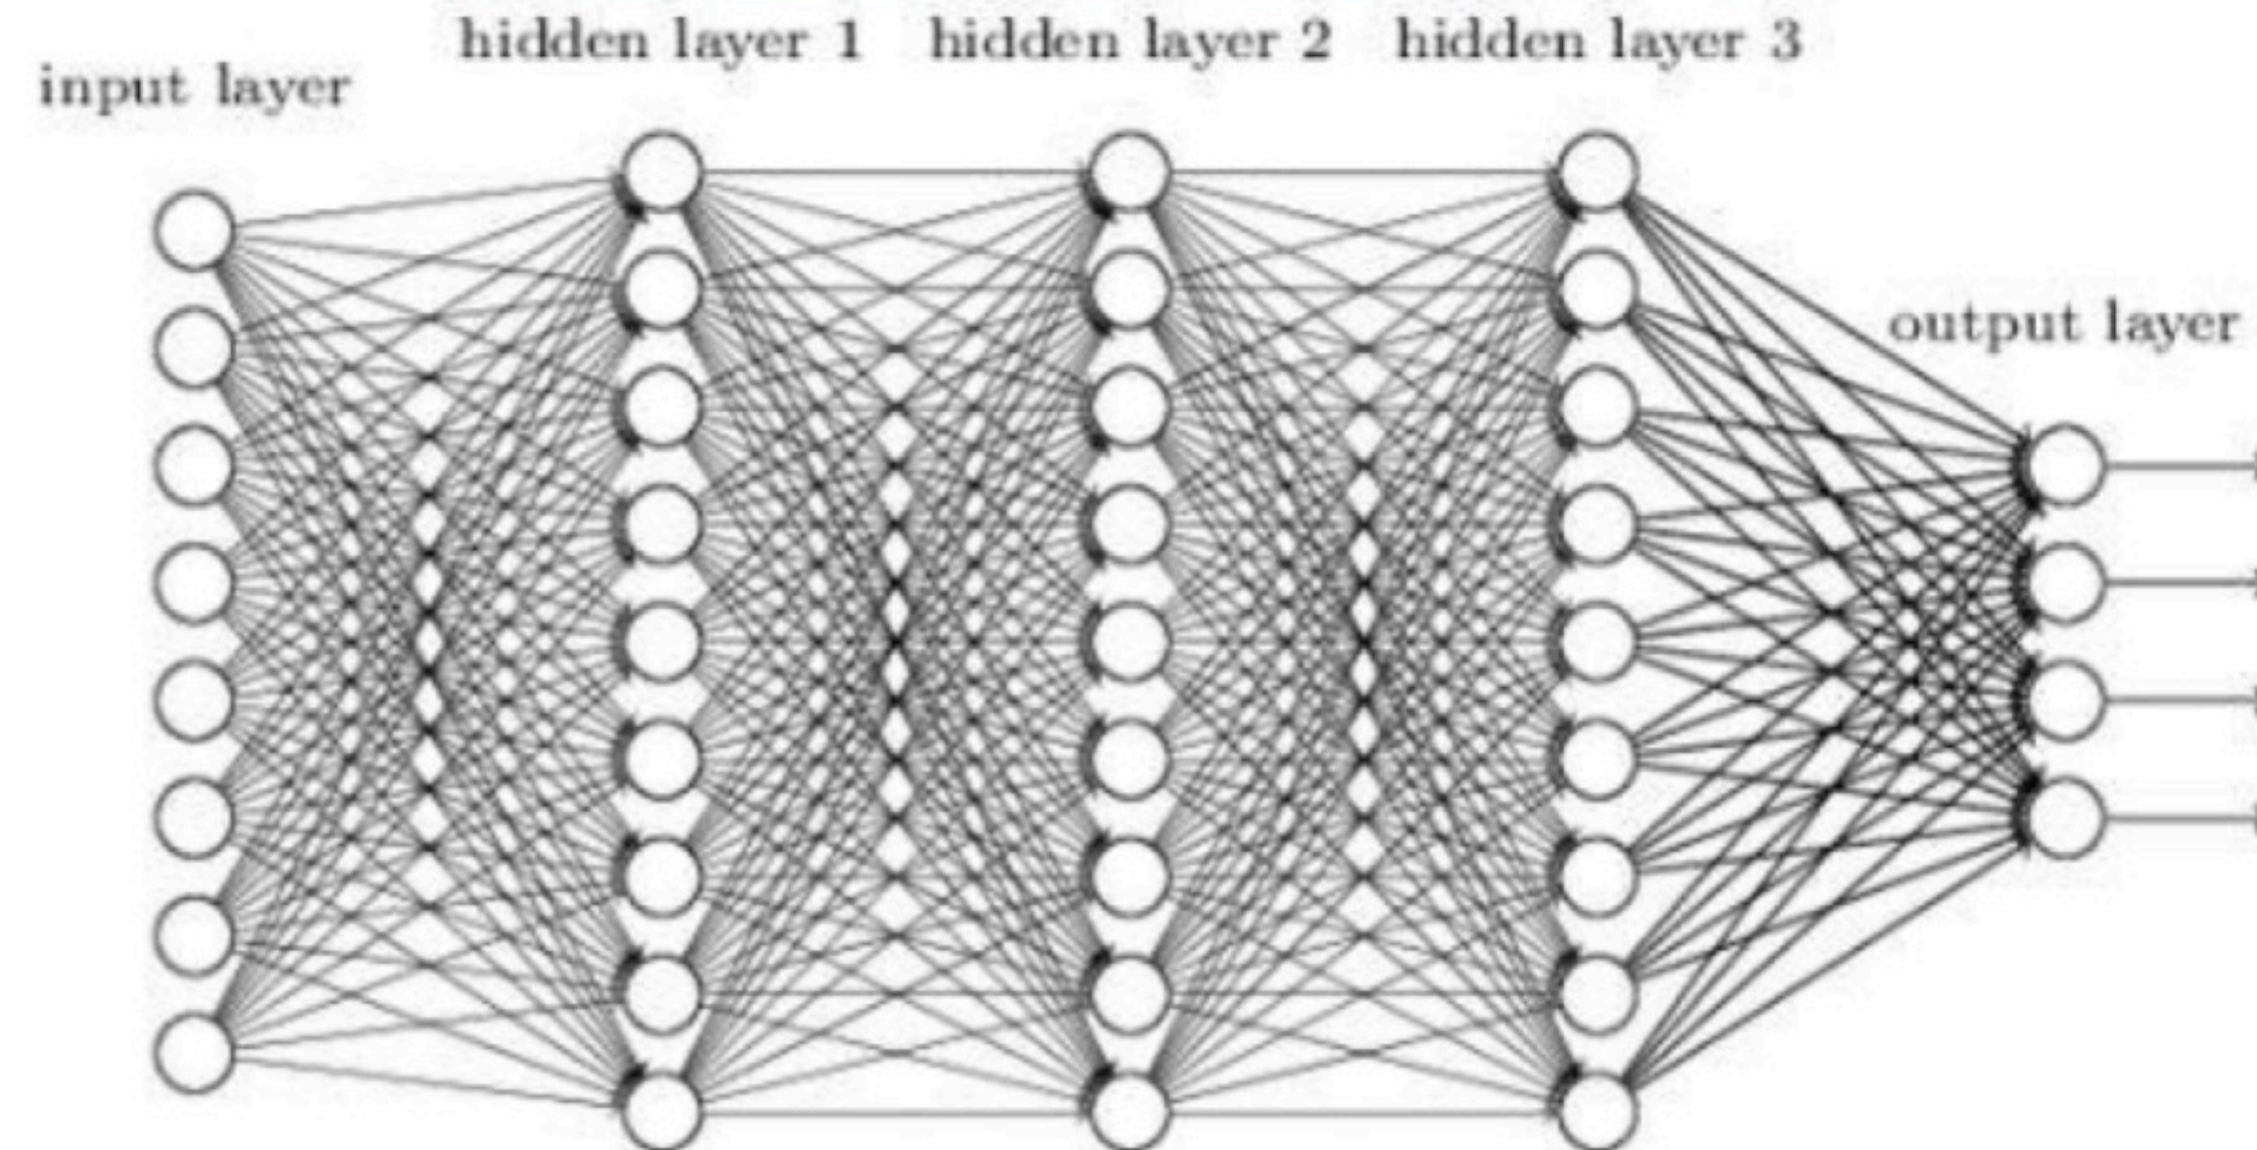
Layers

Fully Connected Layers



Fully Connected Neural Network

deep



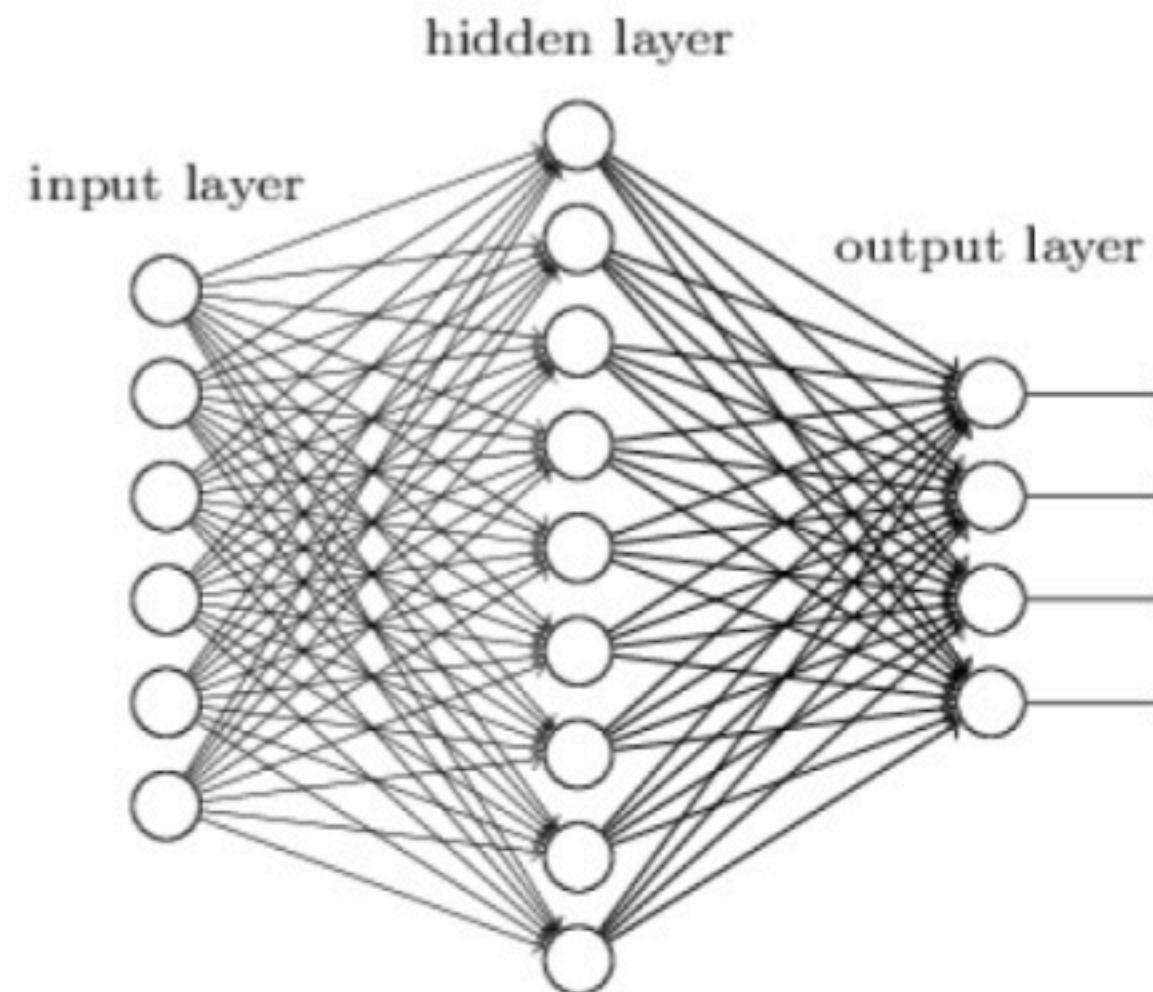
wide



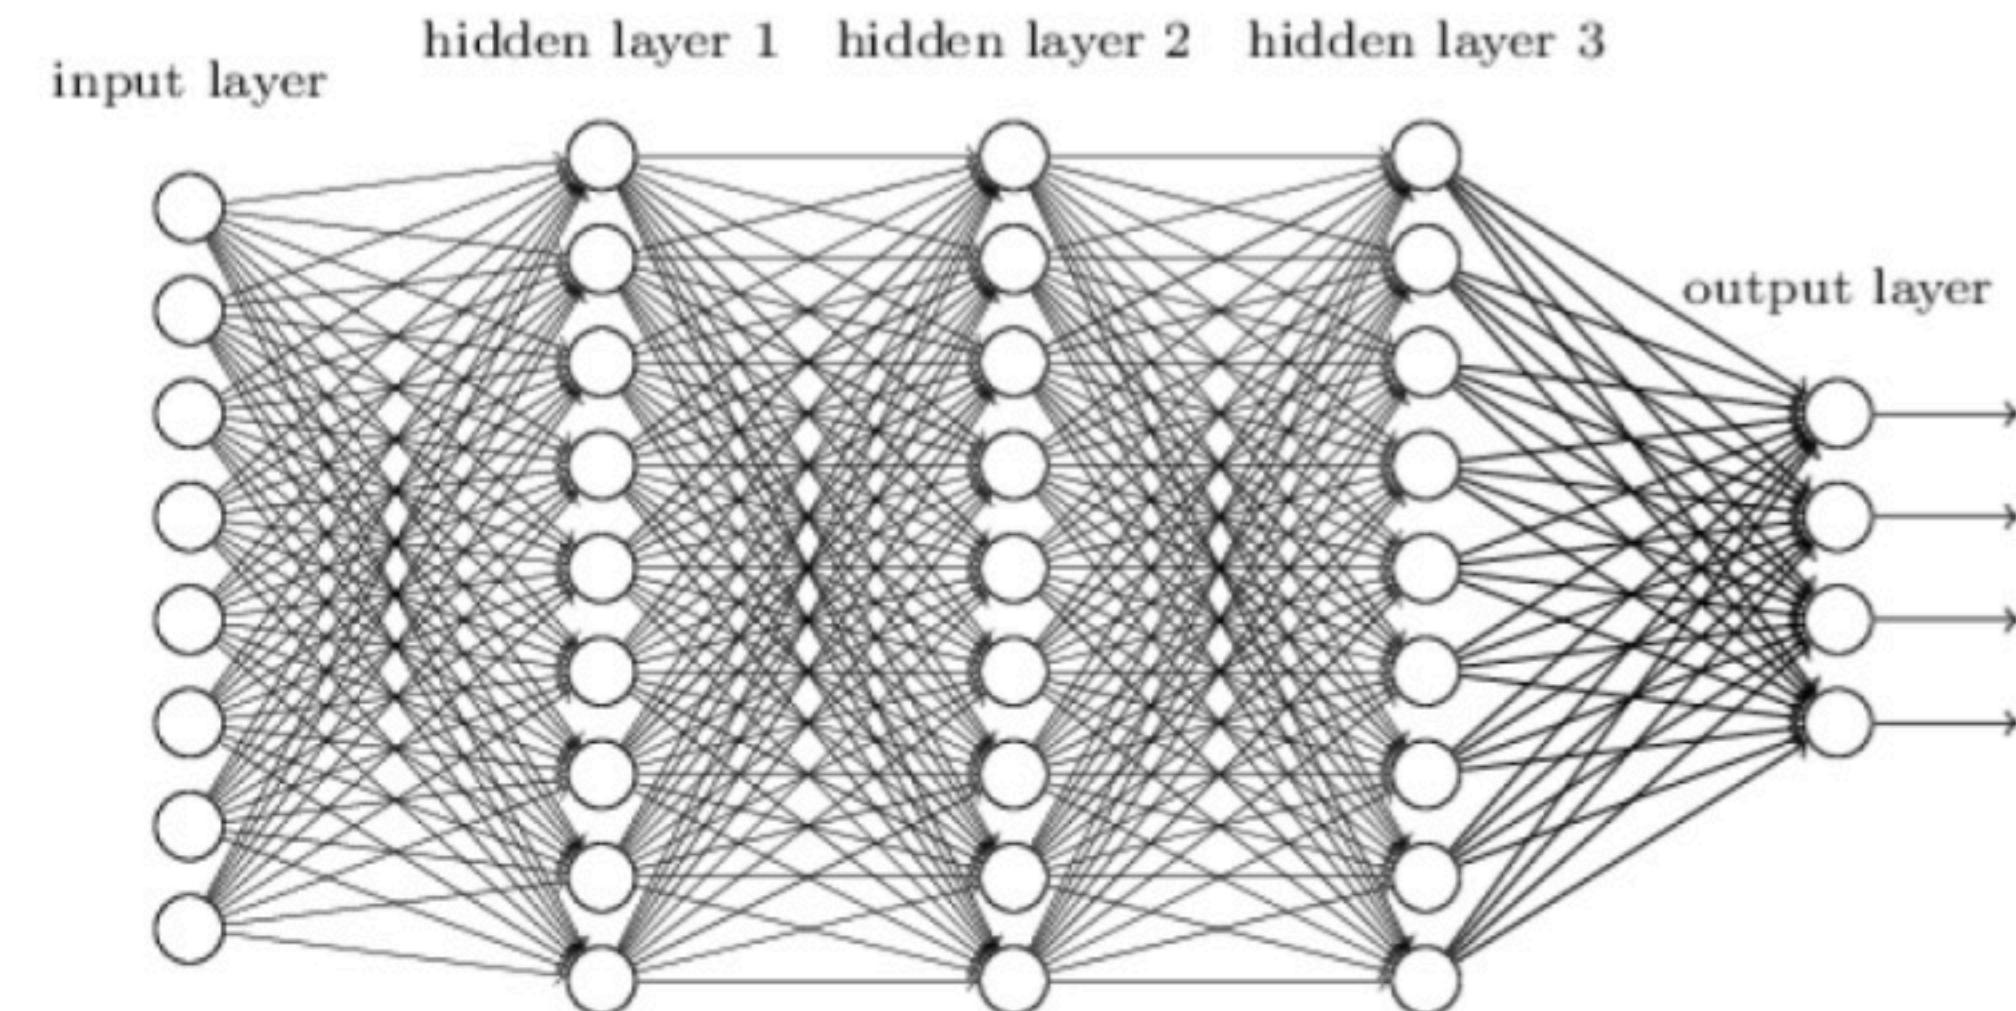
Deep Neural Networks



Neural Network

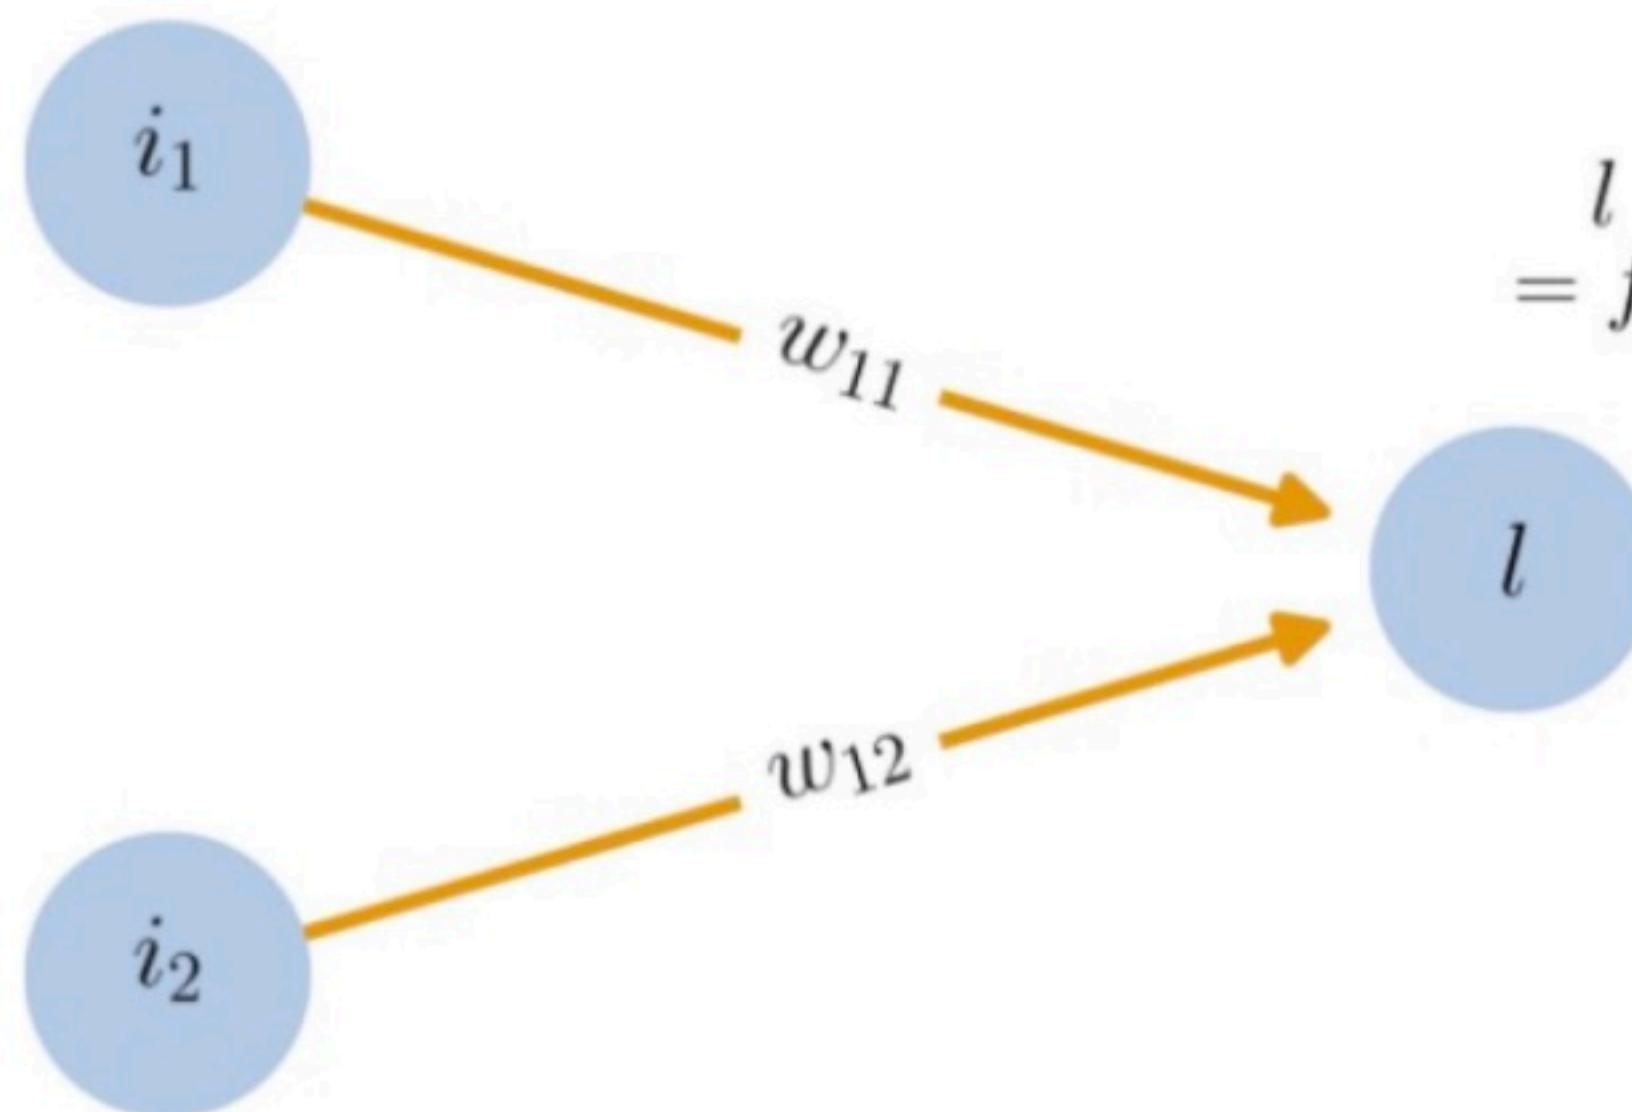


Deep Neural Network



Deep Learning vs Linear Regression

Linear Regression



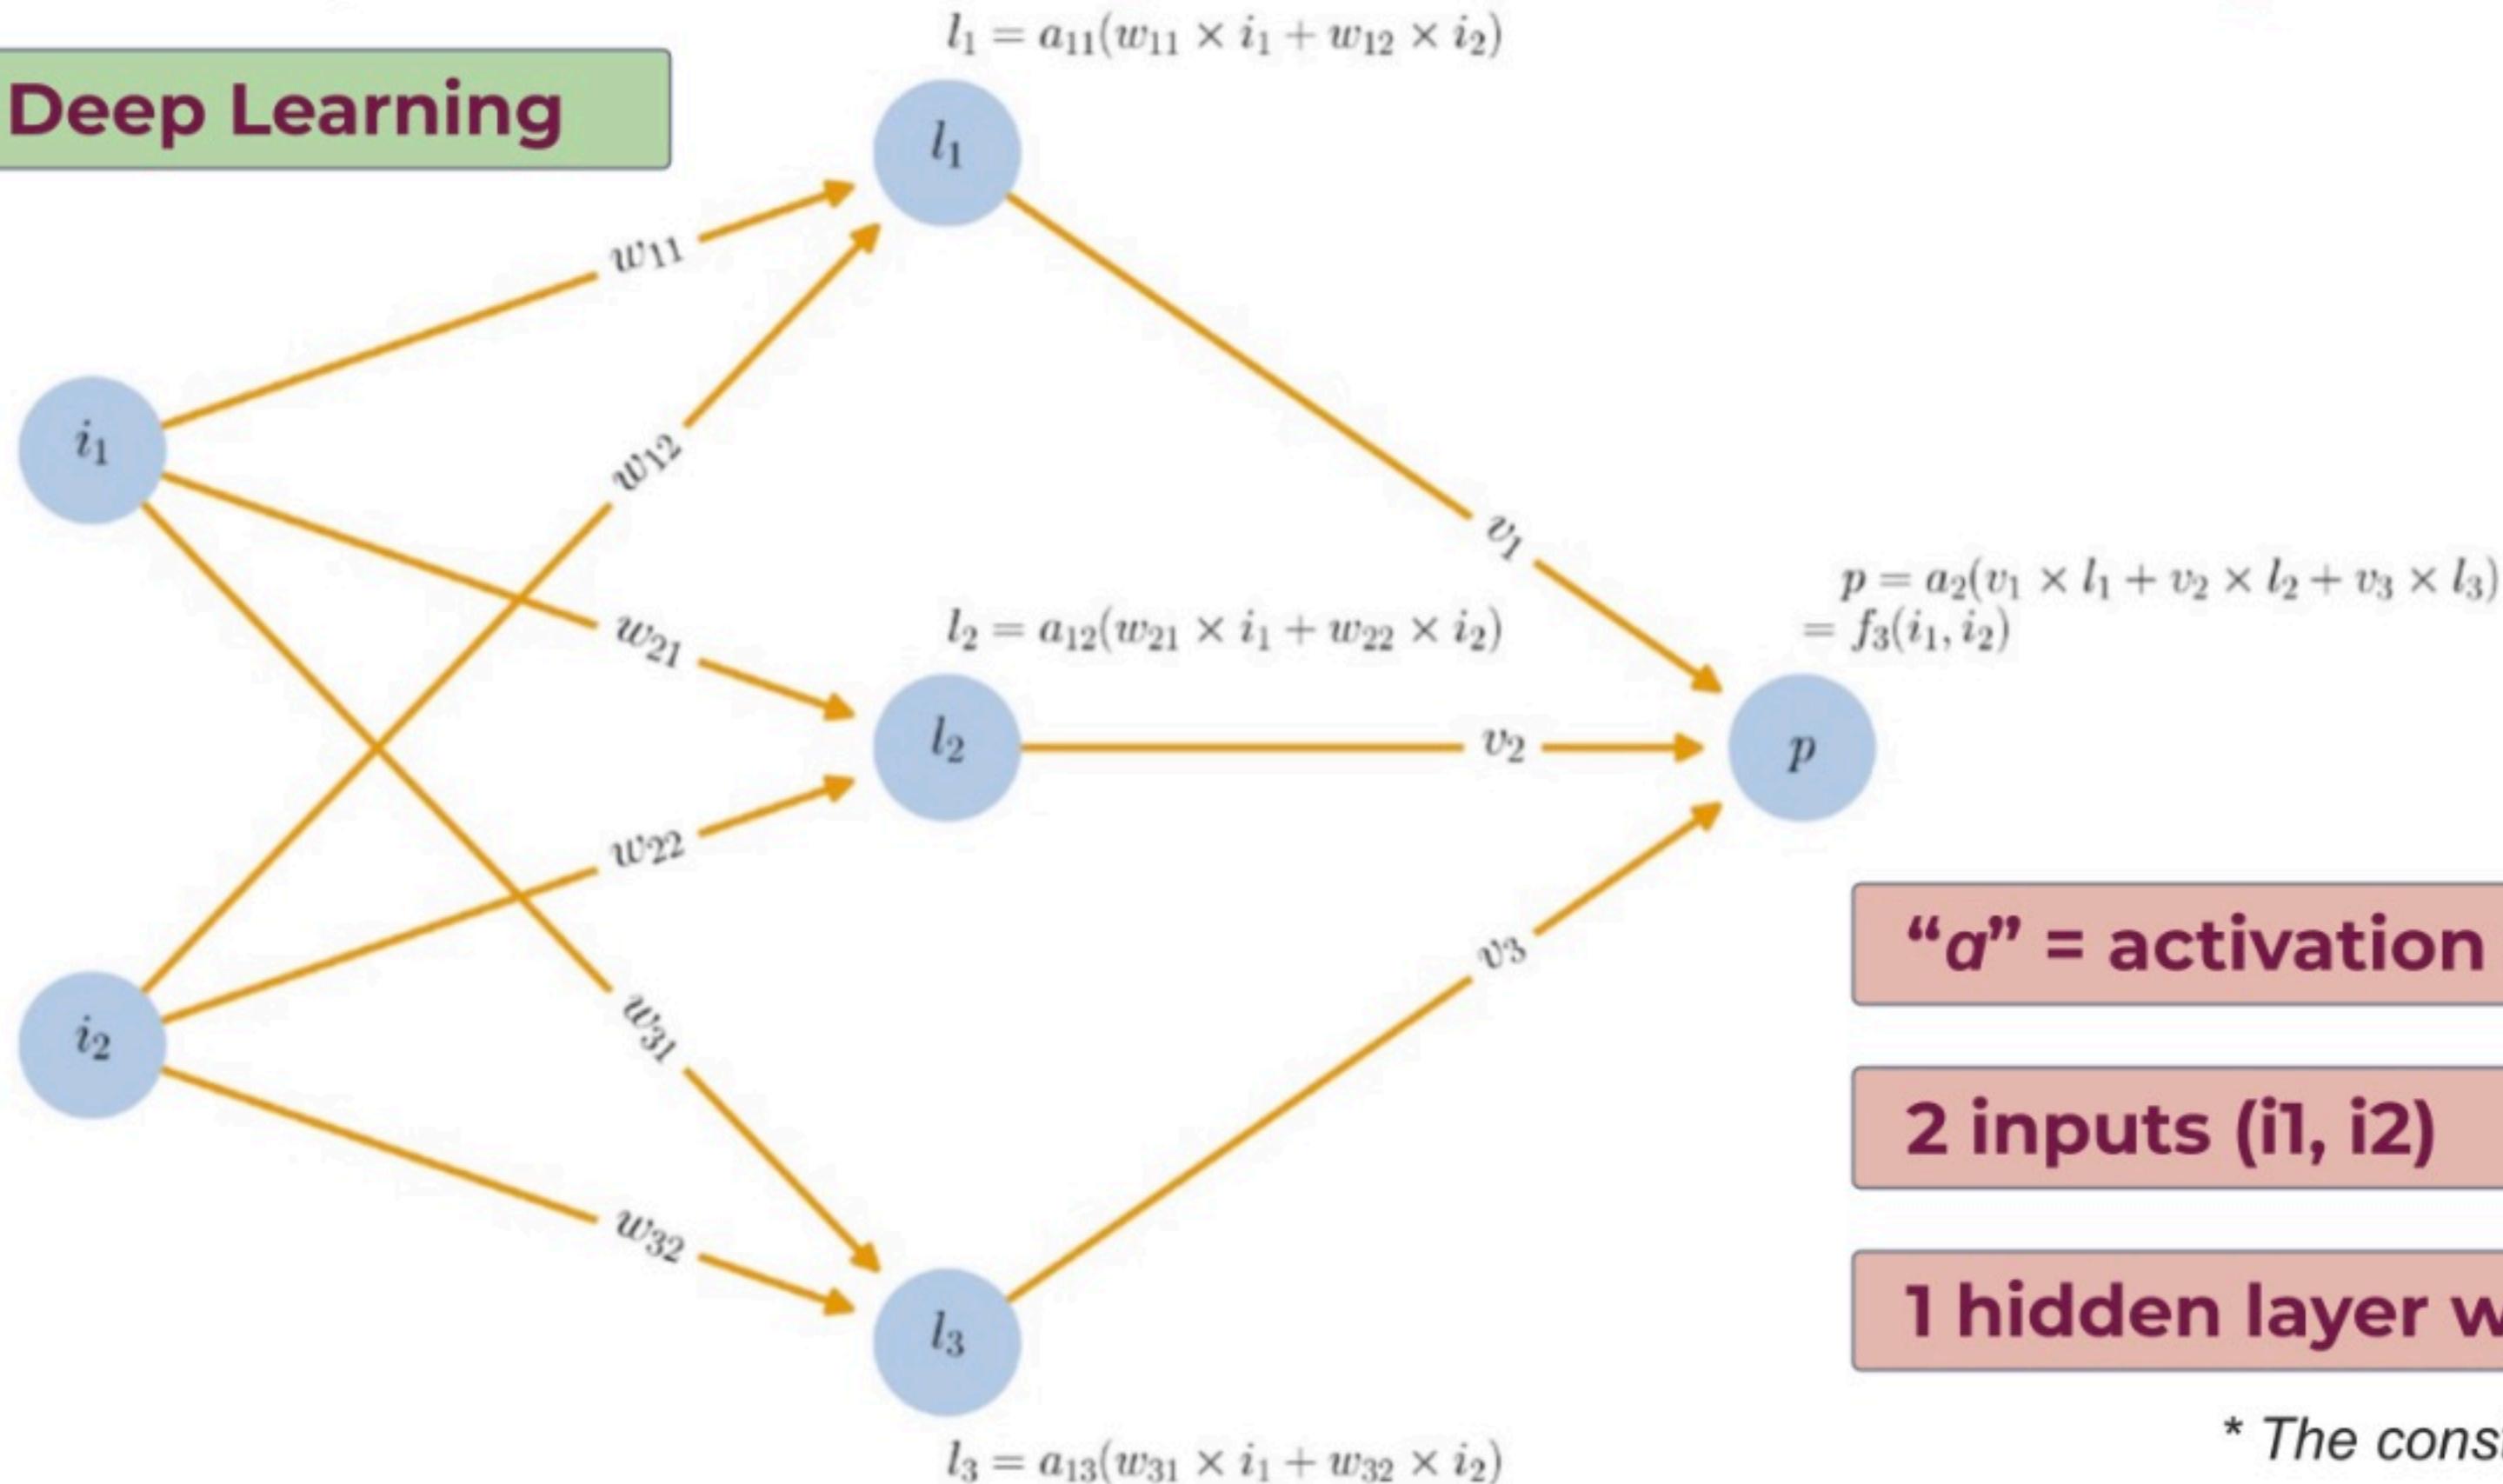
$$l = w_{11} \times i_1 + w_{12} \times i_2 \\ = f_1(i_1, i_2)$$

**prediction of a house price
with two inputs**

*The actual formula is:
 $f_1(i_1, i_2) = w_{11} * i_1 + w_{12} * i_2 + c$.
But the constant (c) is ignored in the picture*

Deep Learning vs Linear Regression

Deep Learning



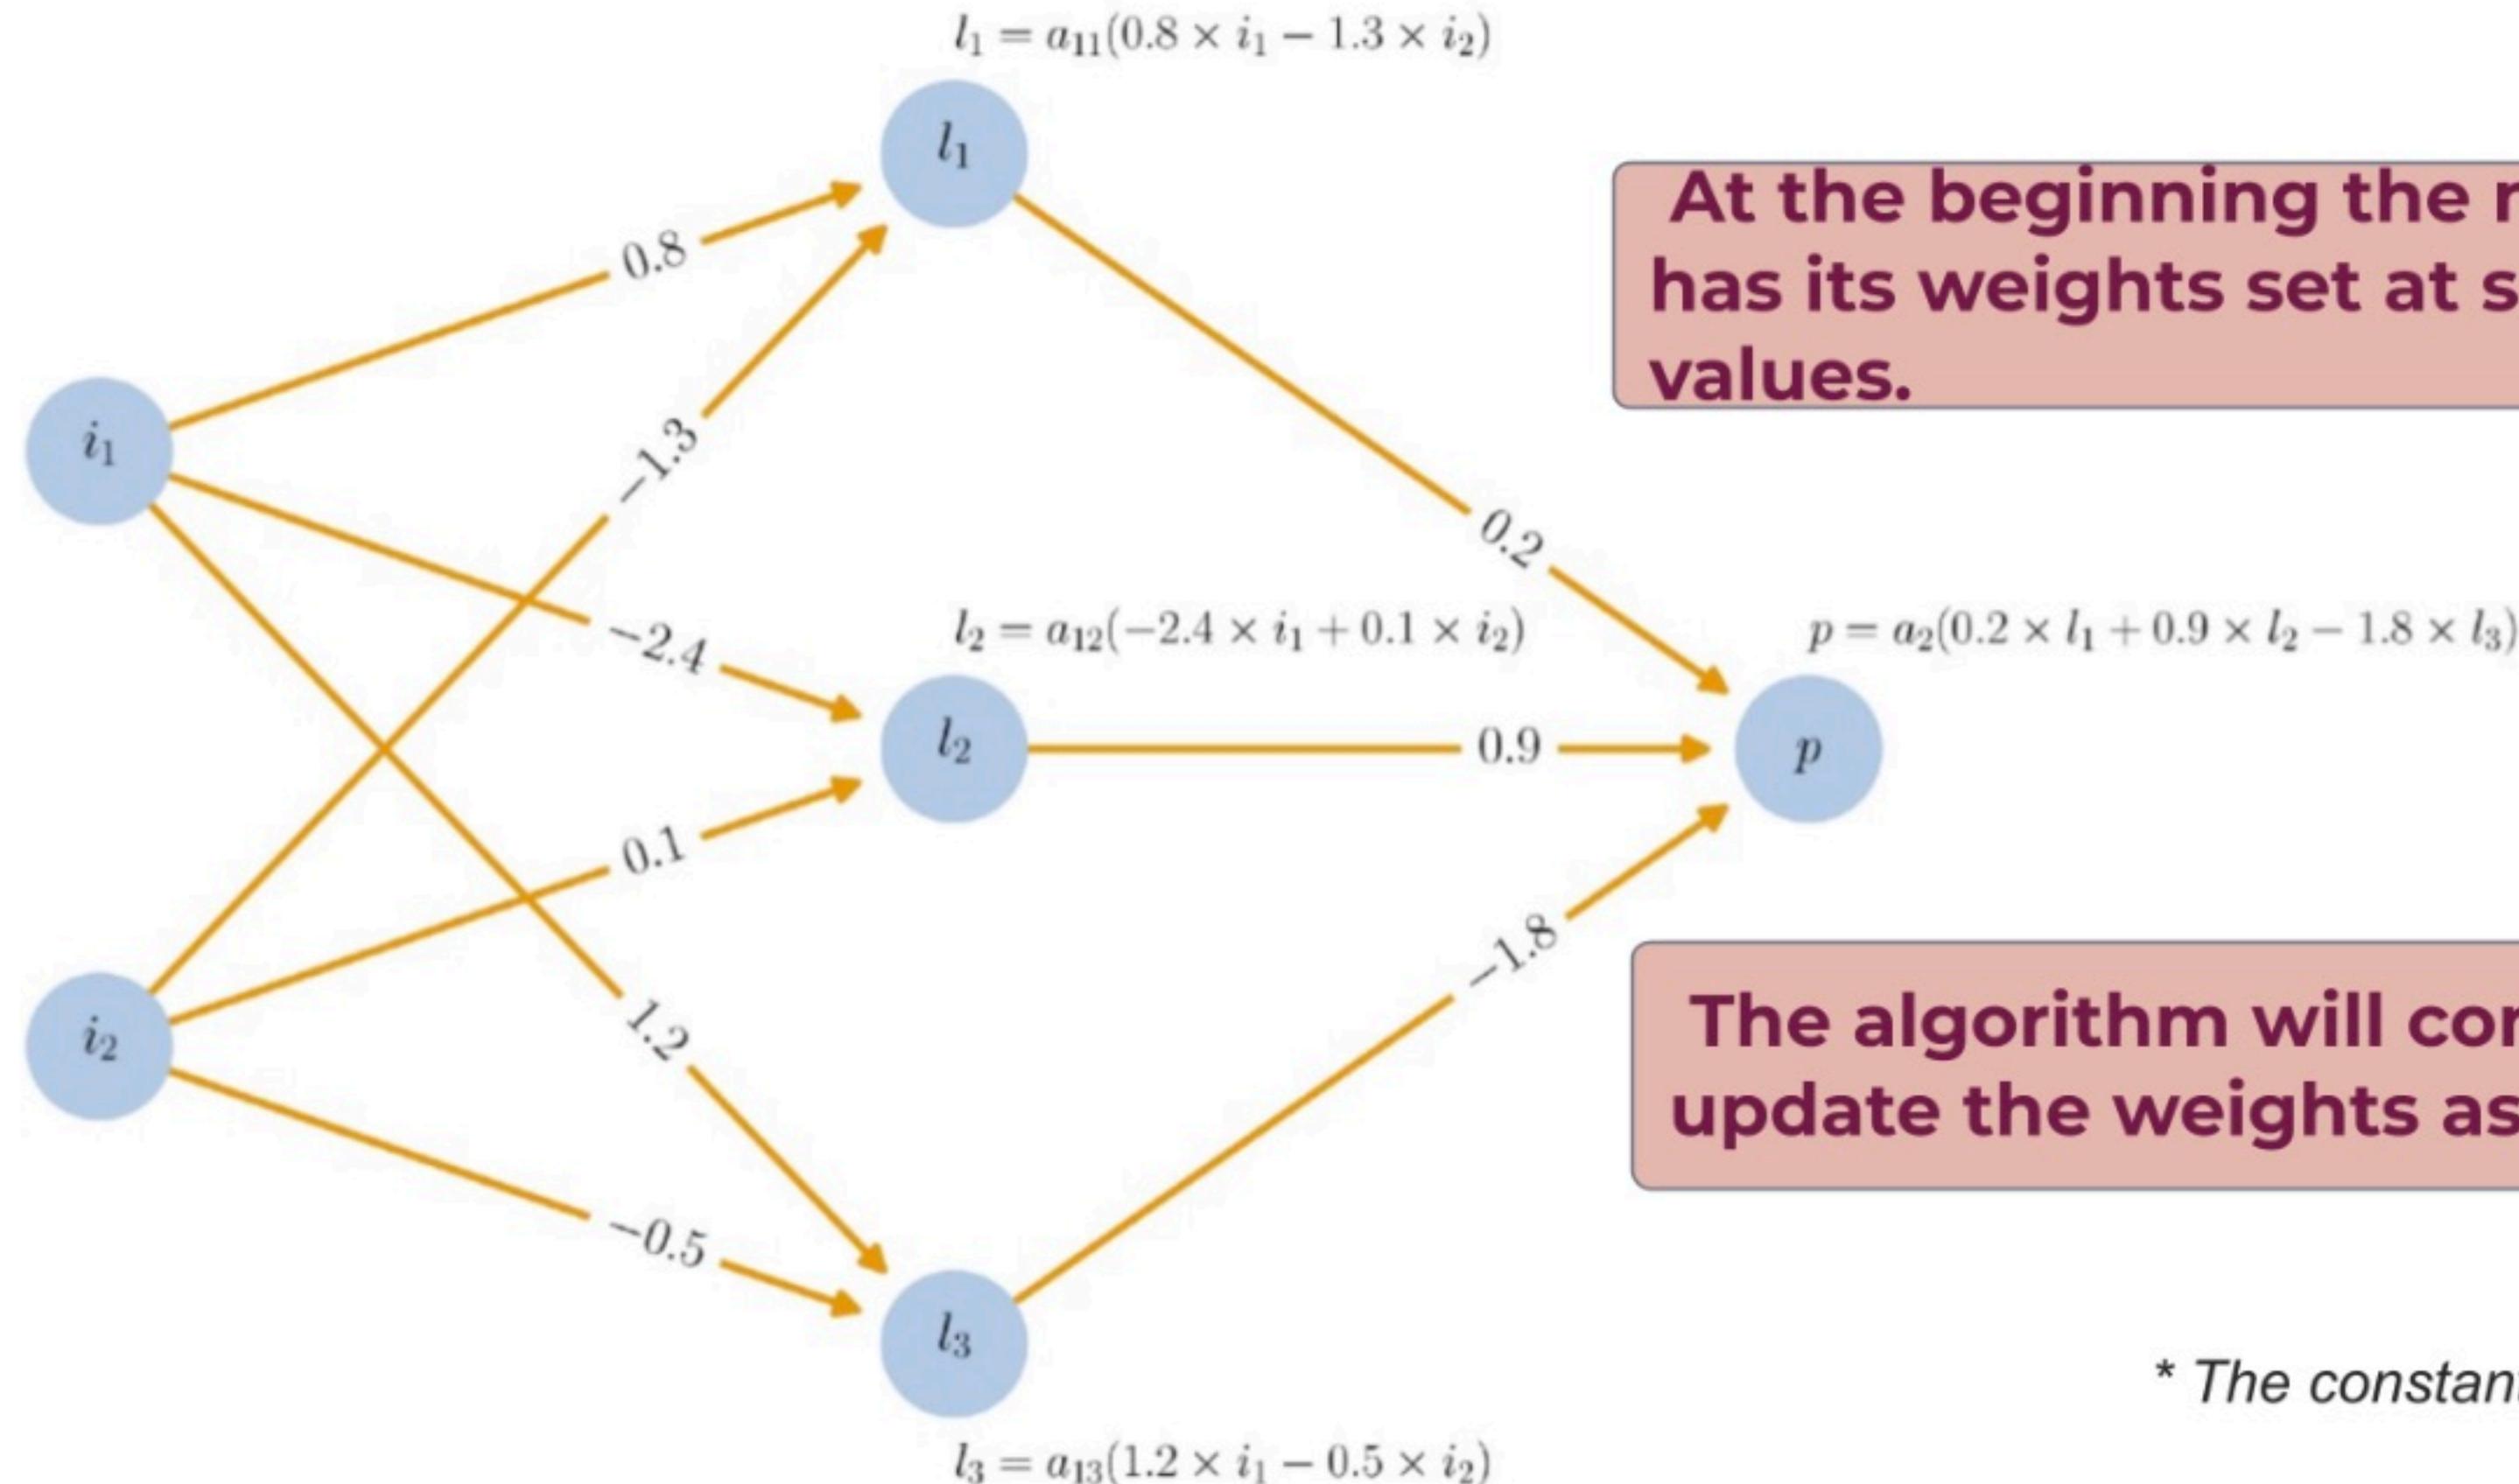
"a" = activation function

2 inputs (i_1, i_2)

1 hidden layer with 3 neurons

* The constants are ignored.

Deep Learning vs Linear Regression

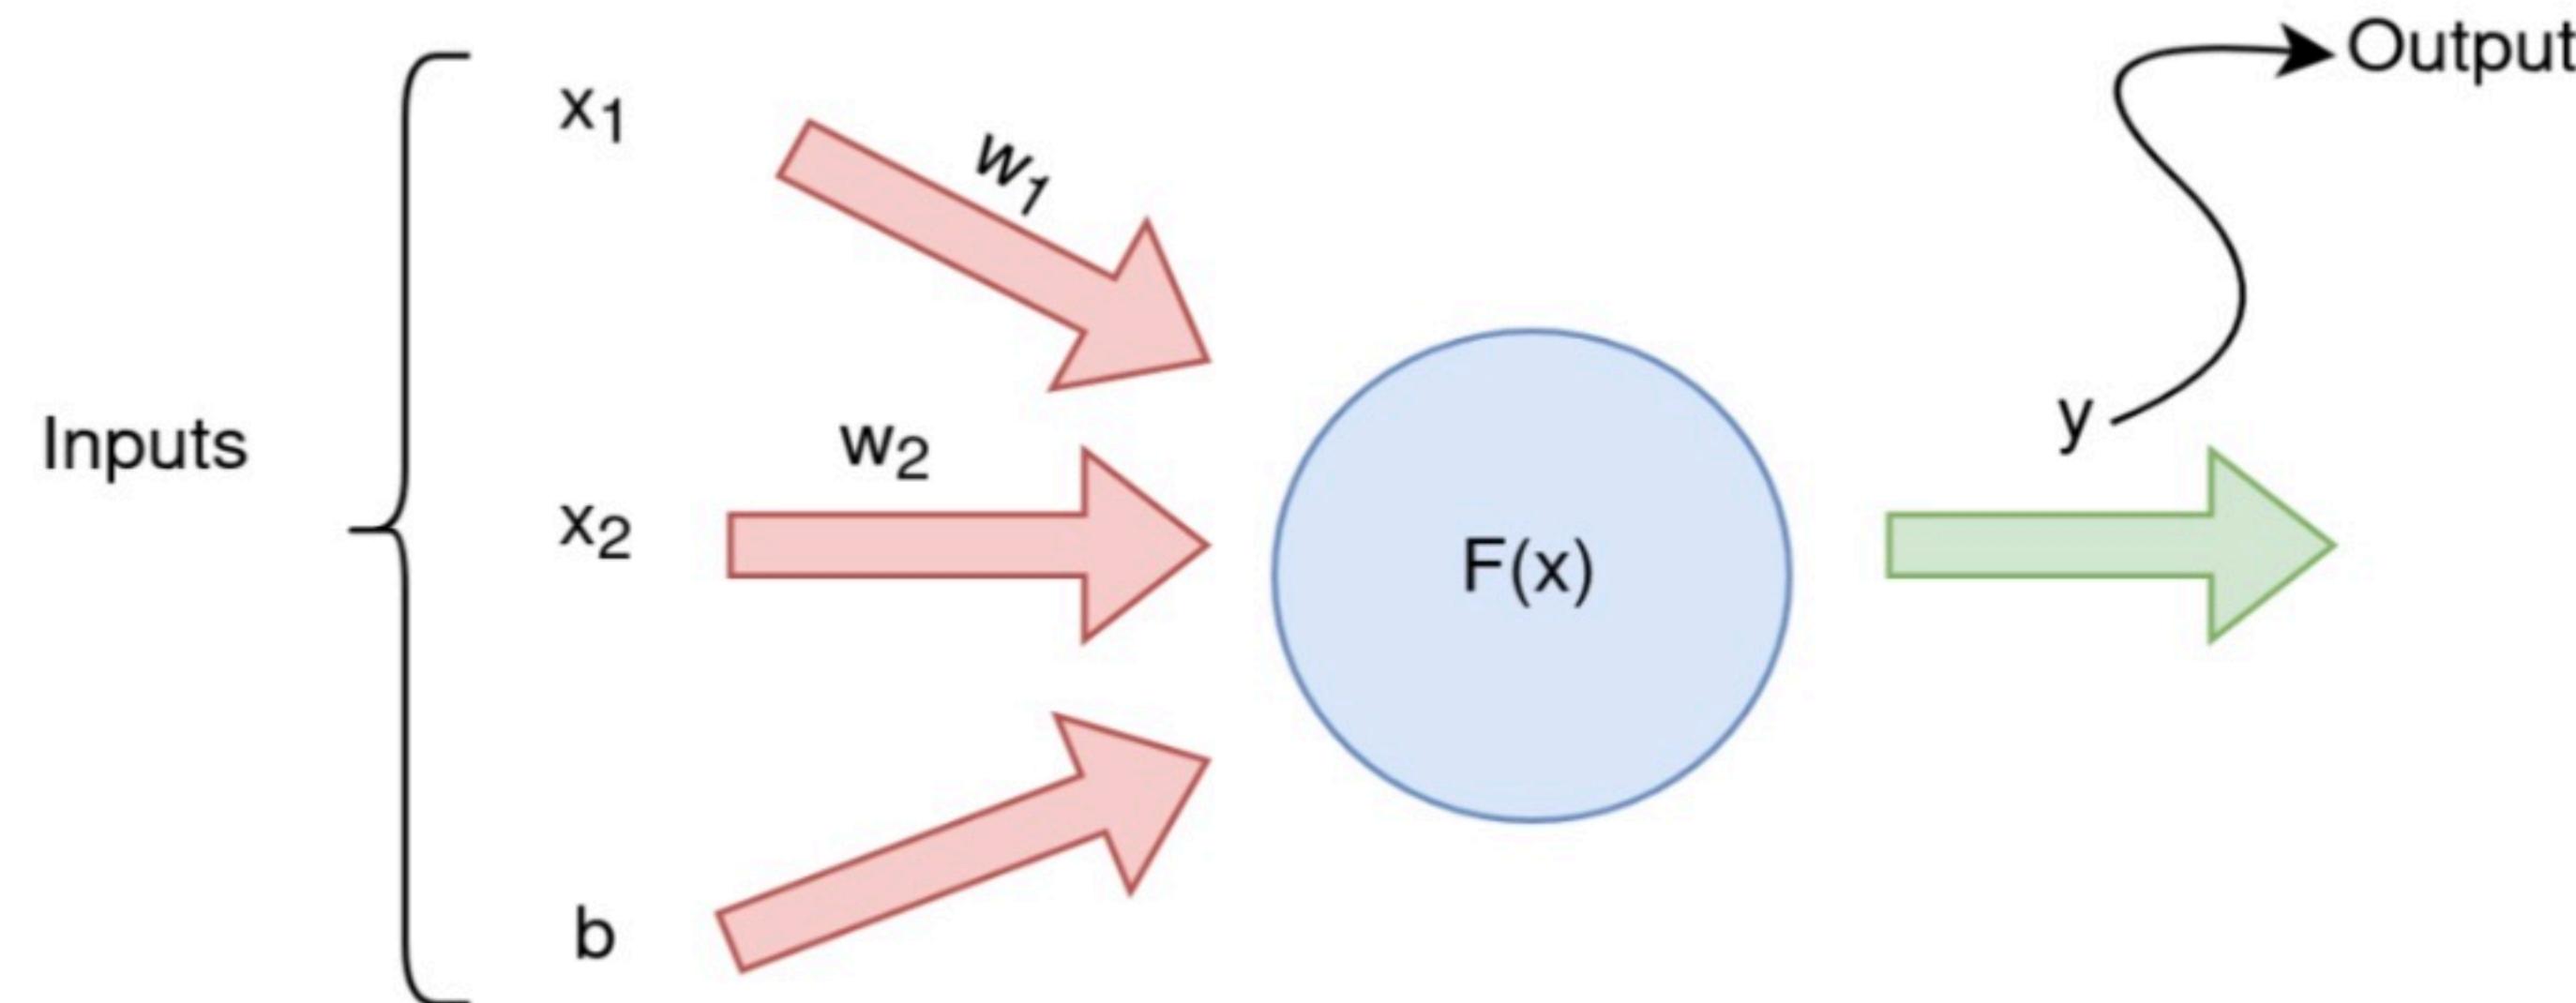




Activation Functions



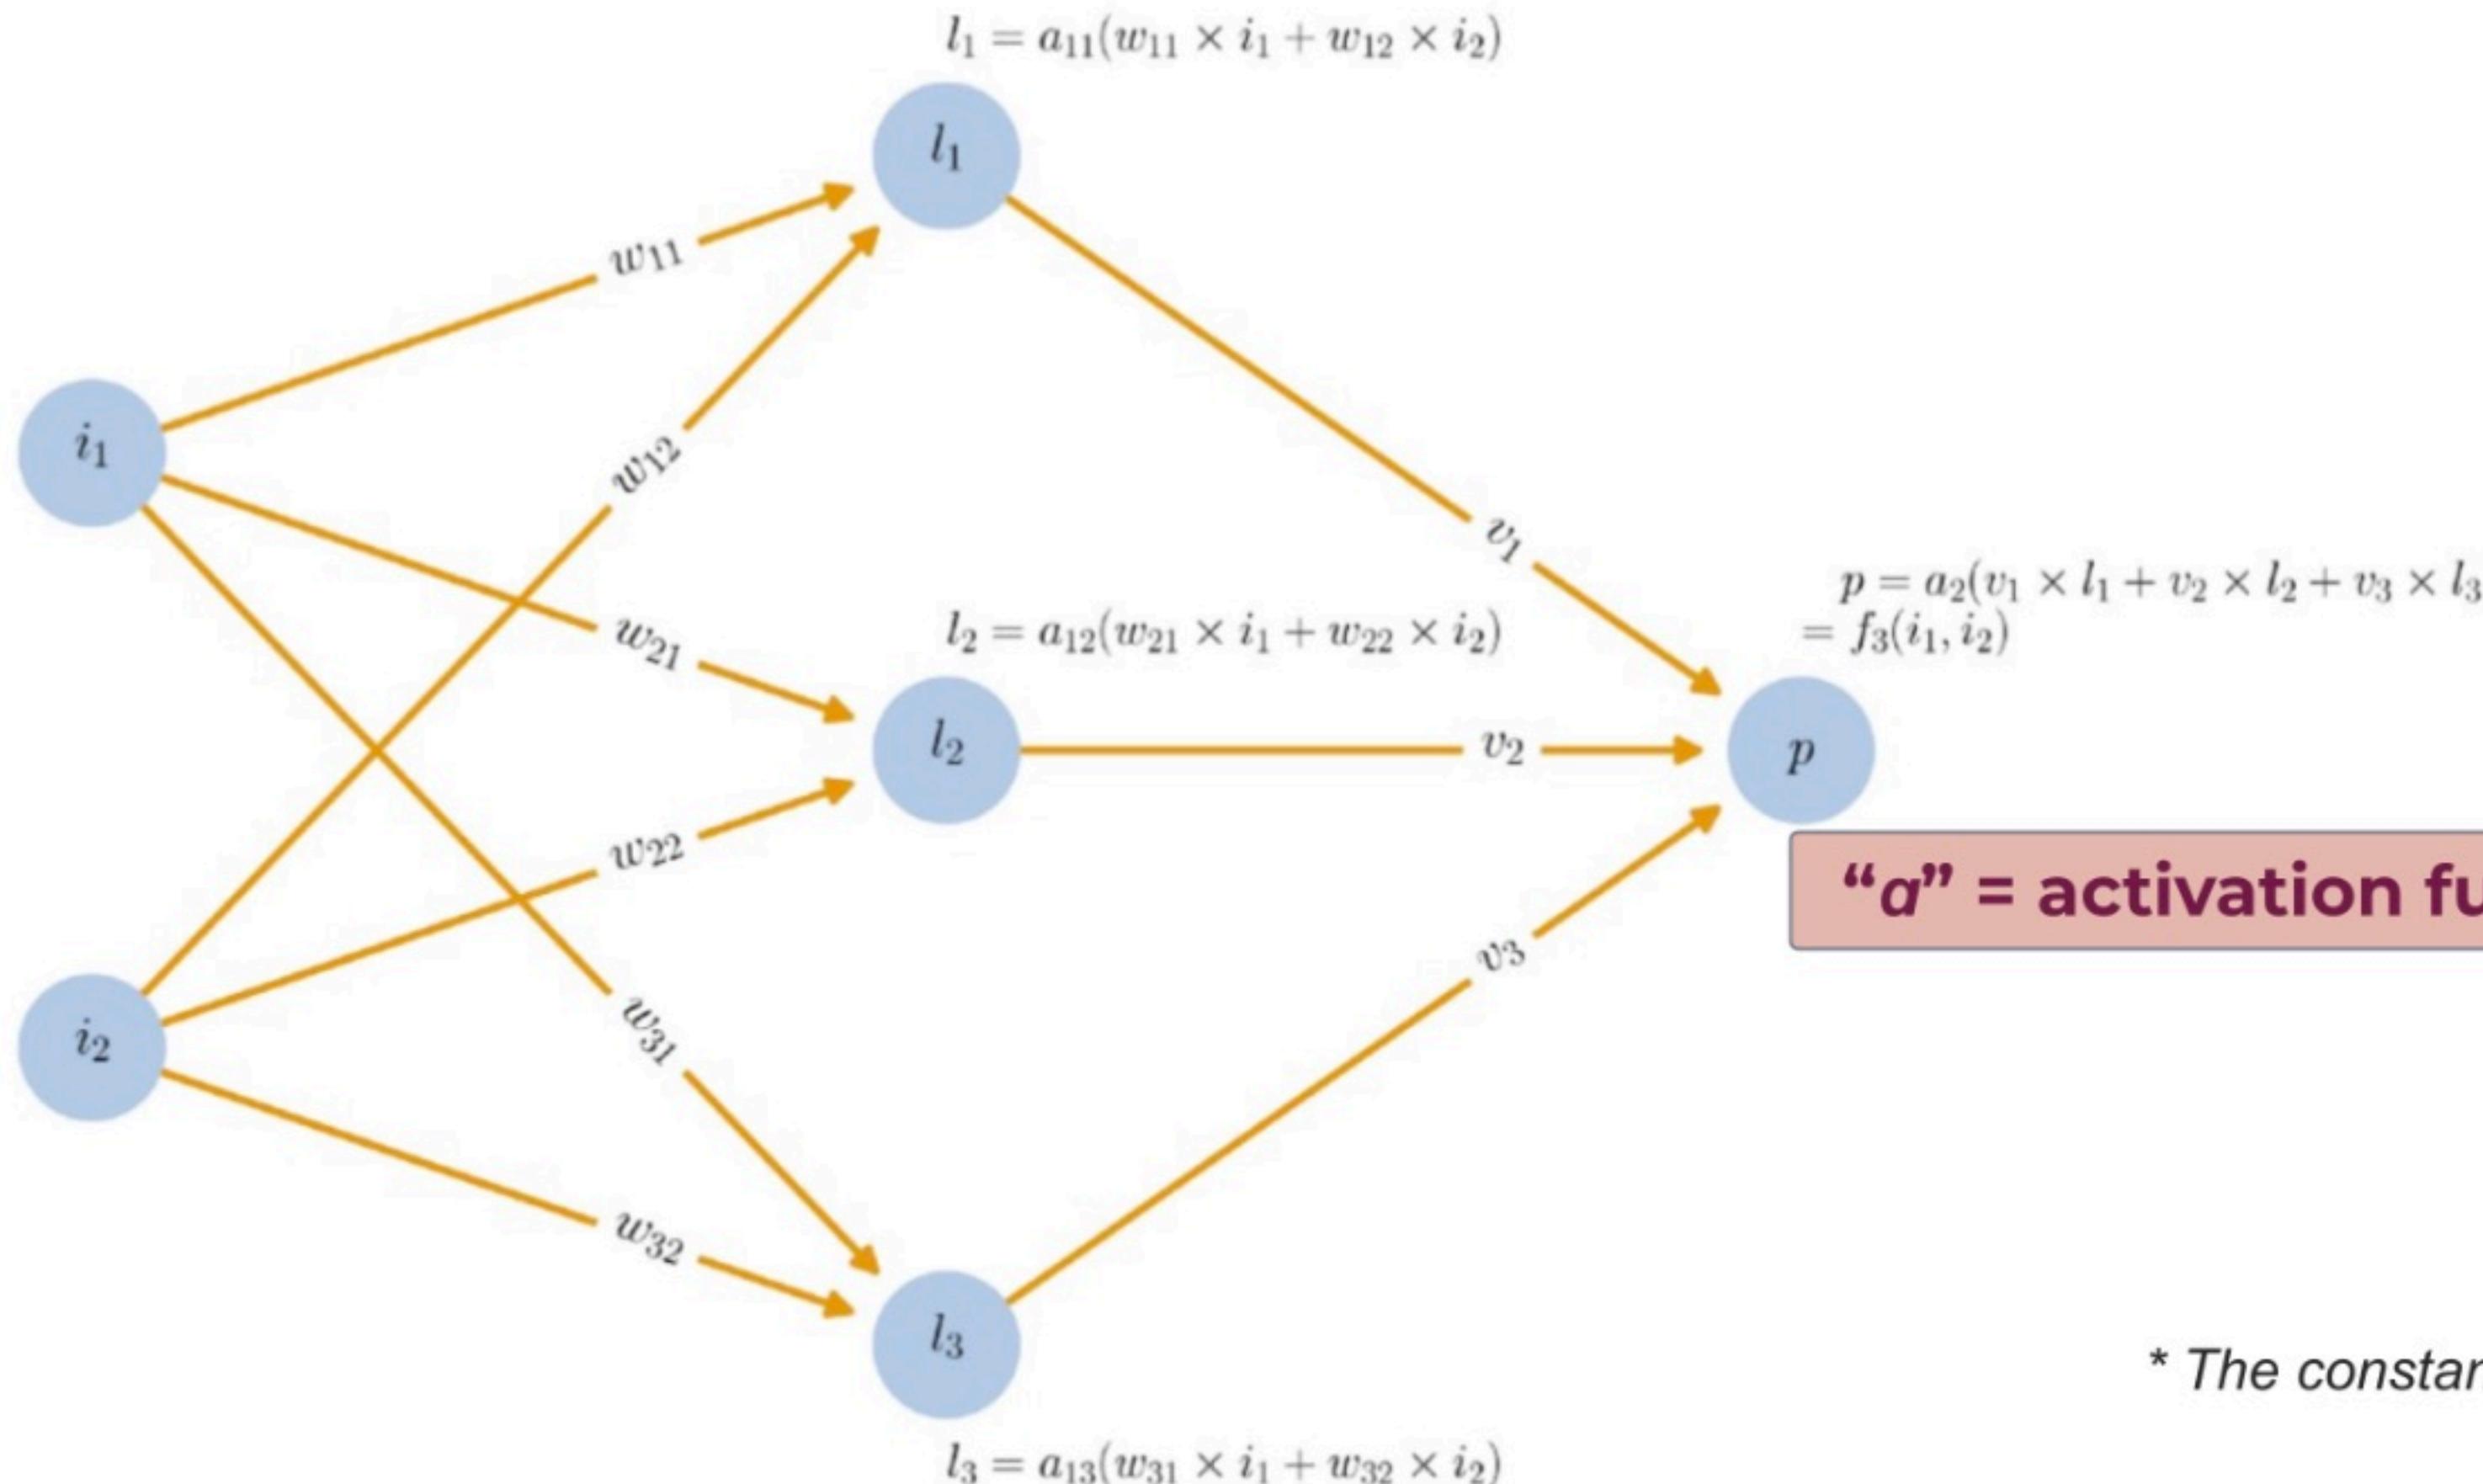
Activation Functions



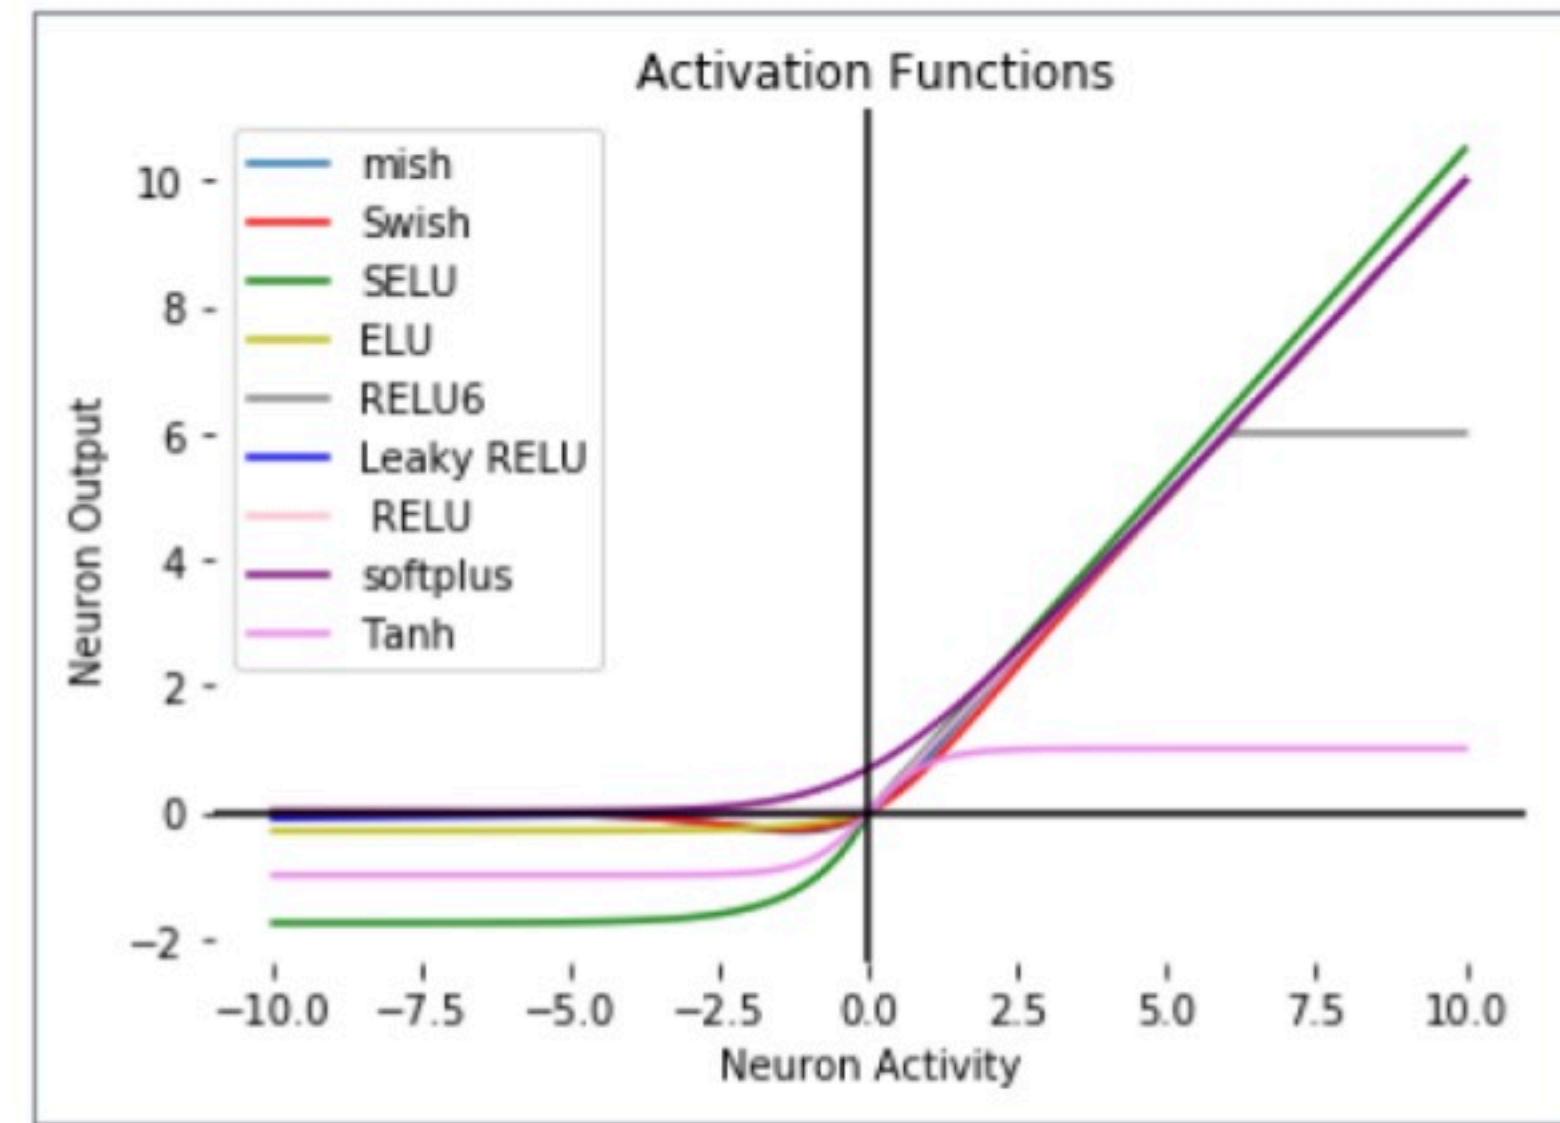
$$\hat{y} = \sum_{i=1}^n x_i w_i + b$$

The value of Y can be anything ranging from -inf to +inf. The neuron really doesn't know the bounds of the value.

Activation Functions



Activation Functions



- ▶ The activation function of a node **defines the output** of that node given an input or set of inputs.
- ▶ Activation functions are **decision making units** of neural networks.

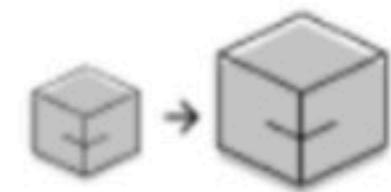
Activation Functions

Role of the Activation function

- ▶ Activation functions decide if a neuron will be **fired or not** and **strength of the signal** outputted by the neuron.
- ▶ Activation functions determine the deep learning model's **accuracy and the computational efficiency** for training a model.
- ▶ Activation functions have to be computationally very efficient as it decides on the network's **convergence speed** or if the network will **ever converge**.

Activation Functions

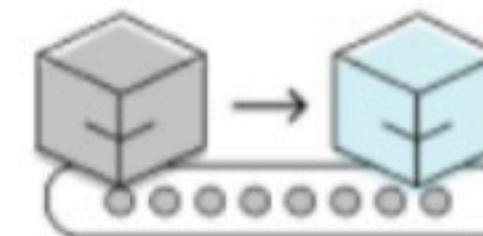
How it works?



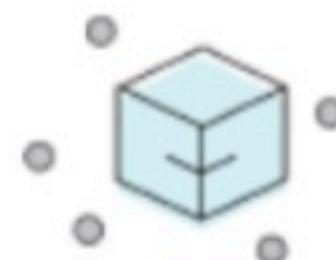
Take input and multiply by the neuron's weight.



Add bias*



Feed the result, x , to the activation function: $f(x)$



Take the output and transmit to the next layer of neurons.

Activation Functions



Type of Activation Functions

Binary Step Function

1

Linear Activation Function

2

Non-Linear Activation Function

3

Sigmoid

Softmax

ReLU

TanH

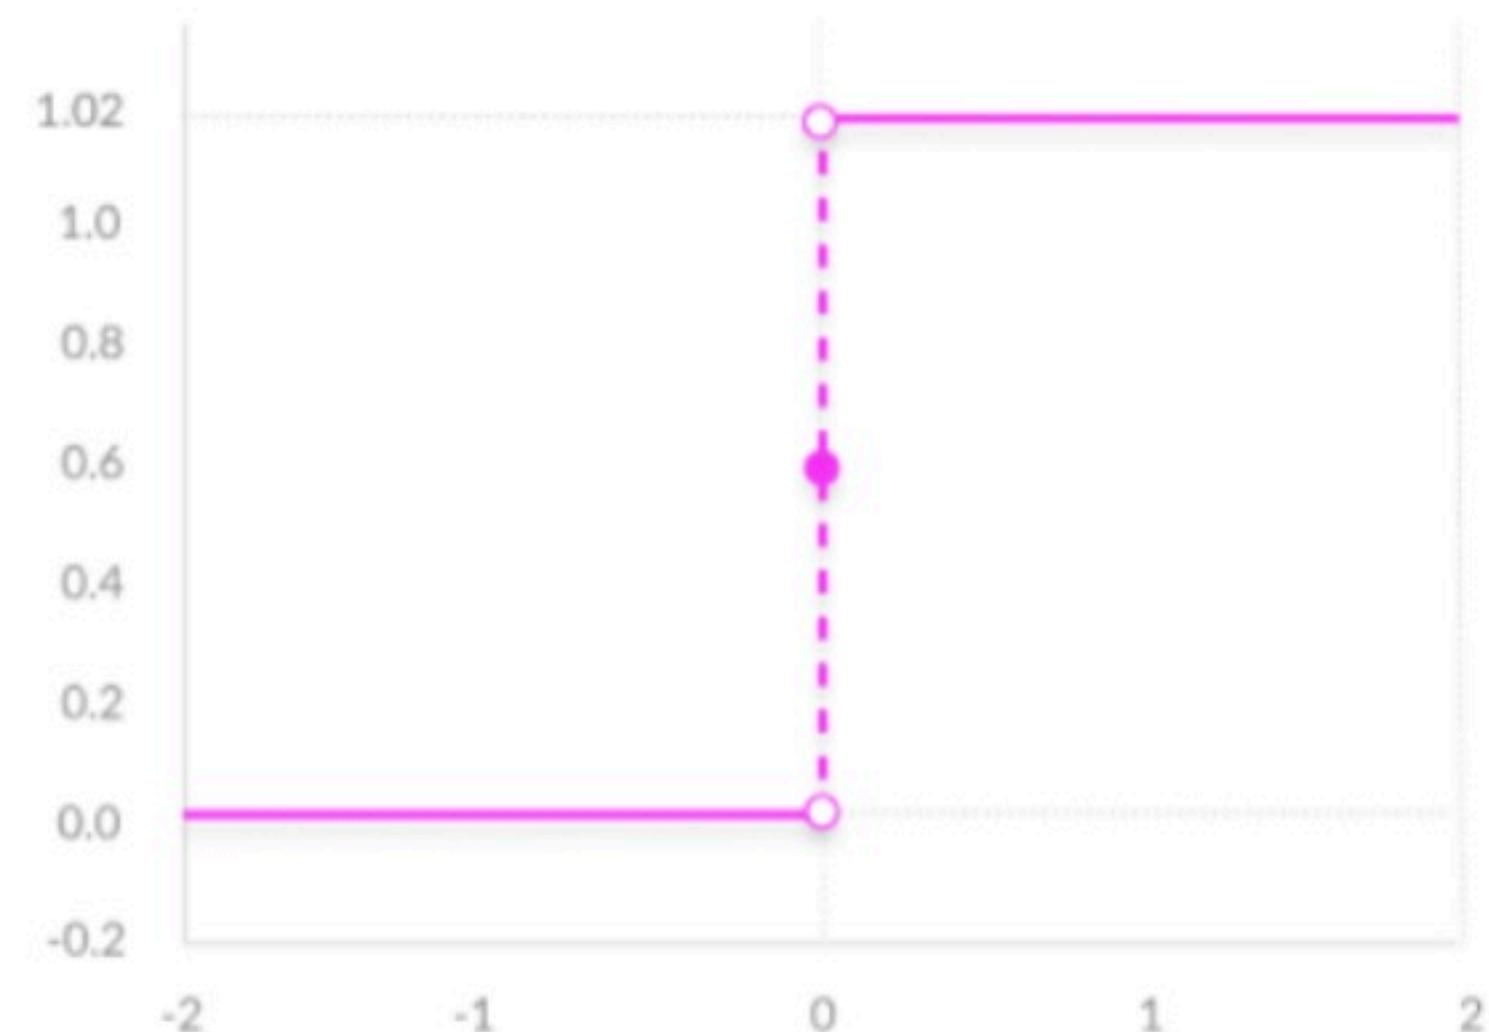
**Leaky
ReLU**

* Modern neural network models
use non-linear activation functions

Type of Activation Functions

Binary Step Function

No multi classification



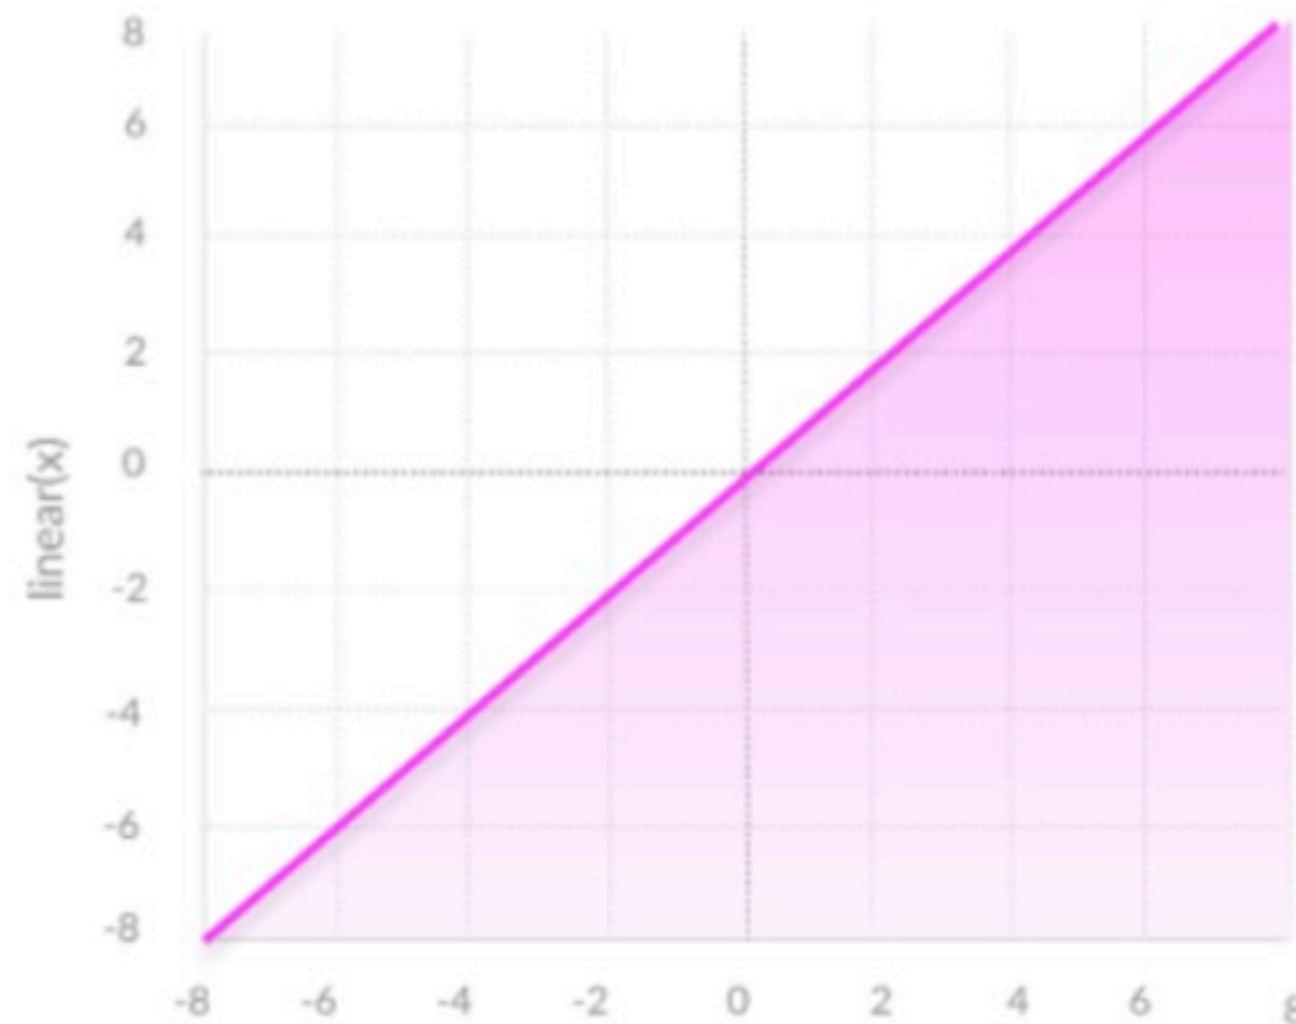
The function produces **binary output**. The function produces 1 (or true) when input **passes threshold limit** whereas it produces 0 (or false) when input does **not pass threshold**.

Type of Activation Functions

Linear Activation Function

No backpropagation

One hidden layer



$$f(x) = x$$

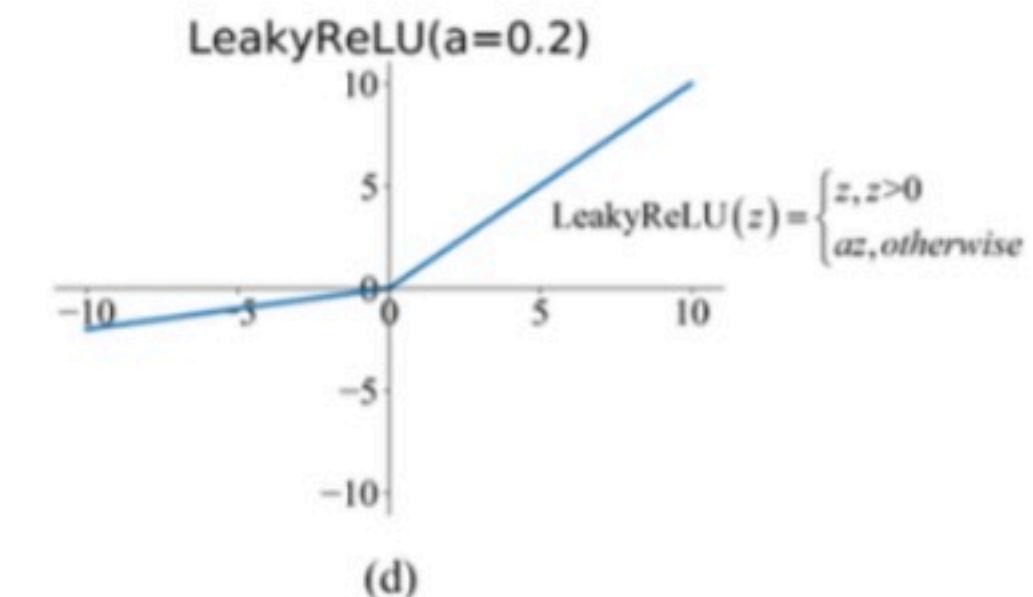
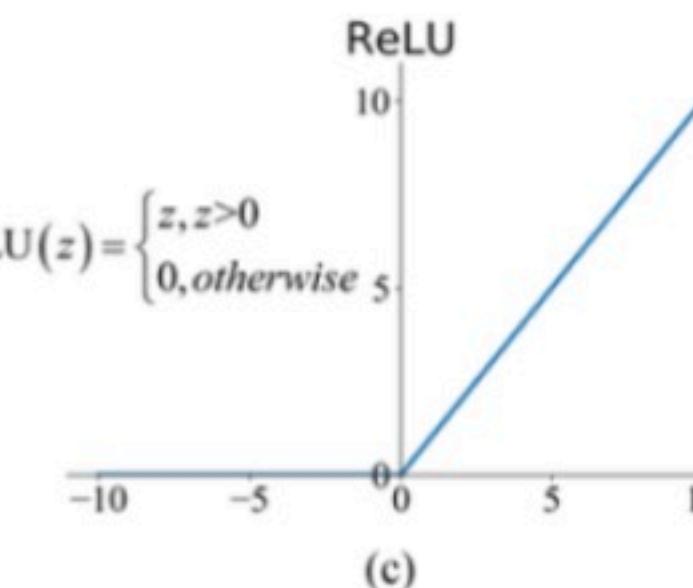
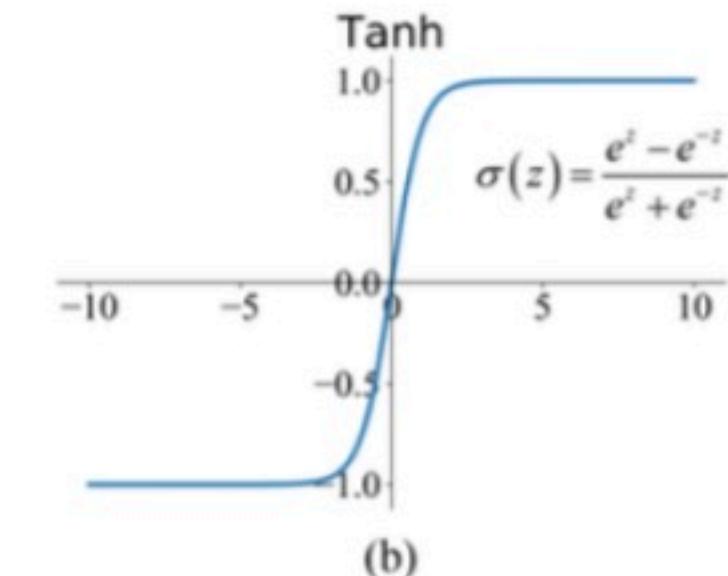
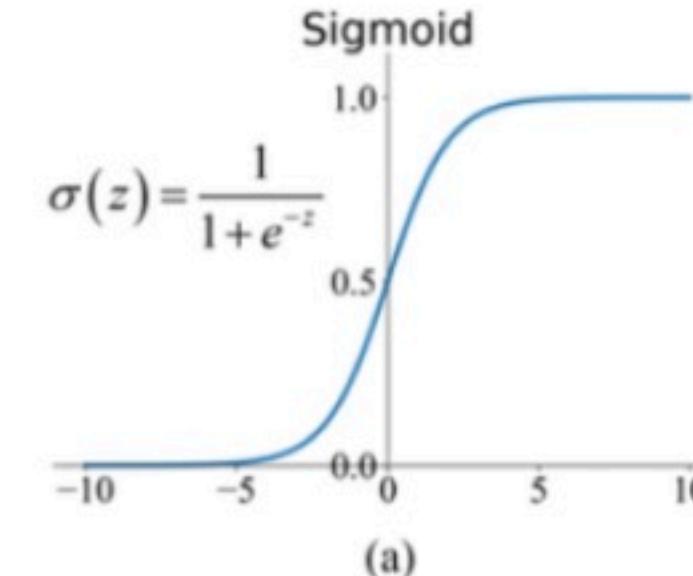
It takes the inputs, multiplied by the weights for each neuron, and creates an output signal proportional to the input.

Type of Activation Functions

Non-Linear Activation Functions

Allow backpropagation

Allow multiple hidden layers



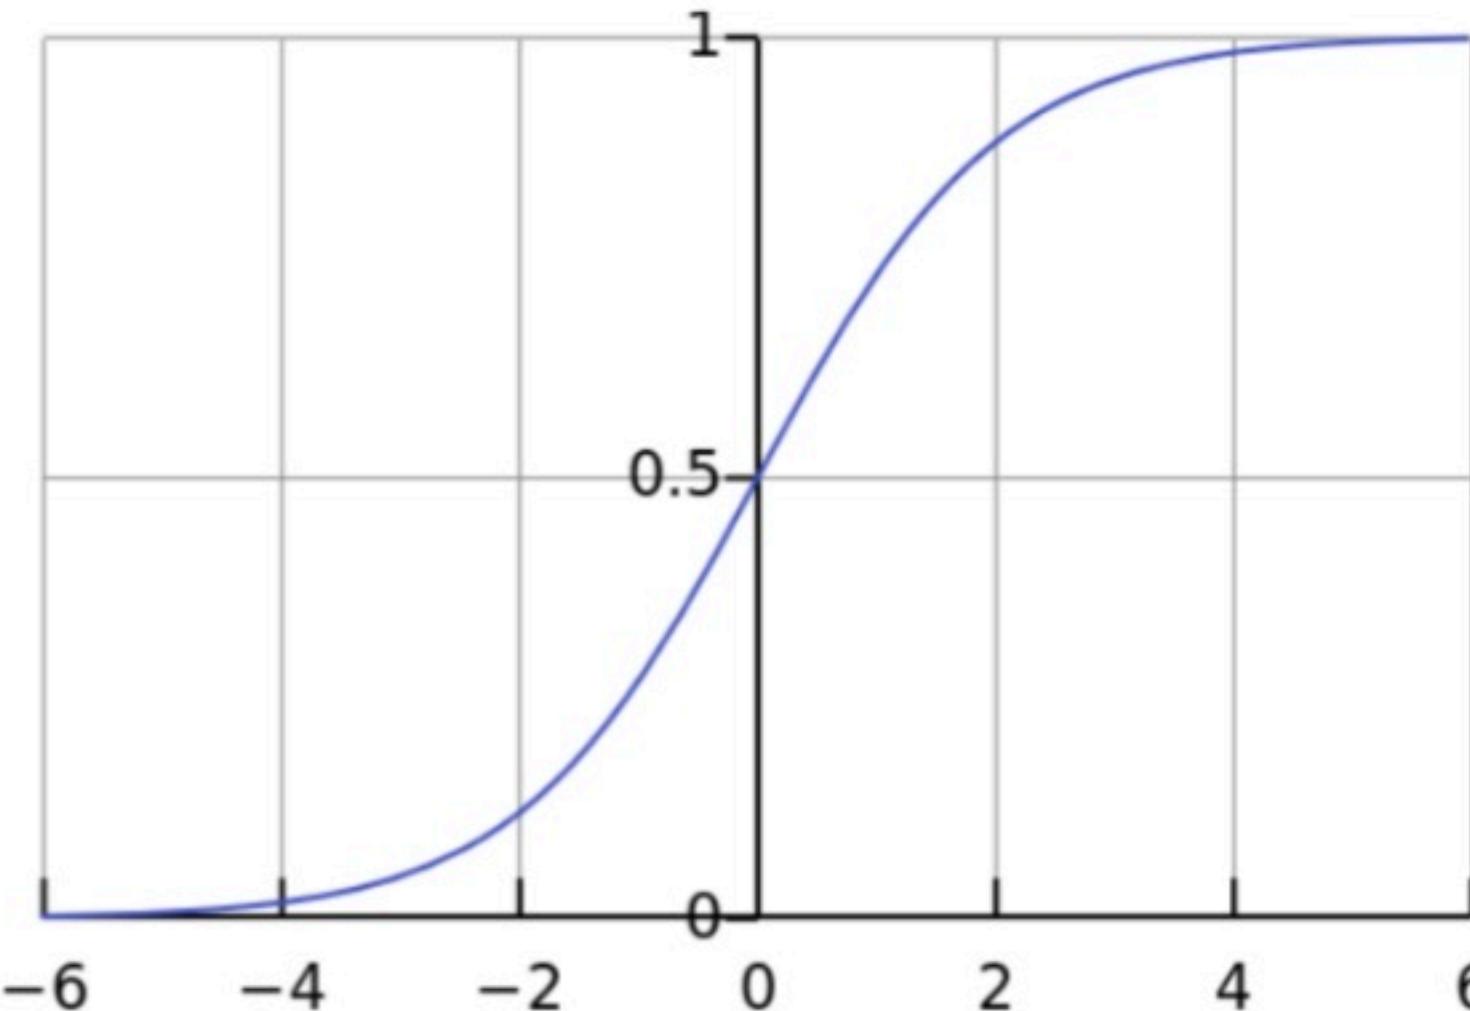
Modern neural network models use non-linear activation functions

Sigmoid Activation Functions



Output values bound between 0 and 1, normalizing the output of each neuron.

$$A = \frac{1}{1+e^{-x}}$$



Pros

- **Smooth gradient**, preventing “jumps” in output values.
- **Clear predictions**—For X above 2 or below -2, tends to bring the Y value (the prediction) to the edge of the curve, very close to 1 or 0. This enables clear predictions.

Cons

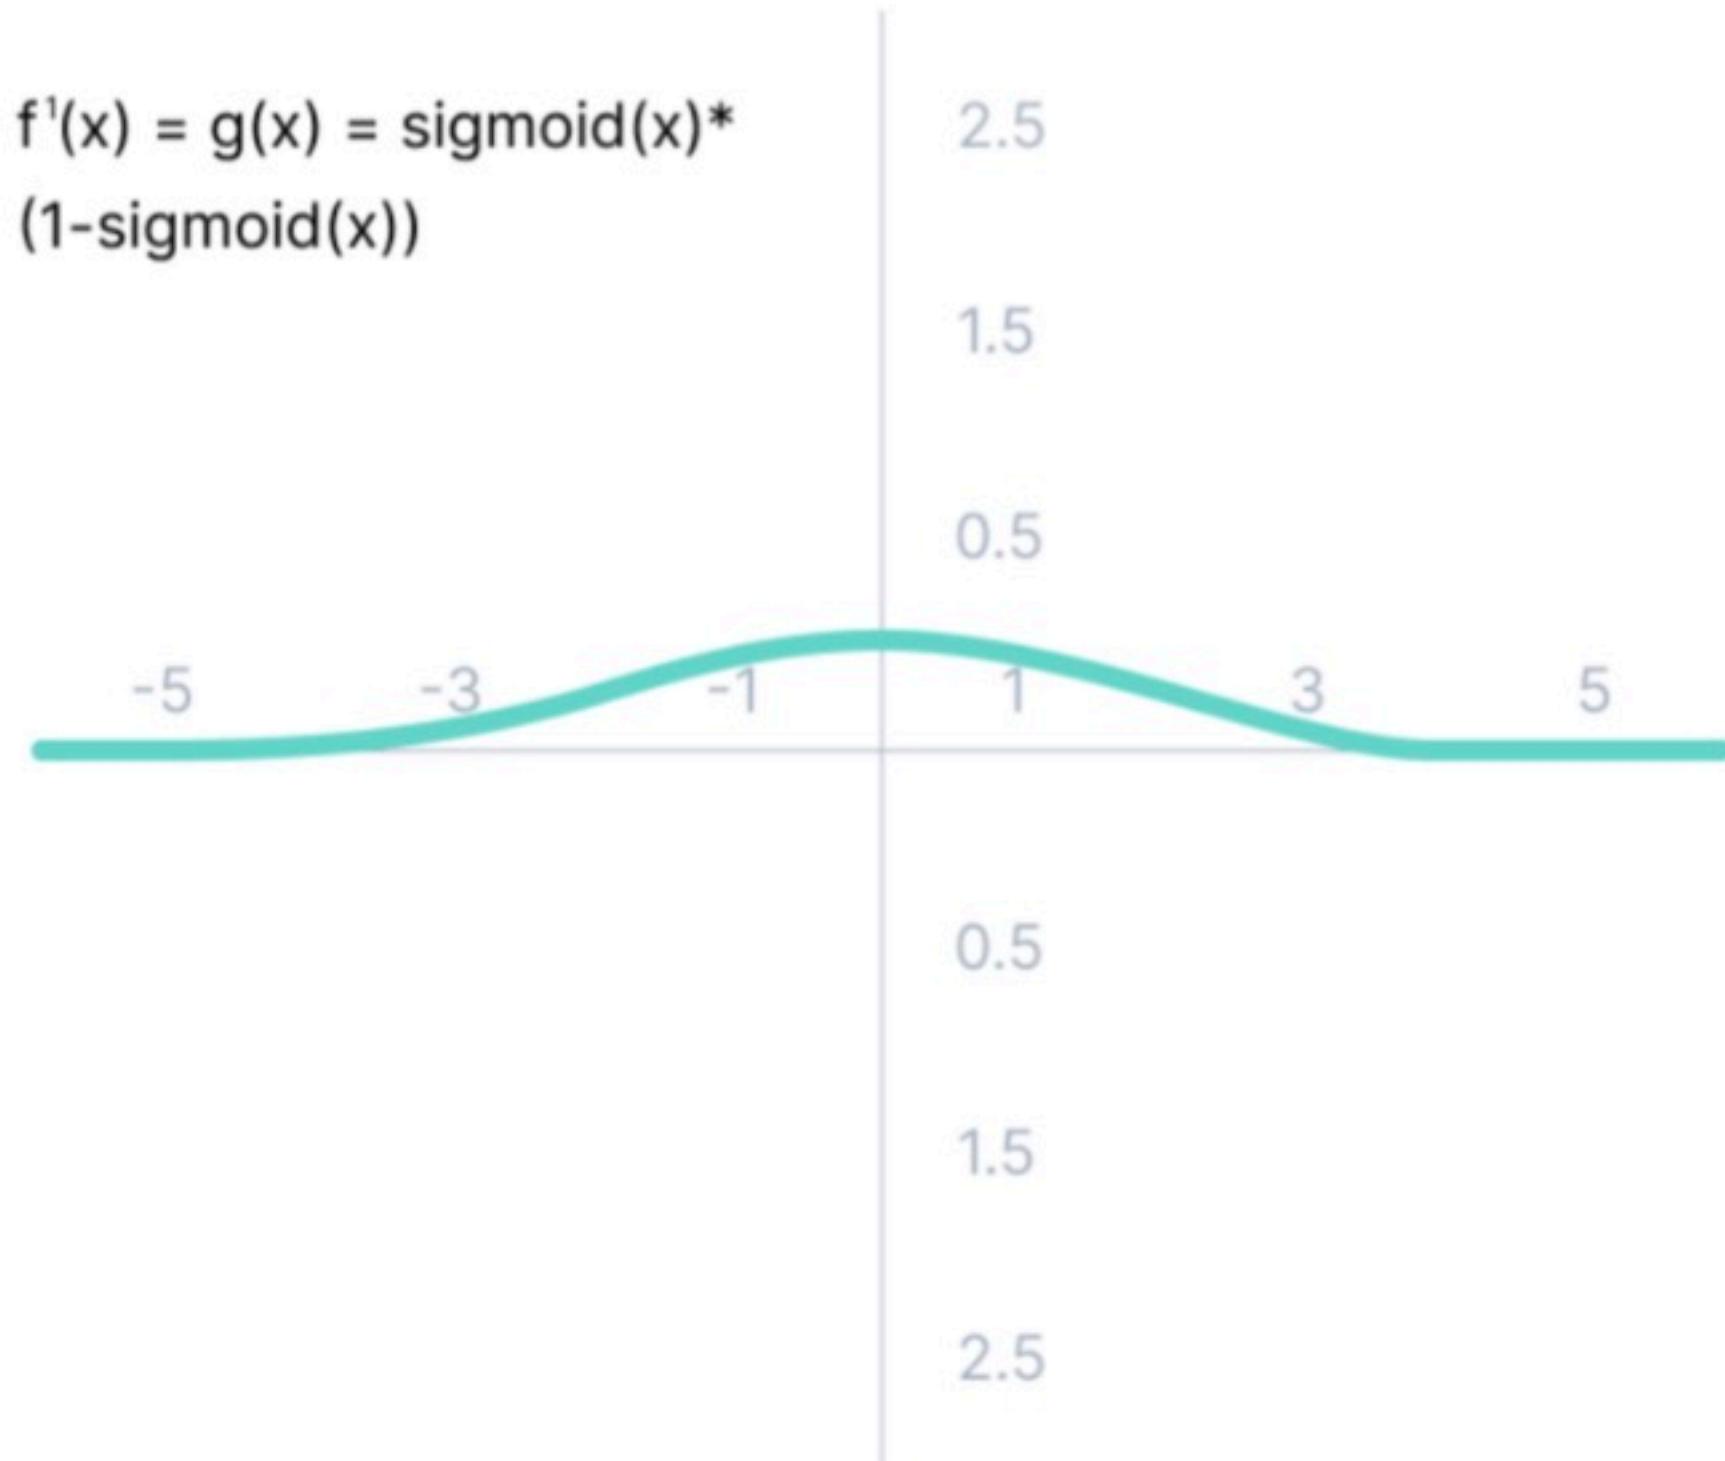
- **Vanishing gradient**—for very high or very low values of X, there is almost no change to the prediction, causing a vanishing gradient problem. This can result in the network refusing to learn further, or being too slow to reach an accurate prediction.
- **Outputs not zero centered**.
- **Computationally expensive**

Sigmoid Activation Functions



Derivative of the Sigmoid Activation Function

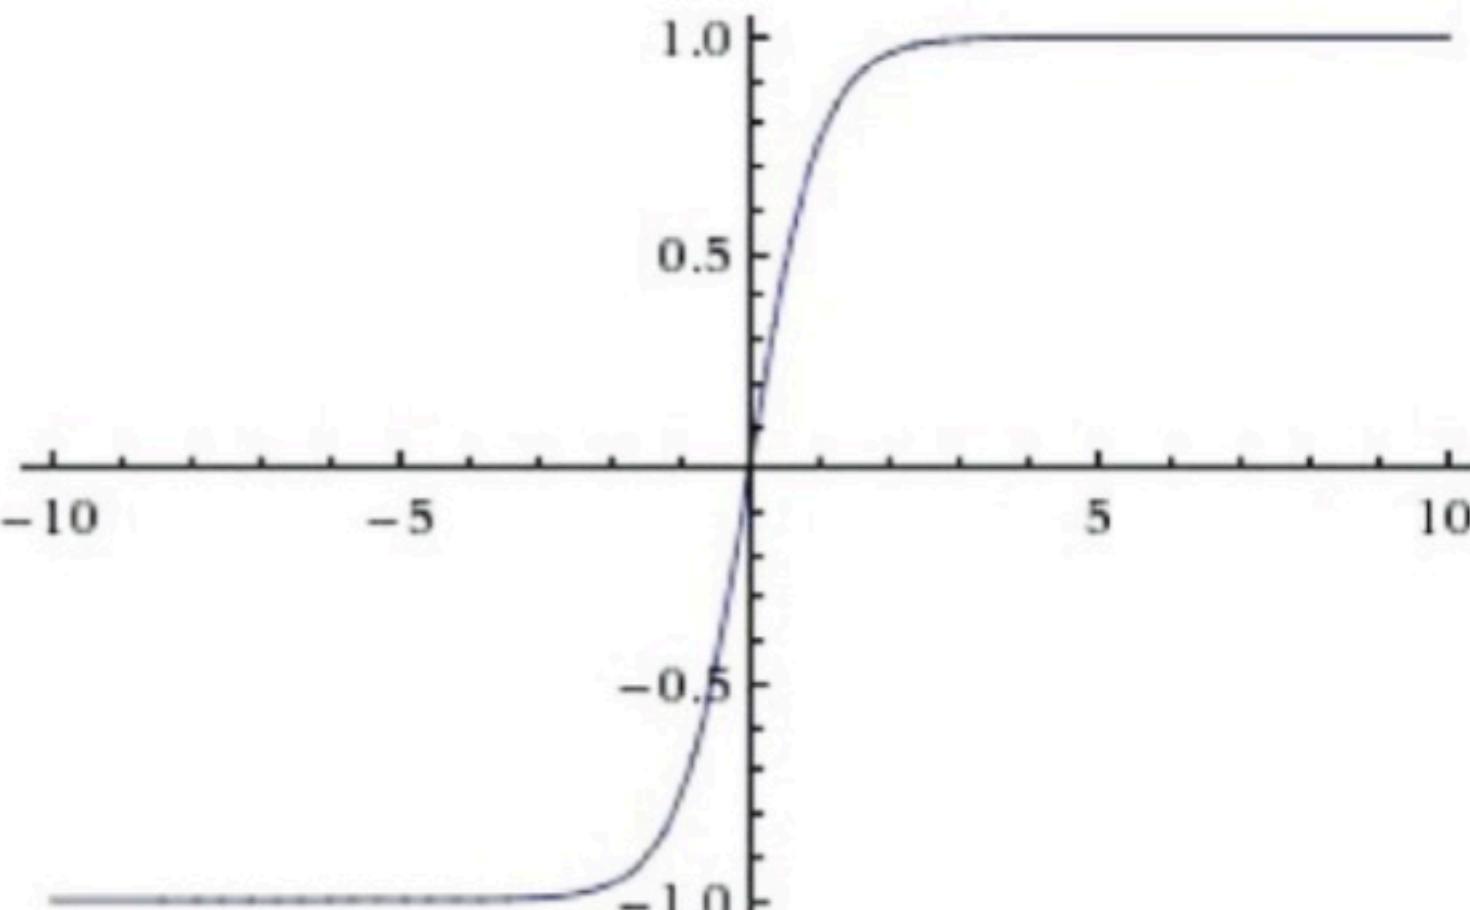
$$f'(x) = g(x) = \text{sigmoid}(x) * (1 - \text{sigmoid}(x))$$



TanH Activation Functions

Output values bound between -1 and 1.

$$f(x) = \tanh(x) = \frac{(e^x - e^{-x})}{(e^x + e^{-x})}$$



Pros

- Same Advantages with Sigmoid.
- More efficient than Sigmoid: Because it has a wider range for faster learning and grading.
- Outputs zero centered.

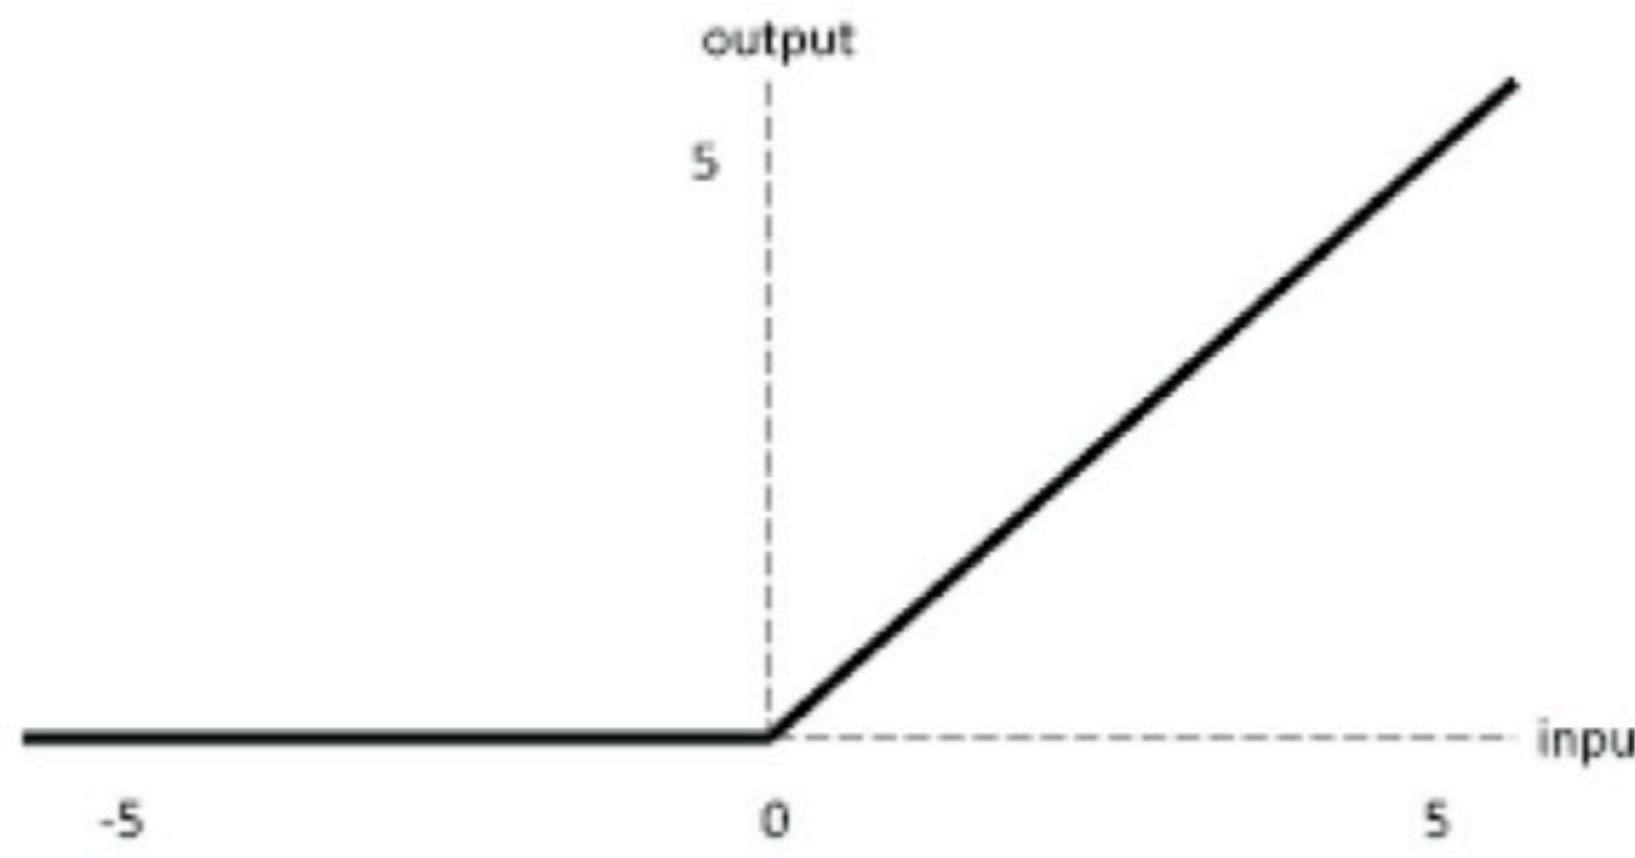
Cons

- Vanishing gradient—for very high or very low values of X, there is almost no change to the prediction, causing a vanishing gradient problem. This can result in the network refusing to learn further, or being too slow to reach an accurate prediction.
- Computationally expensive

ReLU Activation Functions



$$f(x) = \begin{cases} 0 & \text{for } x \leq 0 \\ x & \text{for } x > 0 \end{cases} = \max\{0, x\}$$



Pros

- **Computationally efficient**—allows the network to converge very **quickly**
- **Non-linear**—although it looks like a linear function, ReLU has a derivative function and allows for backpropagation

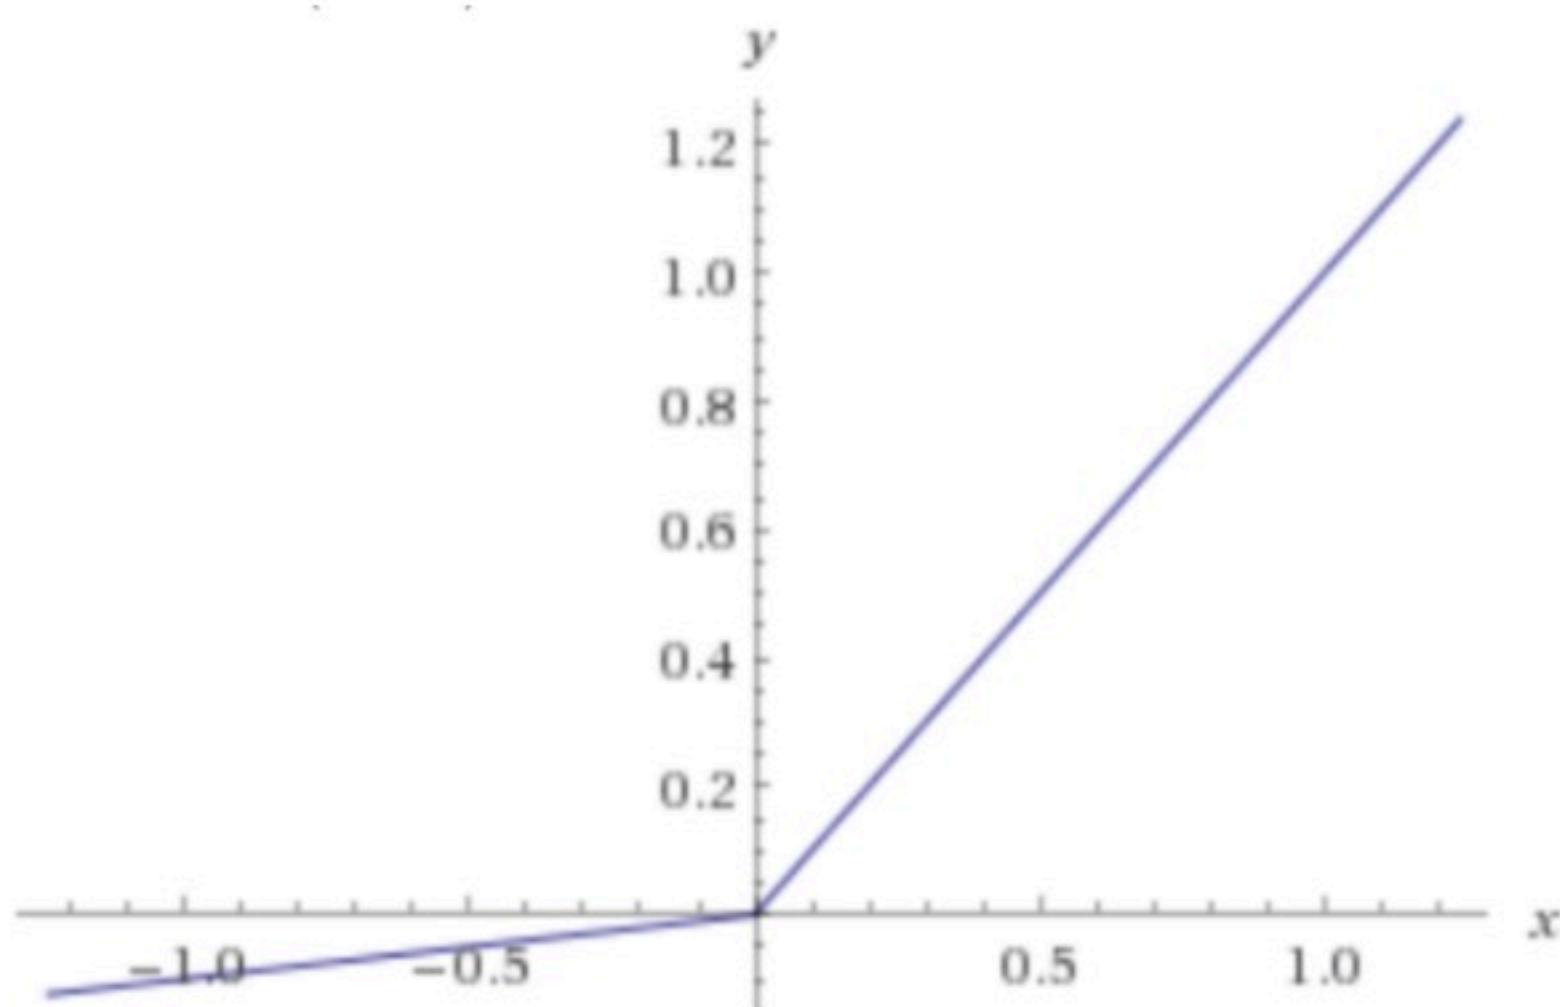
Cons

- **The Dying ReLU problem**—when inputs approach zero, or are negative, the gradient of the function becomes zero, the network cannot perform backpropagation and cannot learn.

Leaky ReLu Activation Functions



$$f(x) = \begin{cases} 0.01x & \text{for } x < 0 \\ x & \text{for } x \geq 0 \end{cases}$$



Pros

- **Prevents dying ReLU problem**—this variation of ReLU has a small positive slope in the negative area, so it does enable backpropagation, even for negative input values
- Otherwise like ReLU

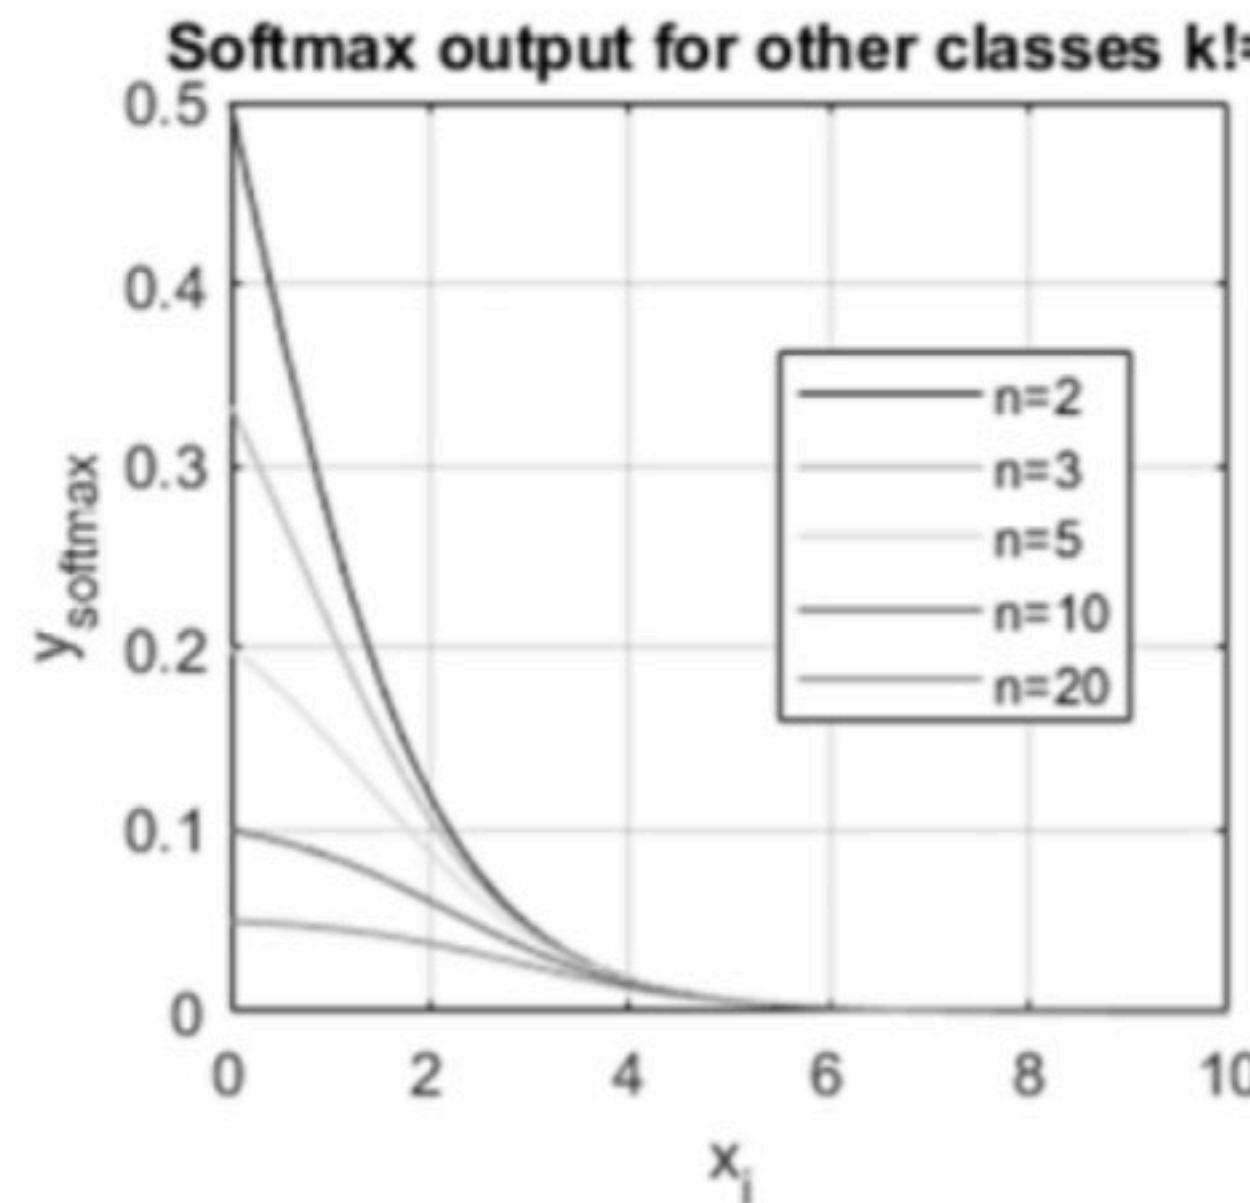
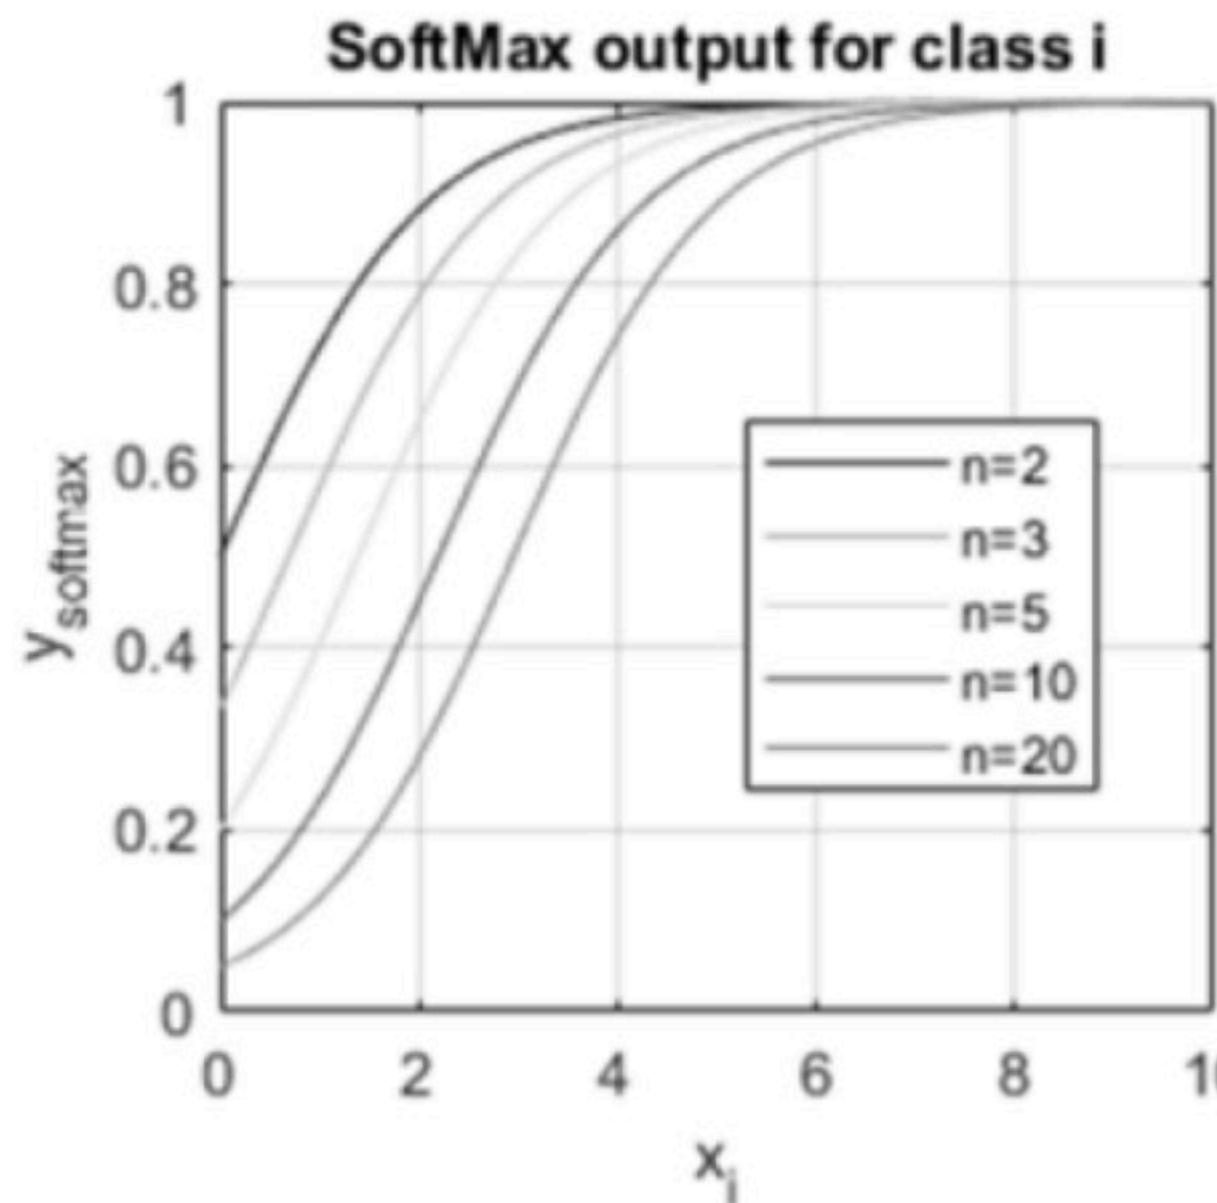
Cons

- **Results not consistent**—leaky ReLU does not provide consistent predictions for negative input values.

Softmax Activation Functions



$$\frac{e^{z_i}}{\sum_{j=1}^K e^{z_j}}$$



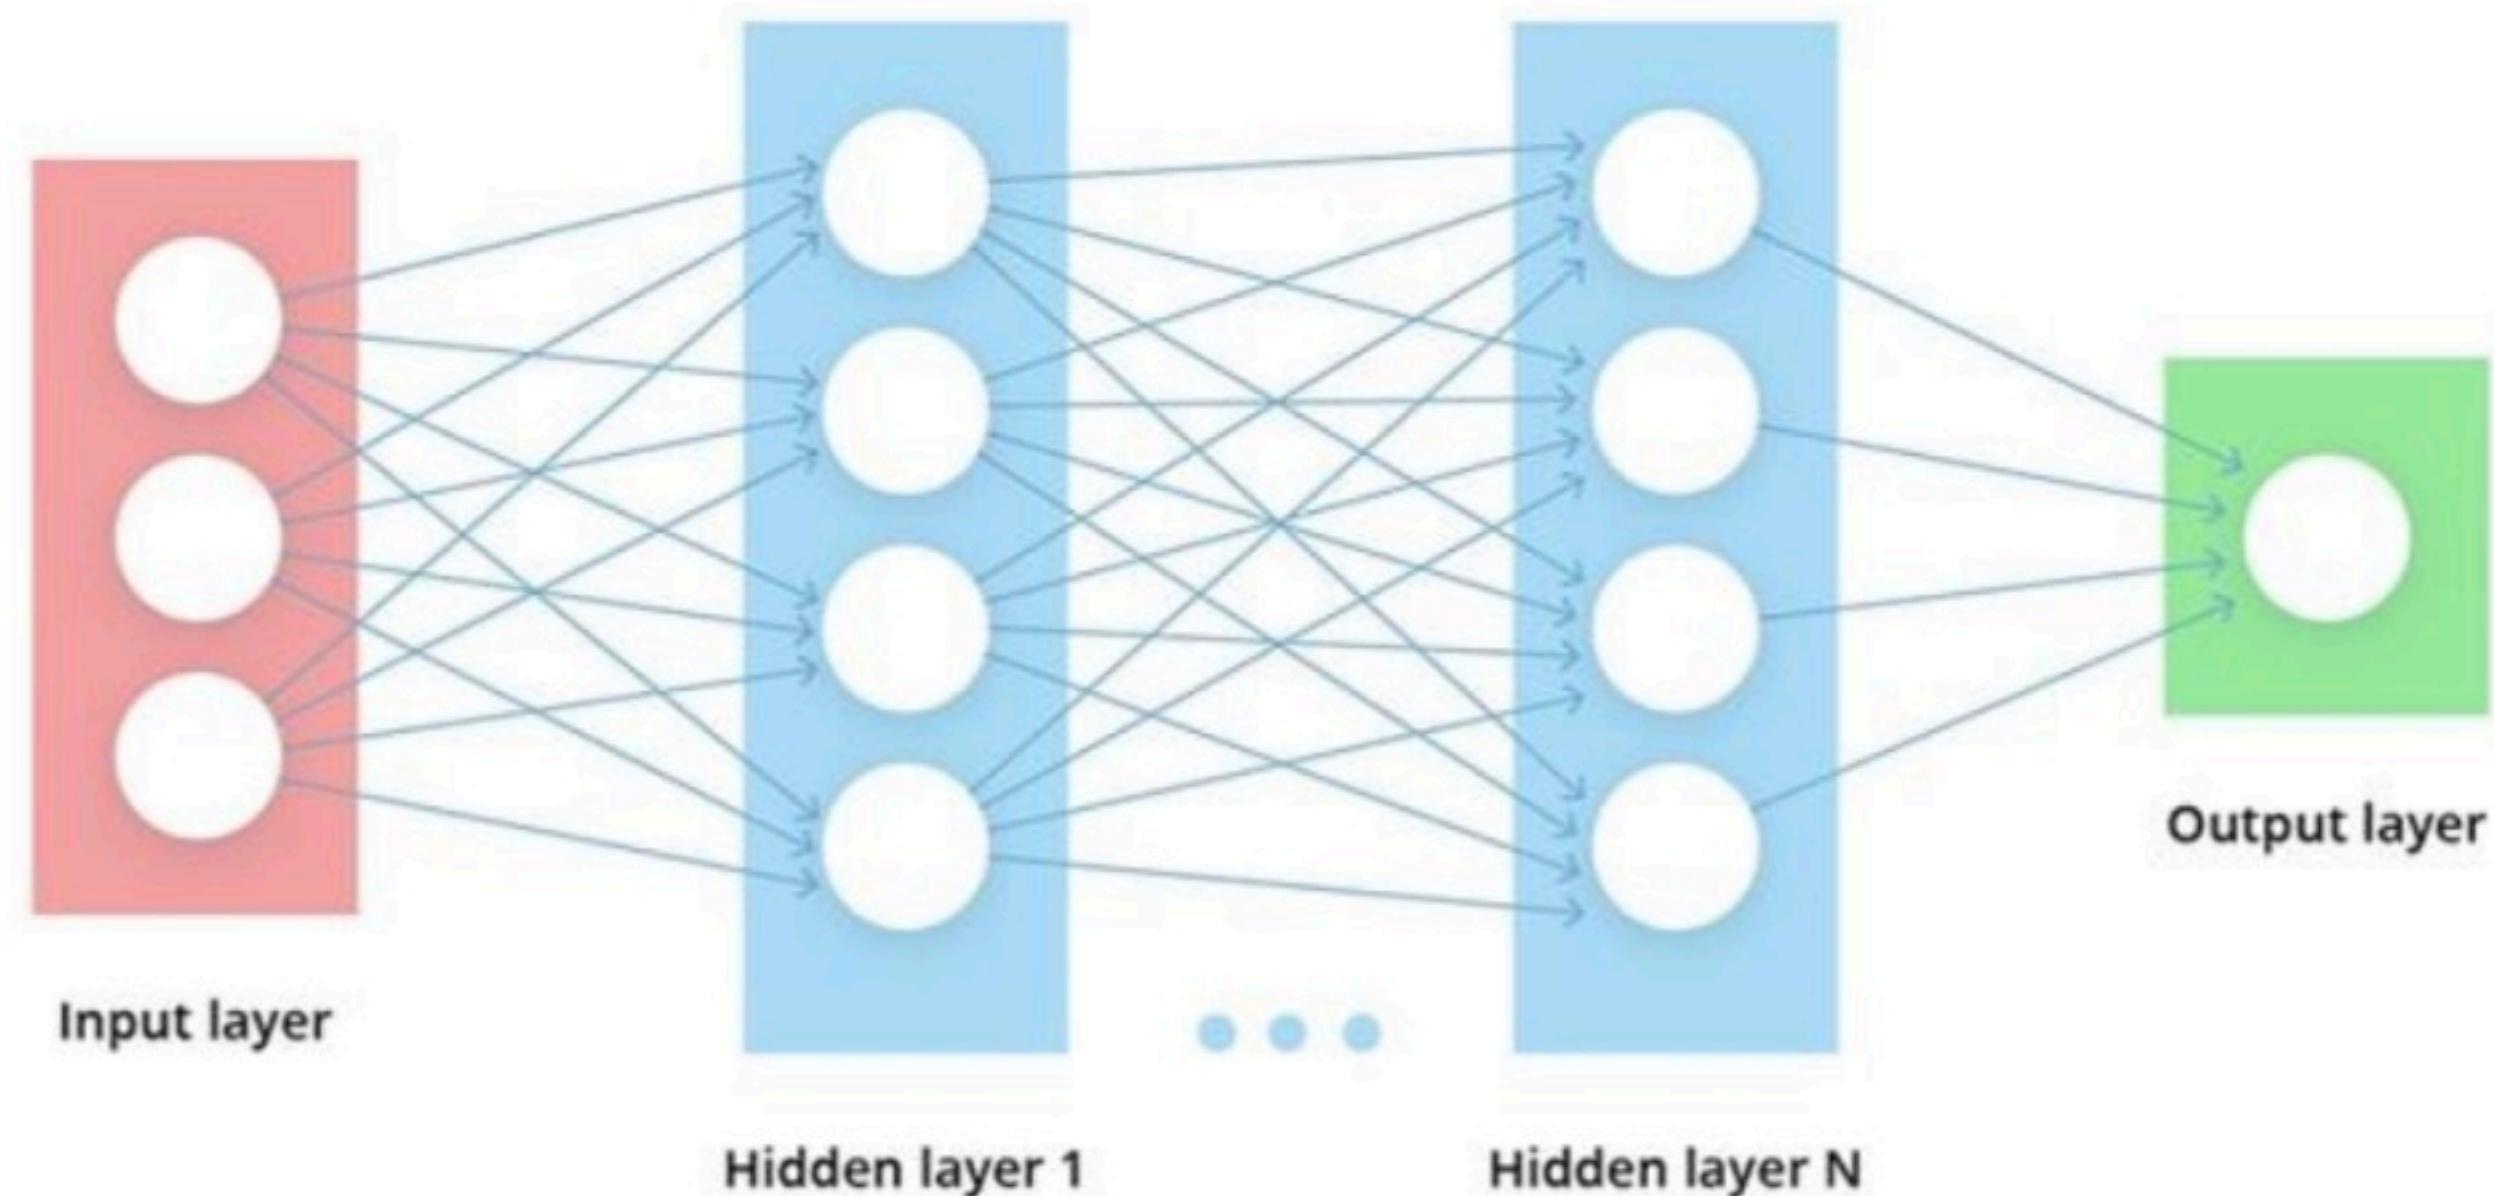
- **Able to handle multiple classes** only one class in other activation functions—normalizes the outputs for each class between 0 and 1
- **Useful for output neurons**—typically Softmax is used only for the output layer



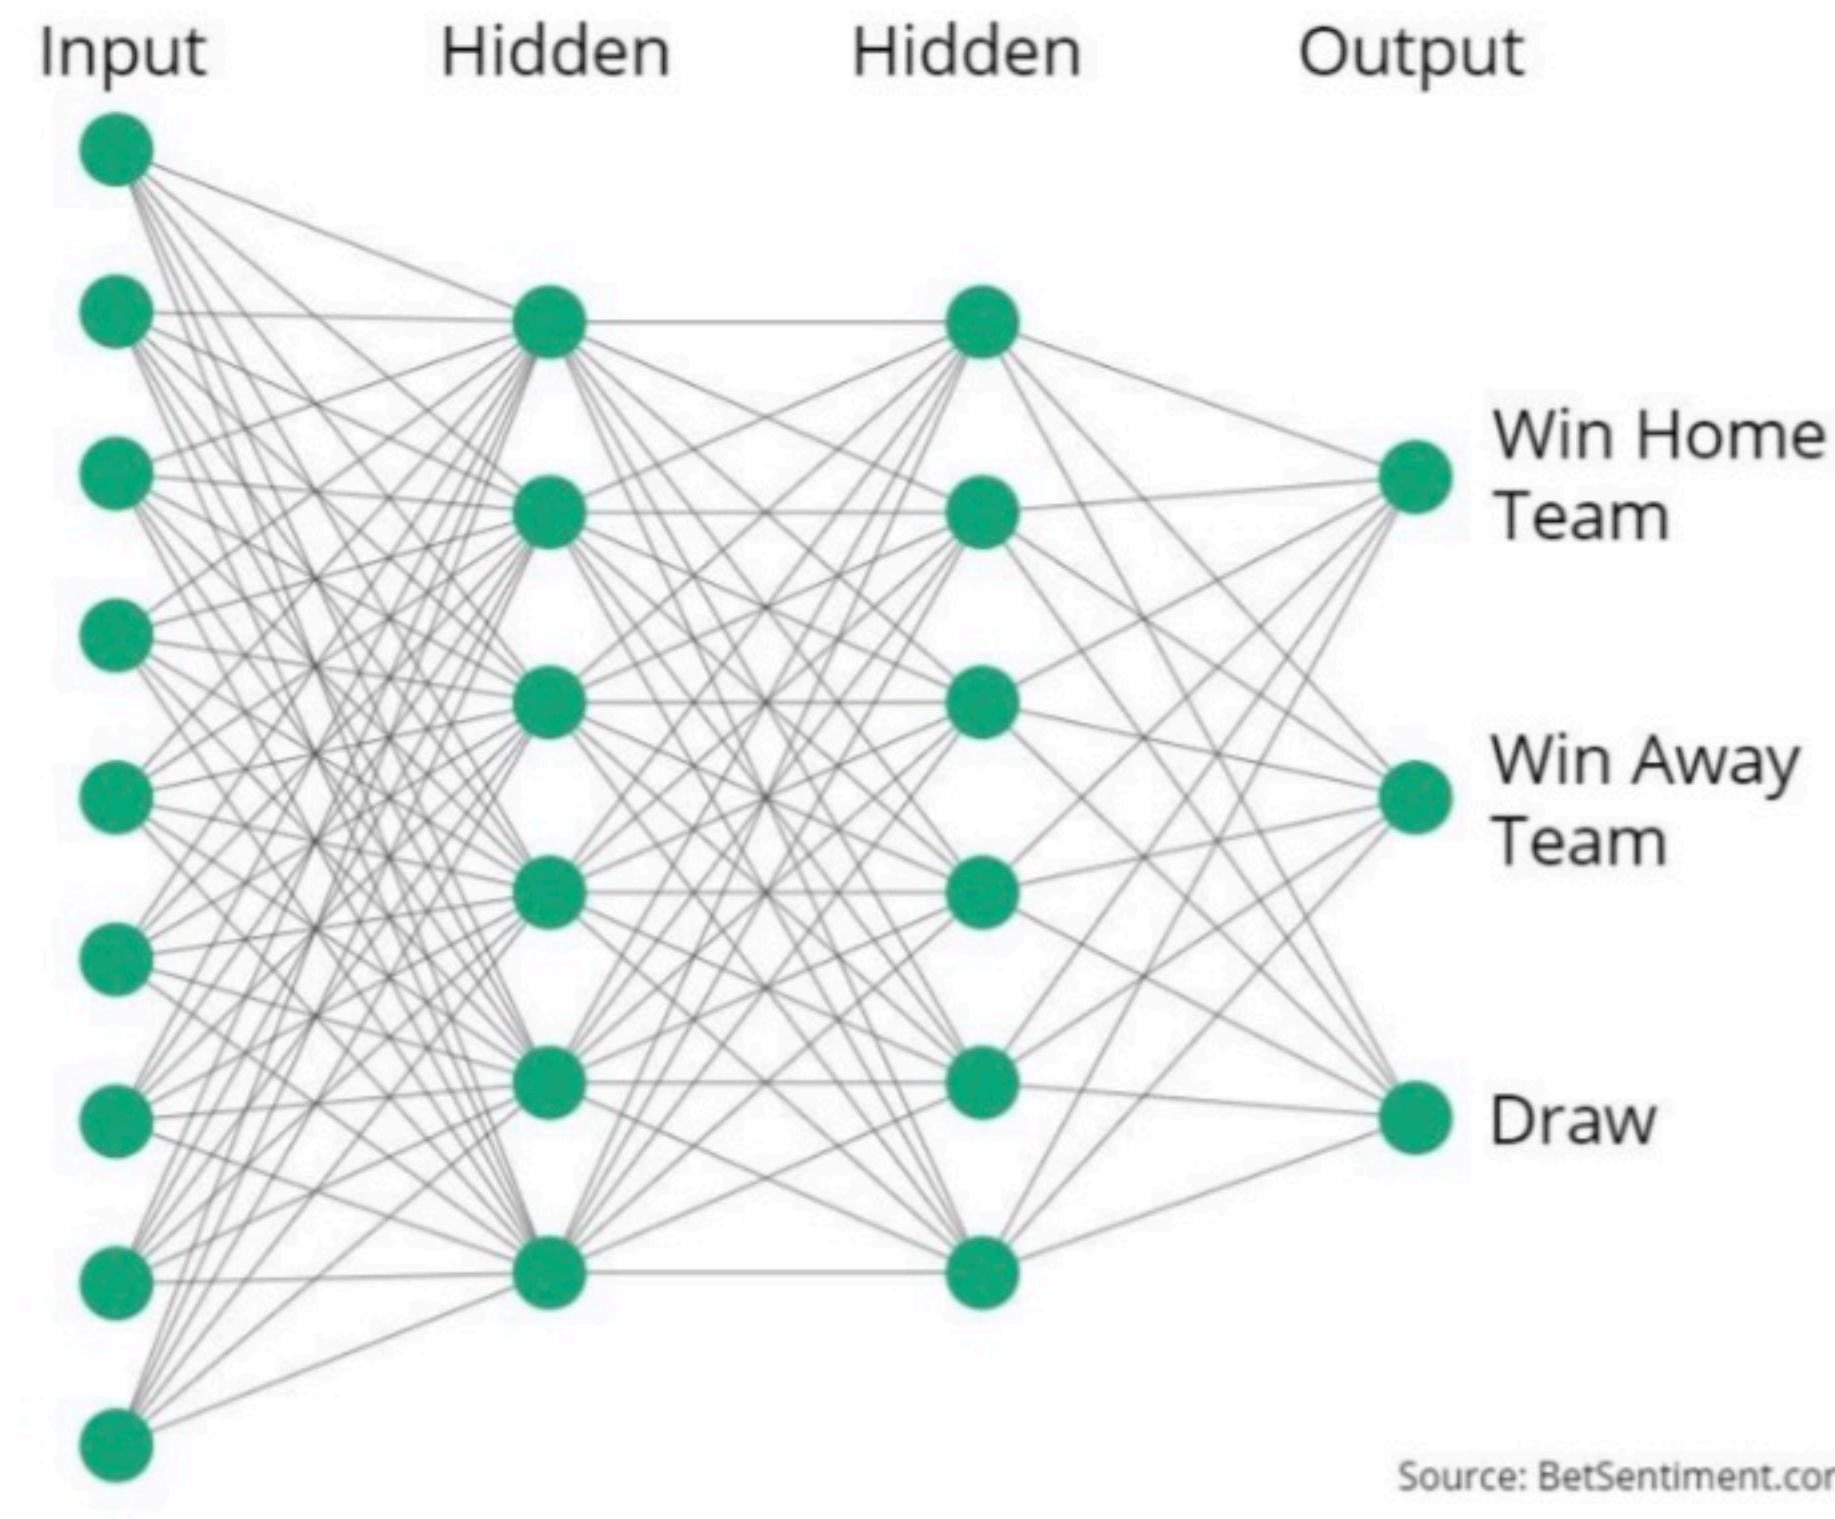
Classification



Binary Classification



Multiclass Classification



Source: BetSentiment.com

Multiclass Classification



Multi Classes

Mutually Exclusive Classes

Data Point 1	RED
Data Point 2	GREEN
Data Point 3	BLUE
...	...
Data Point N	RED

Non-Exclusive Classes

Data Point 1	A,B
Data Point 2	A
Data Point 3	C,B
...	...
Data Point N	B

Softmax Activation Function

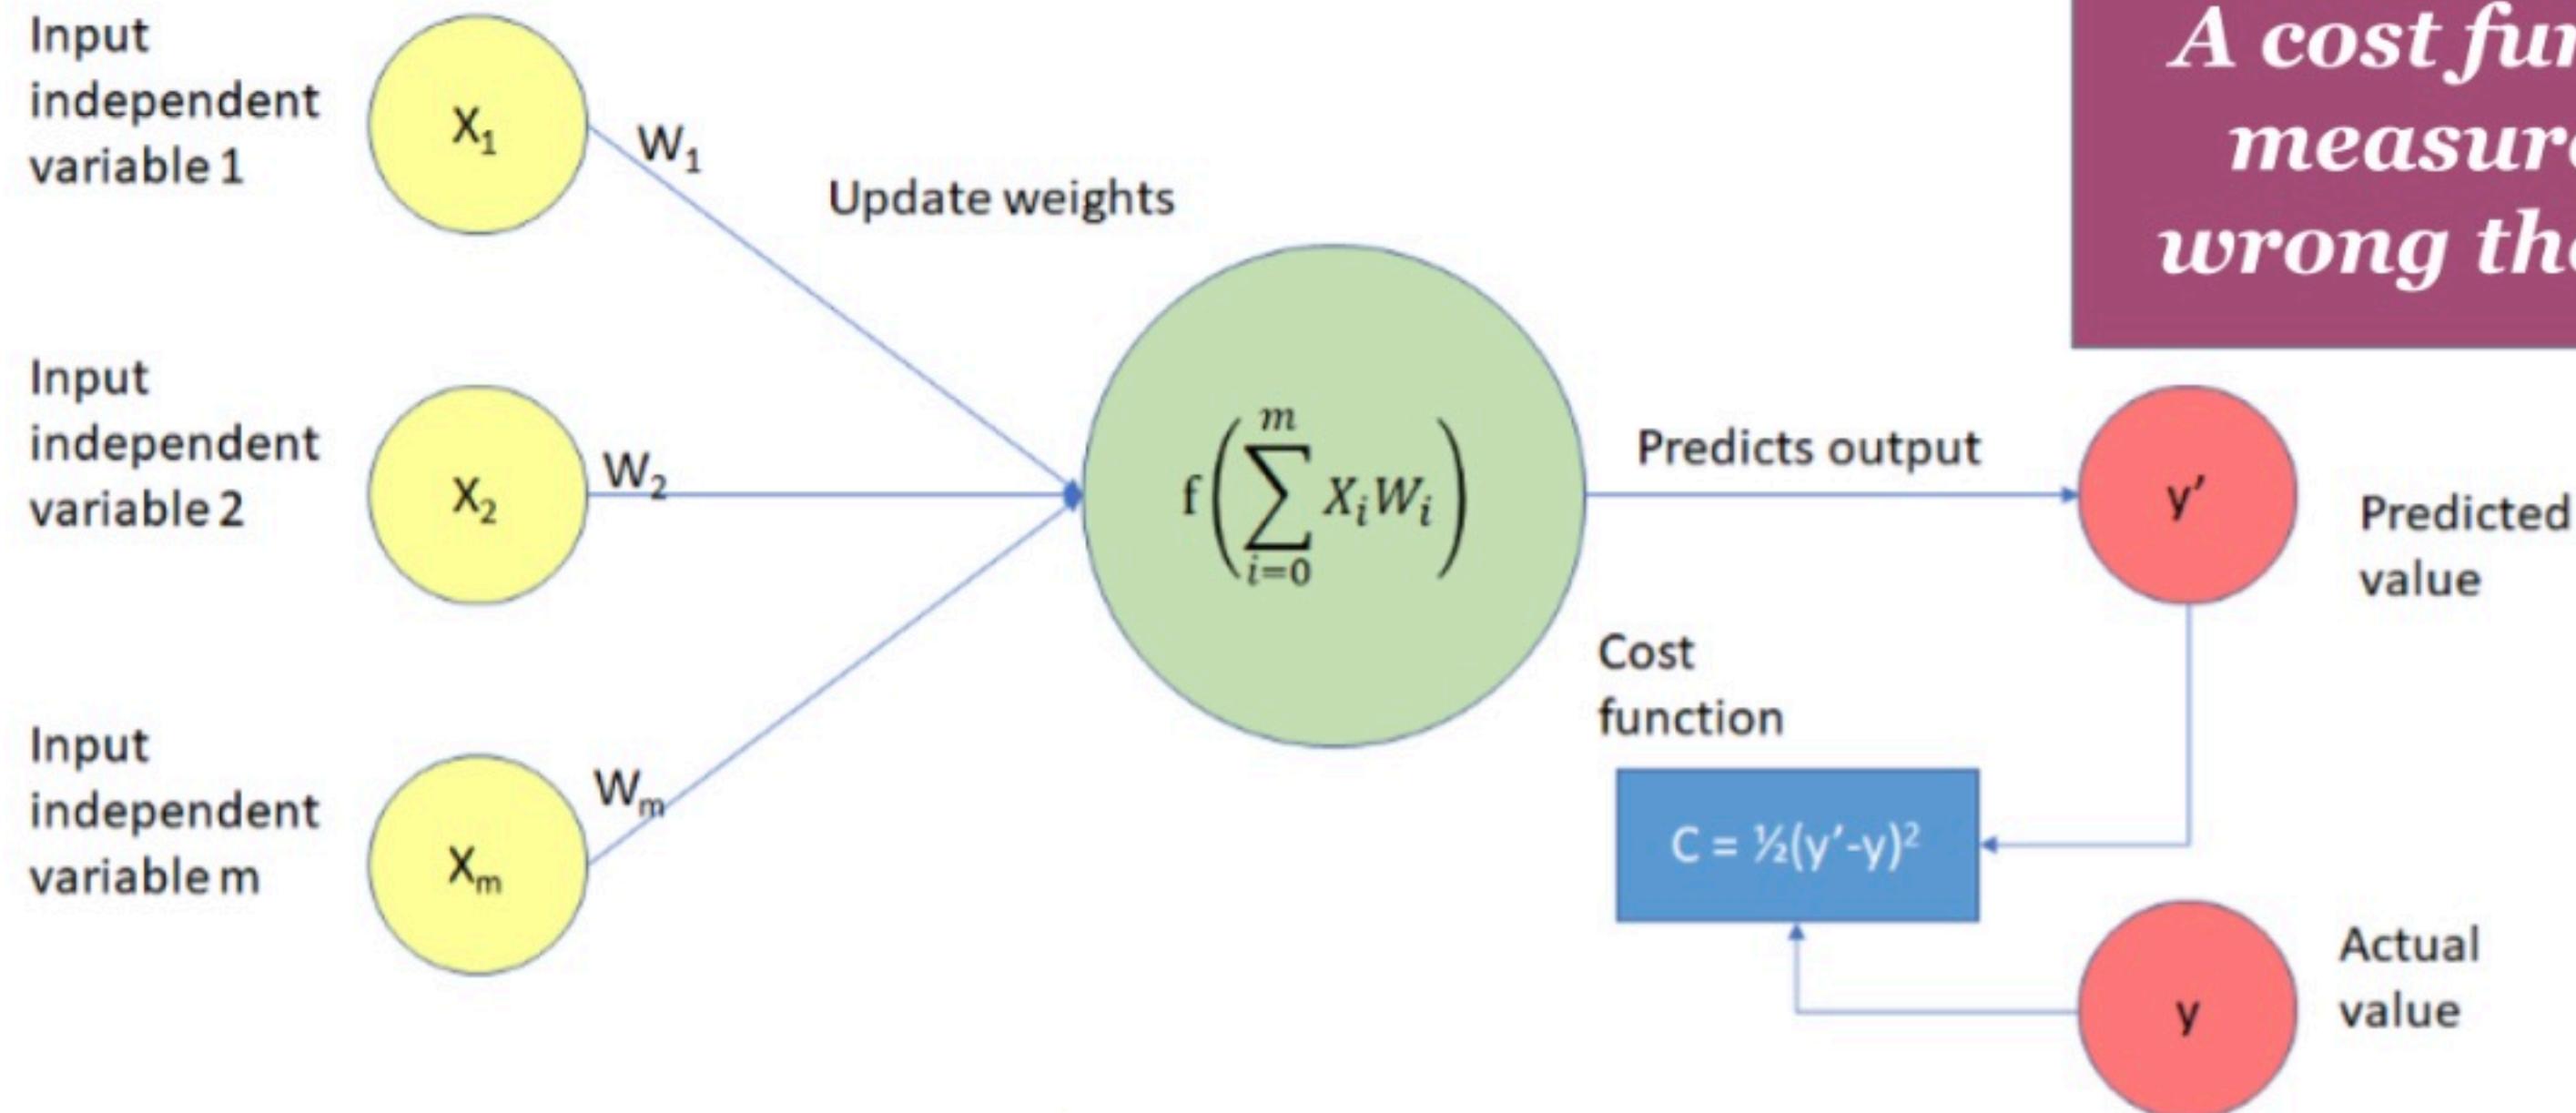
Sigmoid Activation Function



Cost Function



Cost Function (Loss Function)



Cost function $J = \sum_i \frac{1}{2} (y - y')^2$

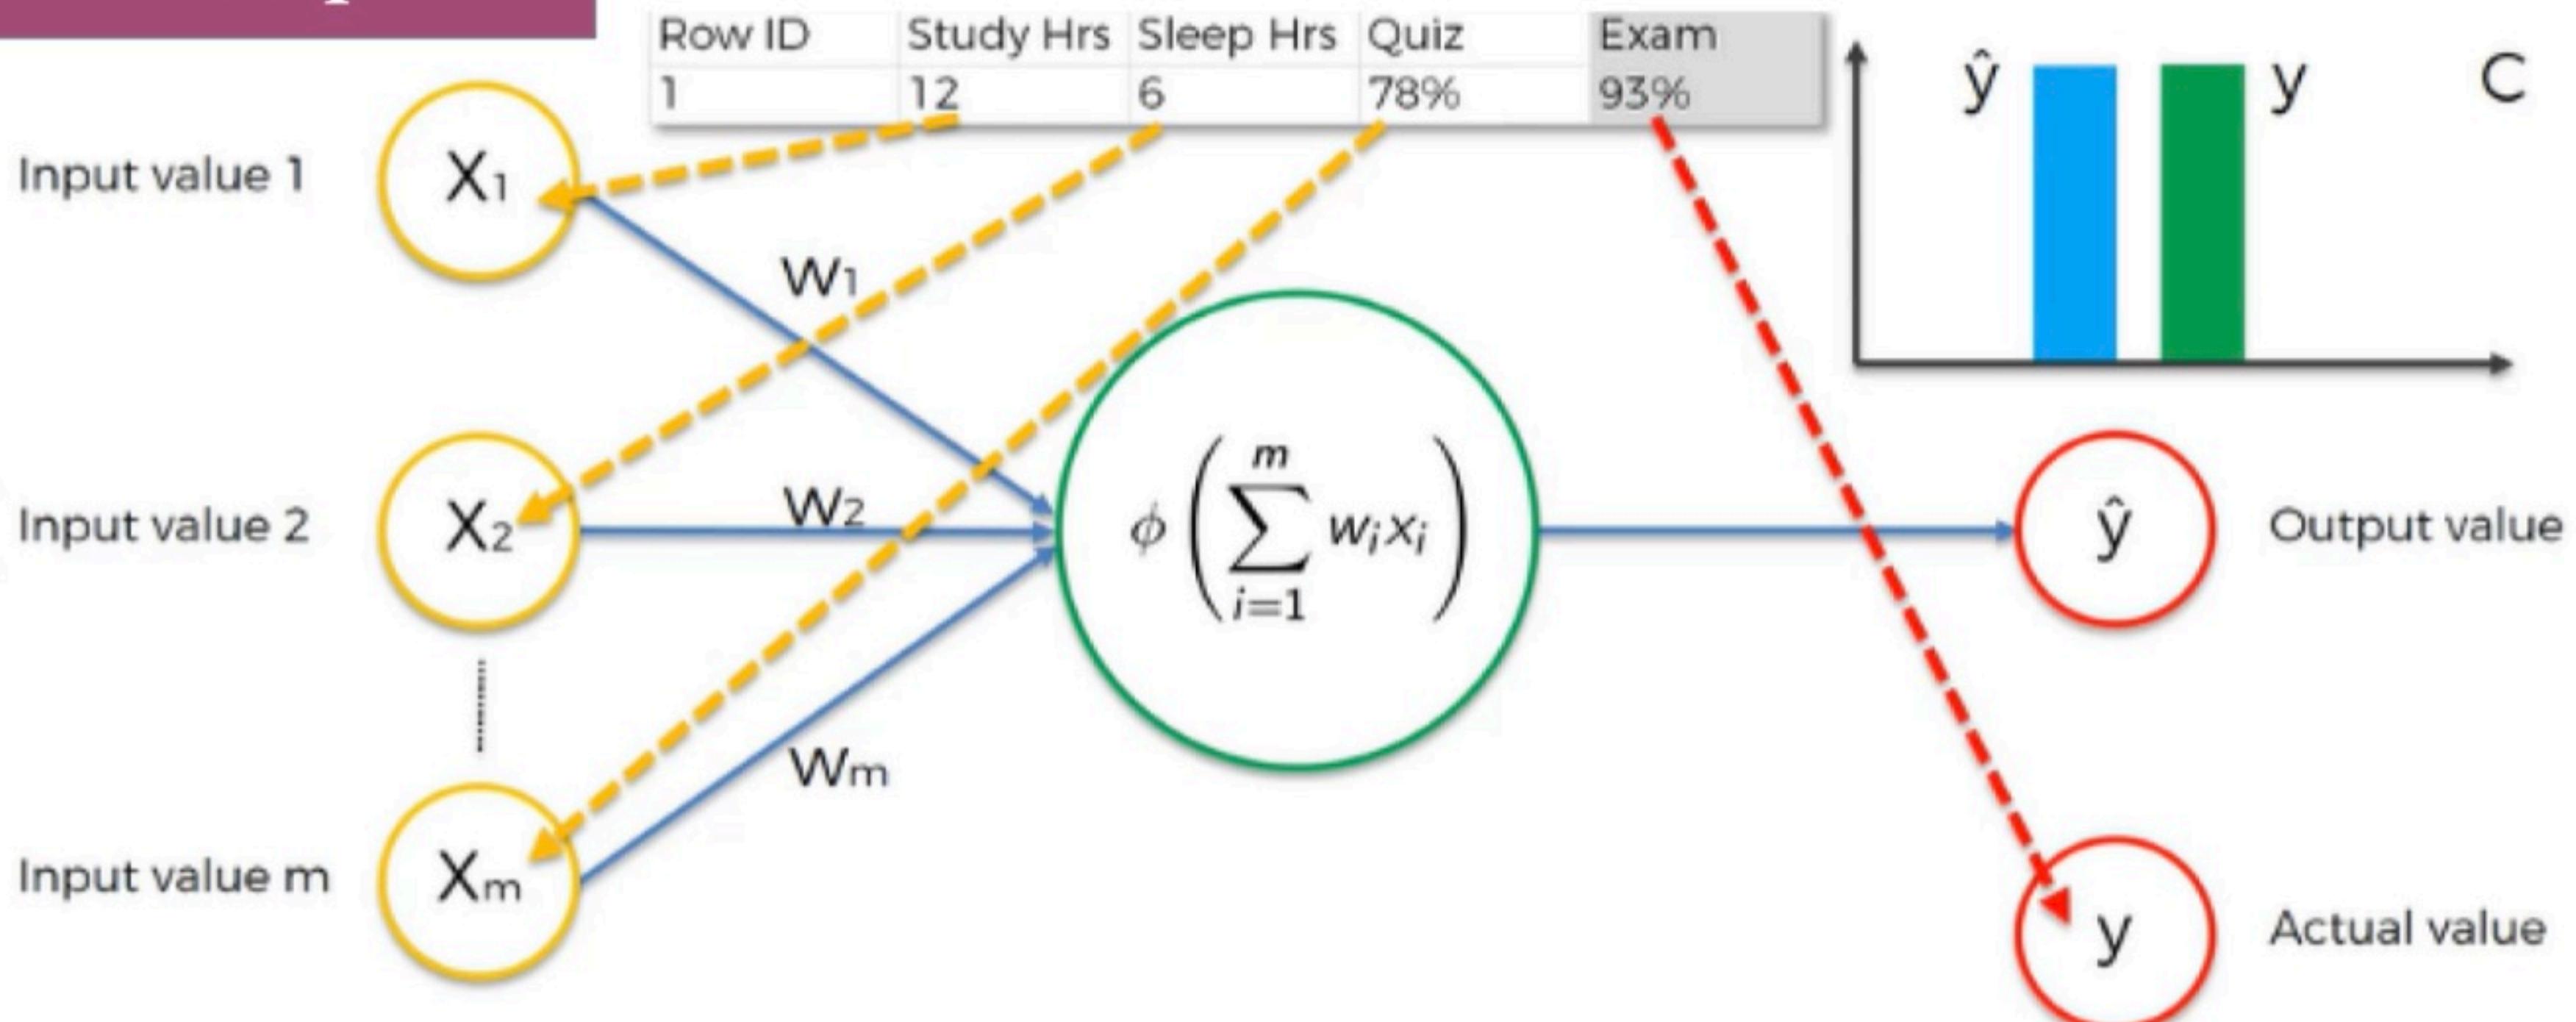
Sum over all samples true value predicted value

Cost function (J) = $1/m$ (Sum of Loss error for 'm' examples)

A **cost function** is a measure of how wrong the model is

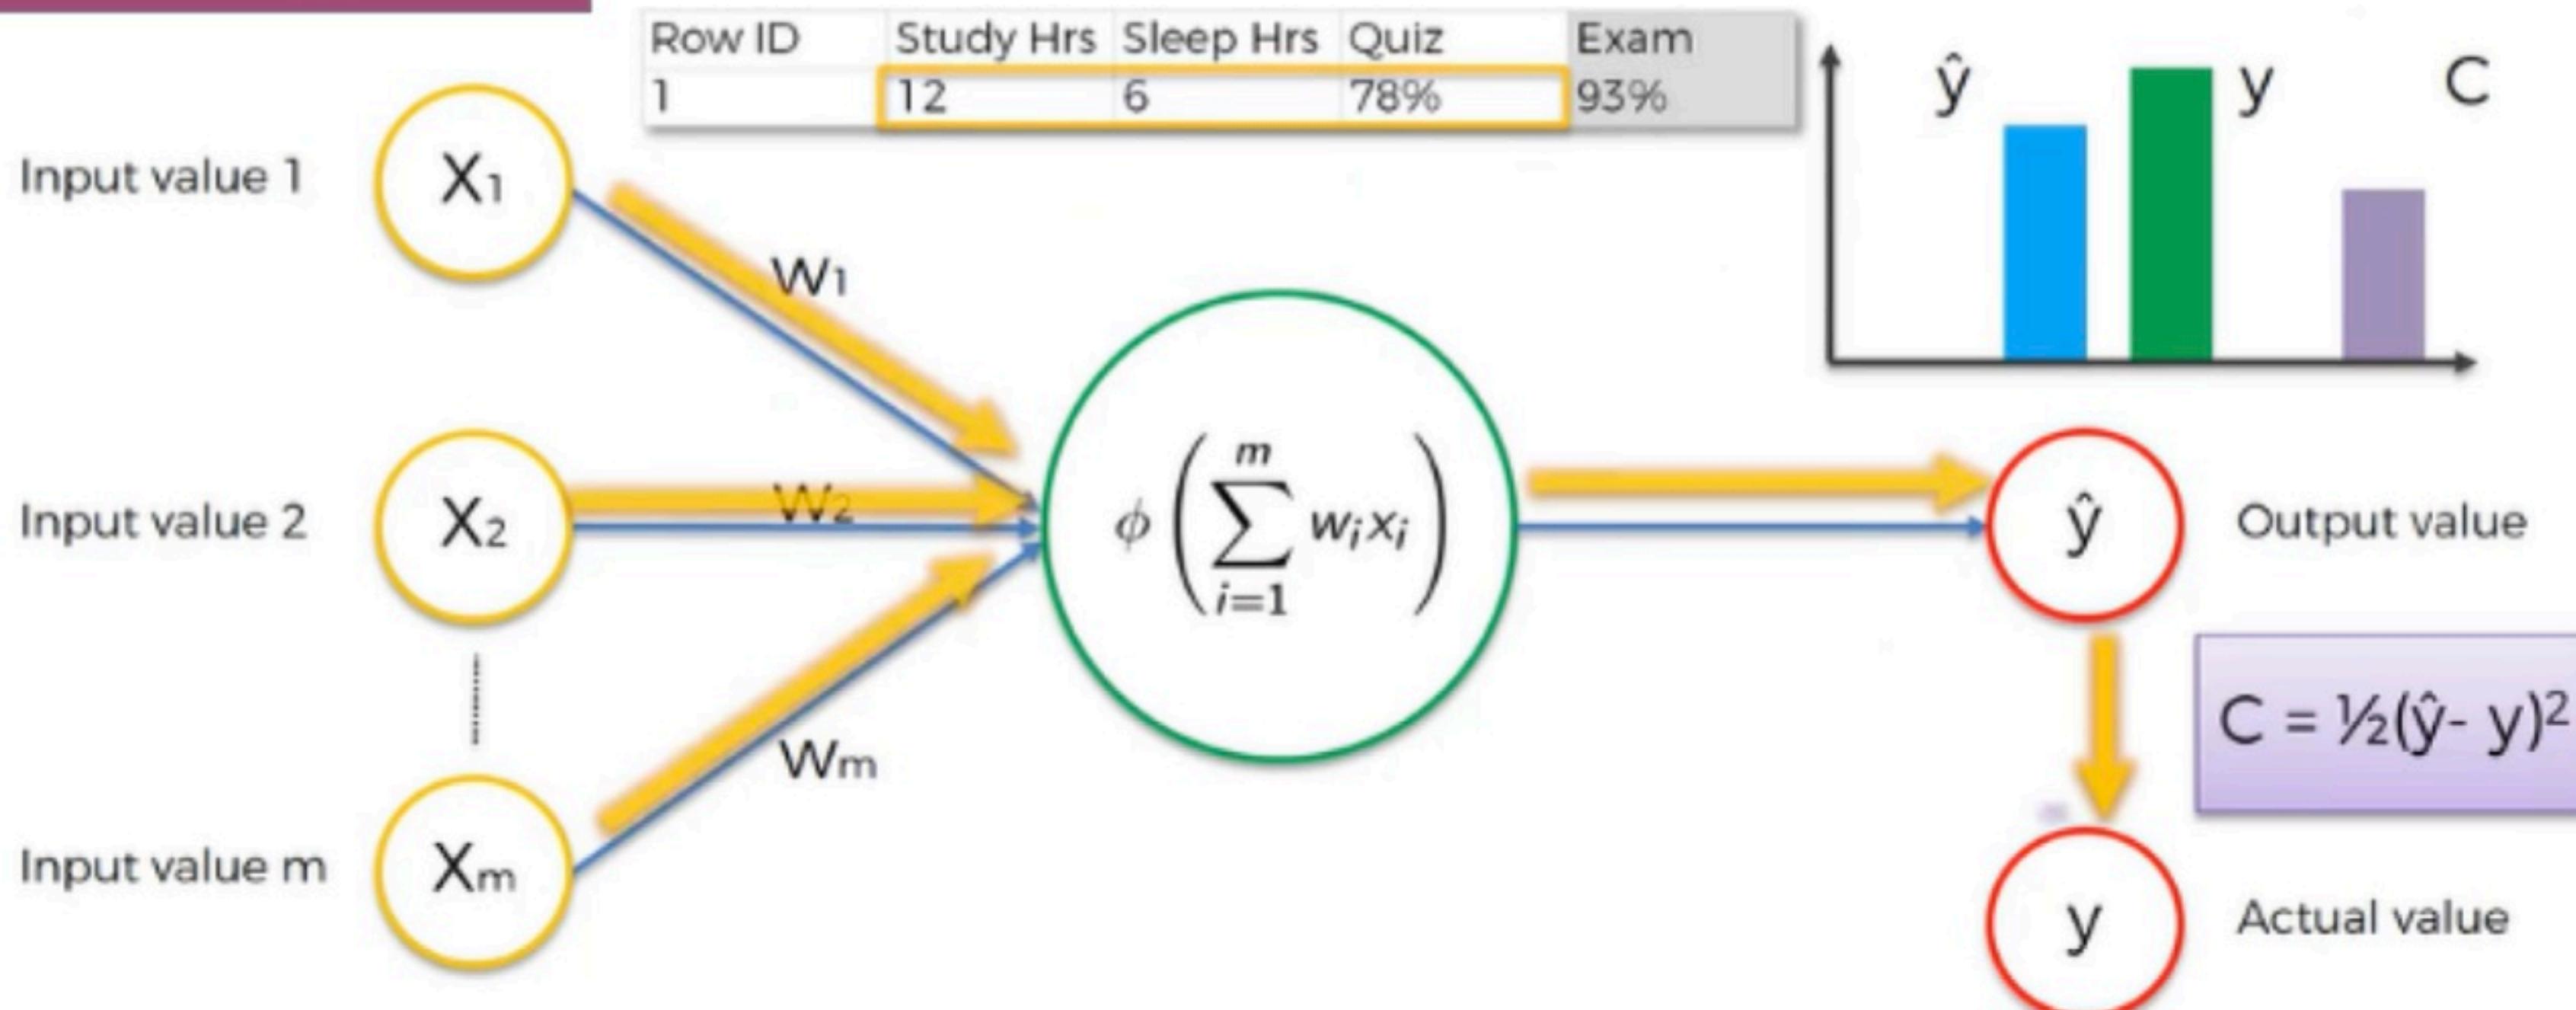
Cost Function (Loss Function)

Example



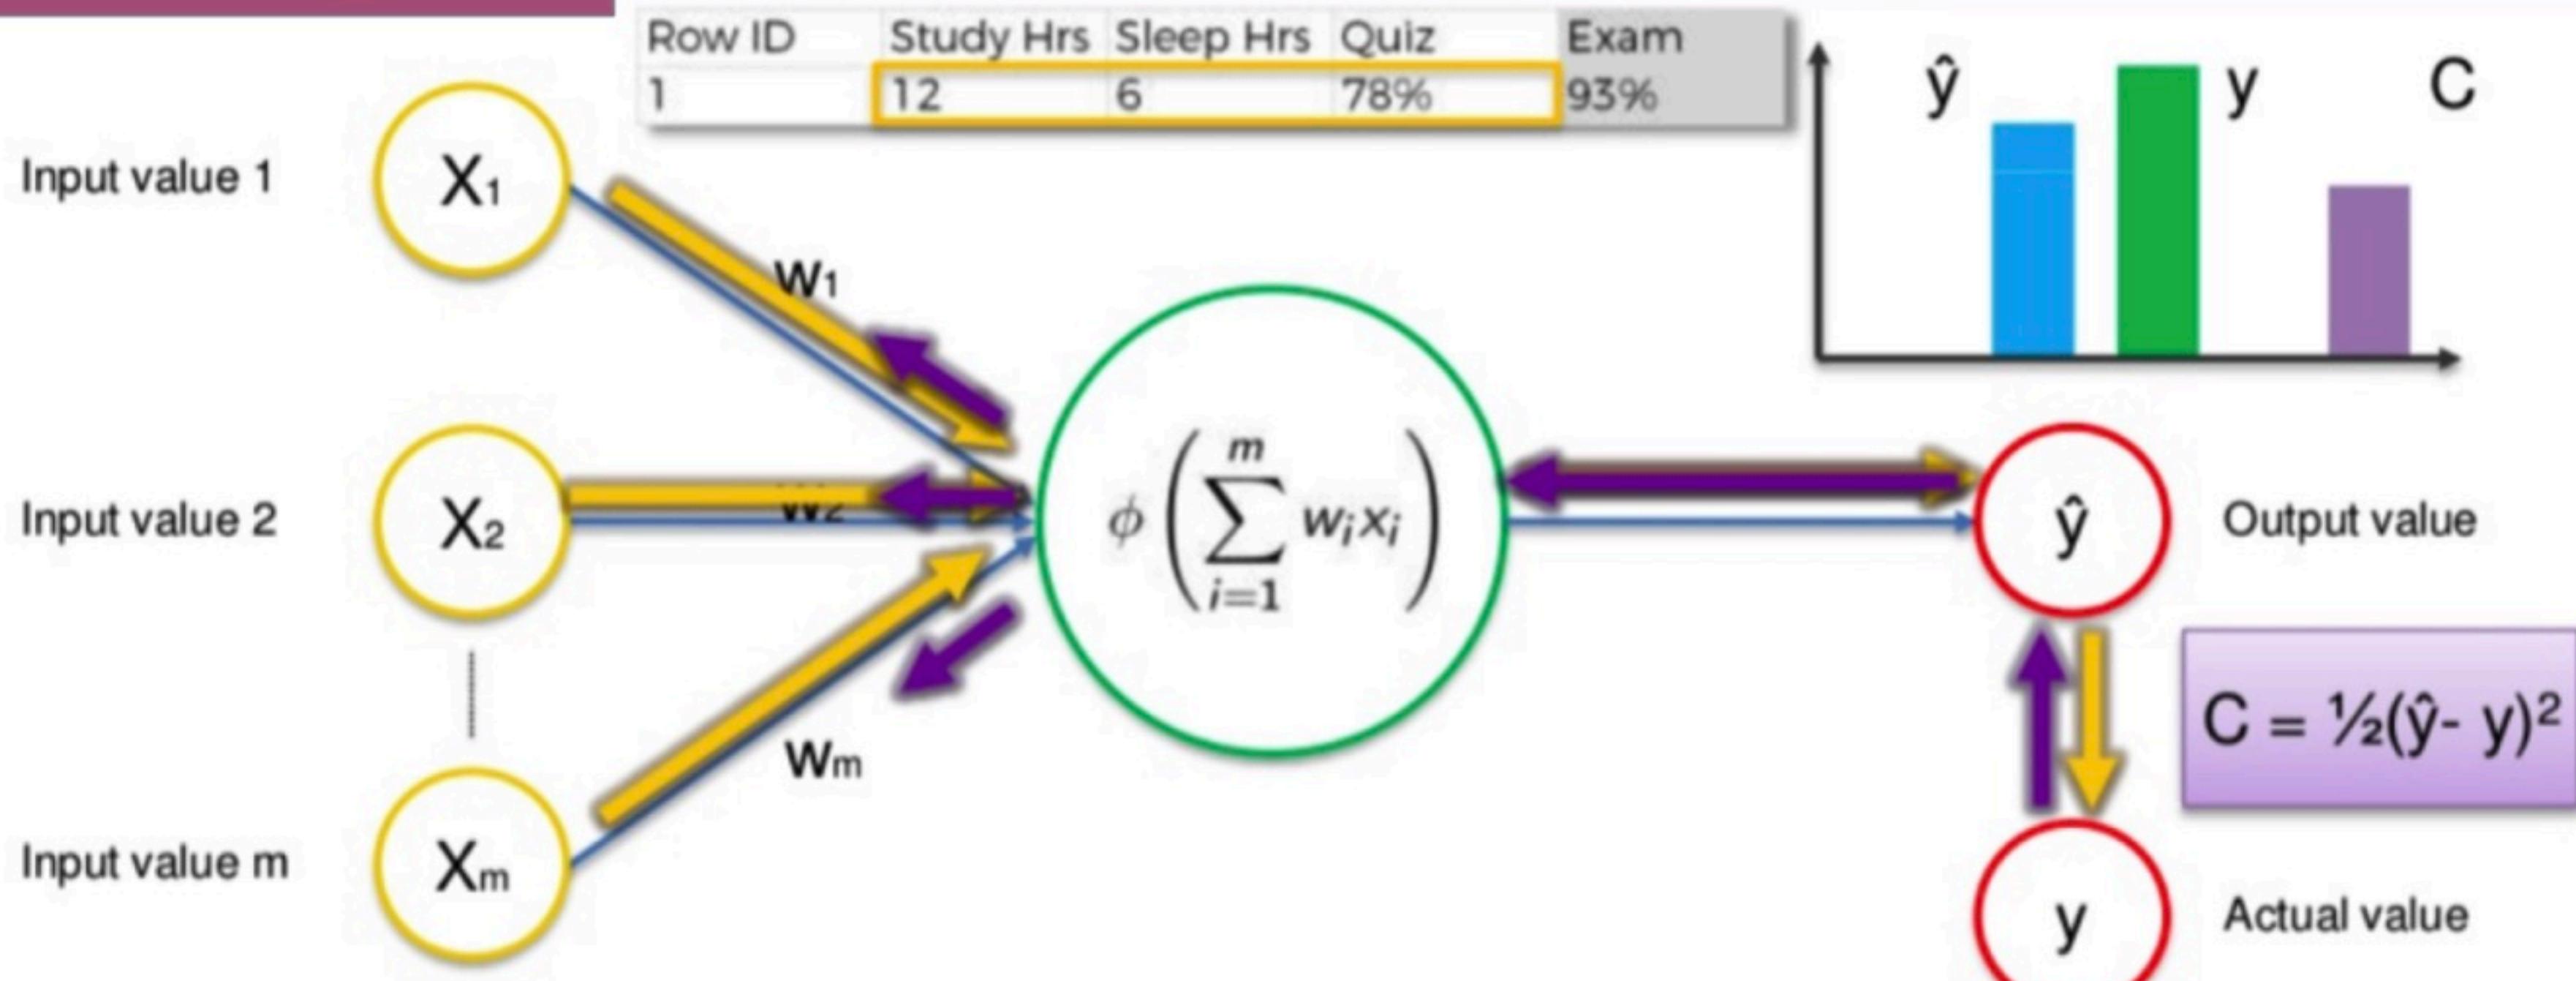
Cost Function (Loss Function)

Example



Cost Function (Loss Function)

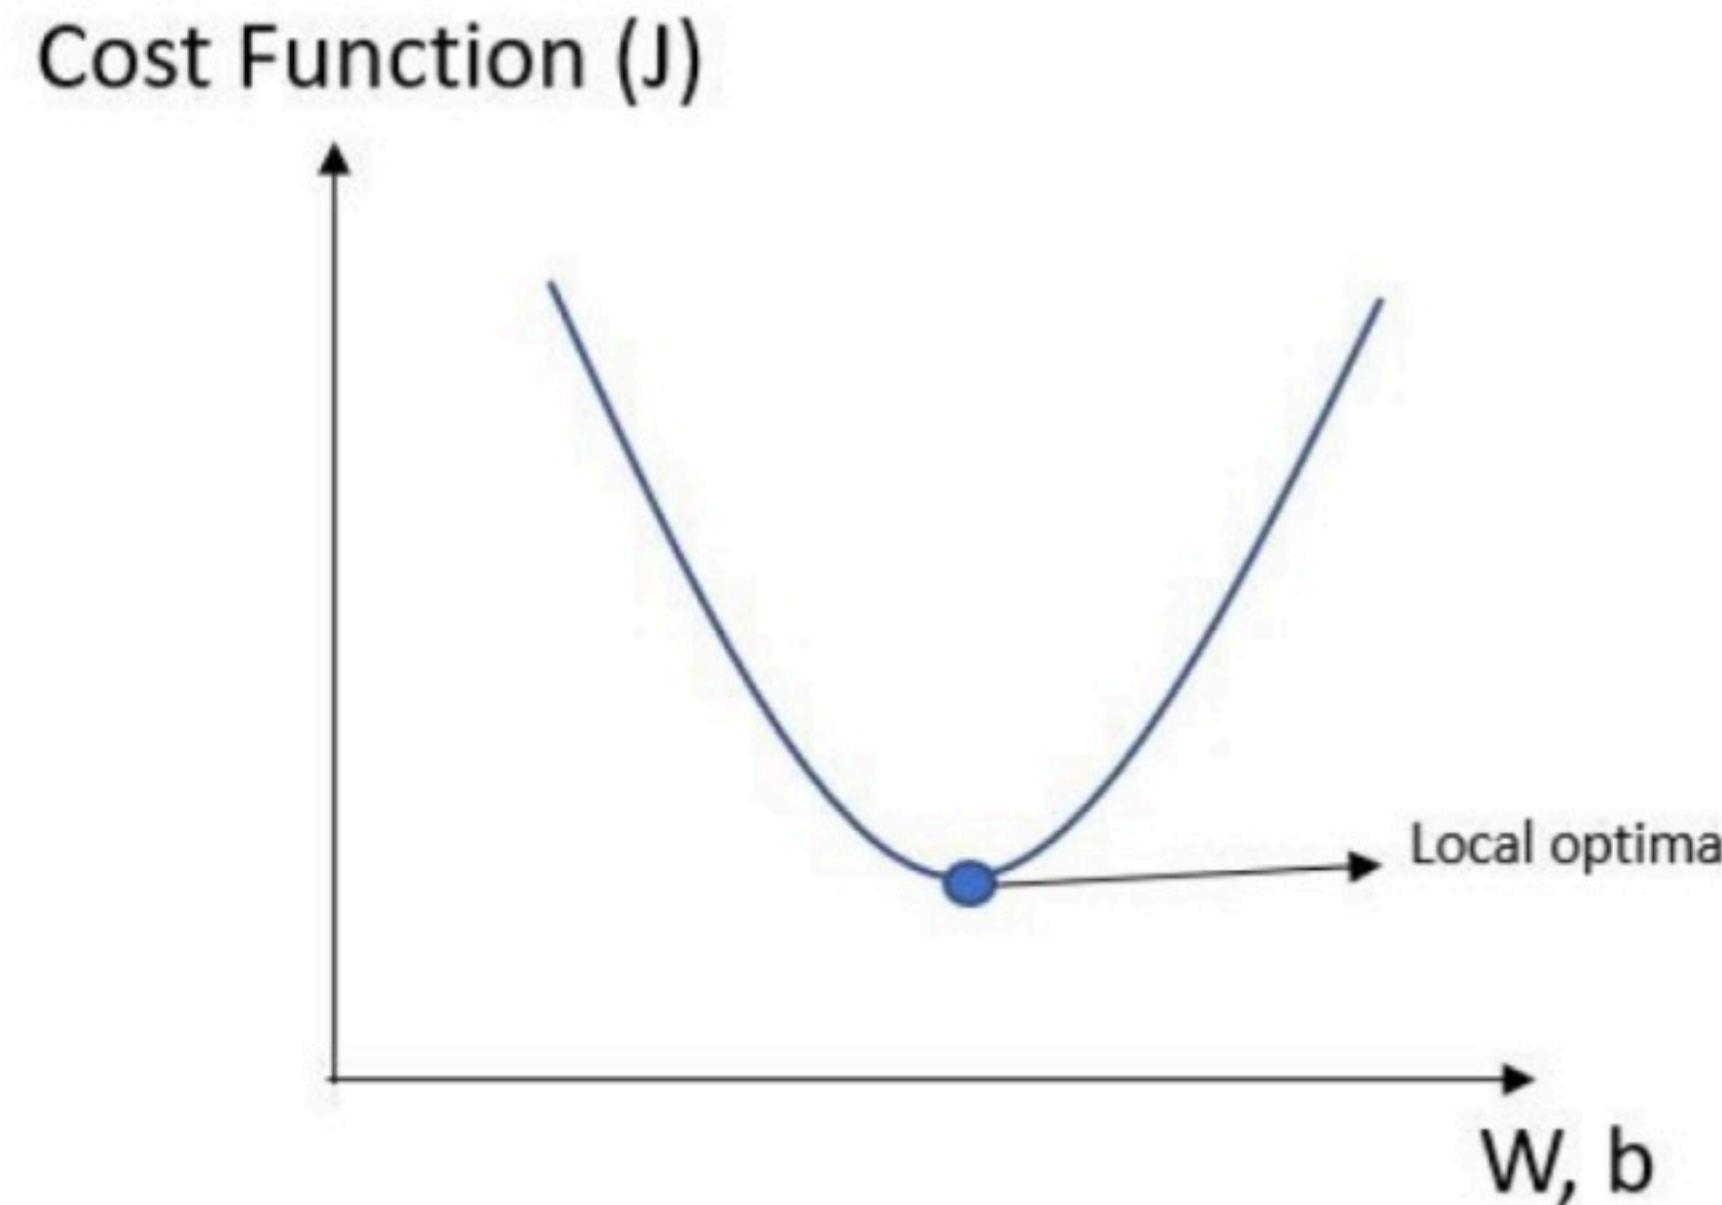
Example



Cost Function (Loss Function)

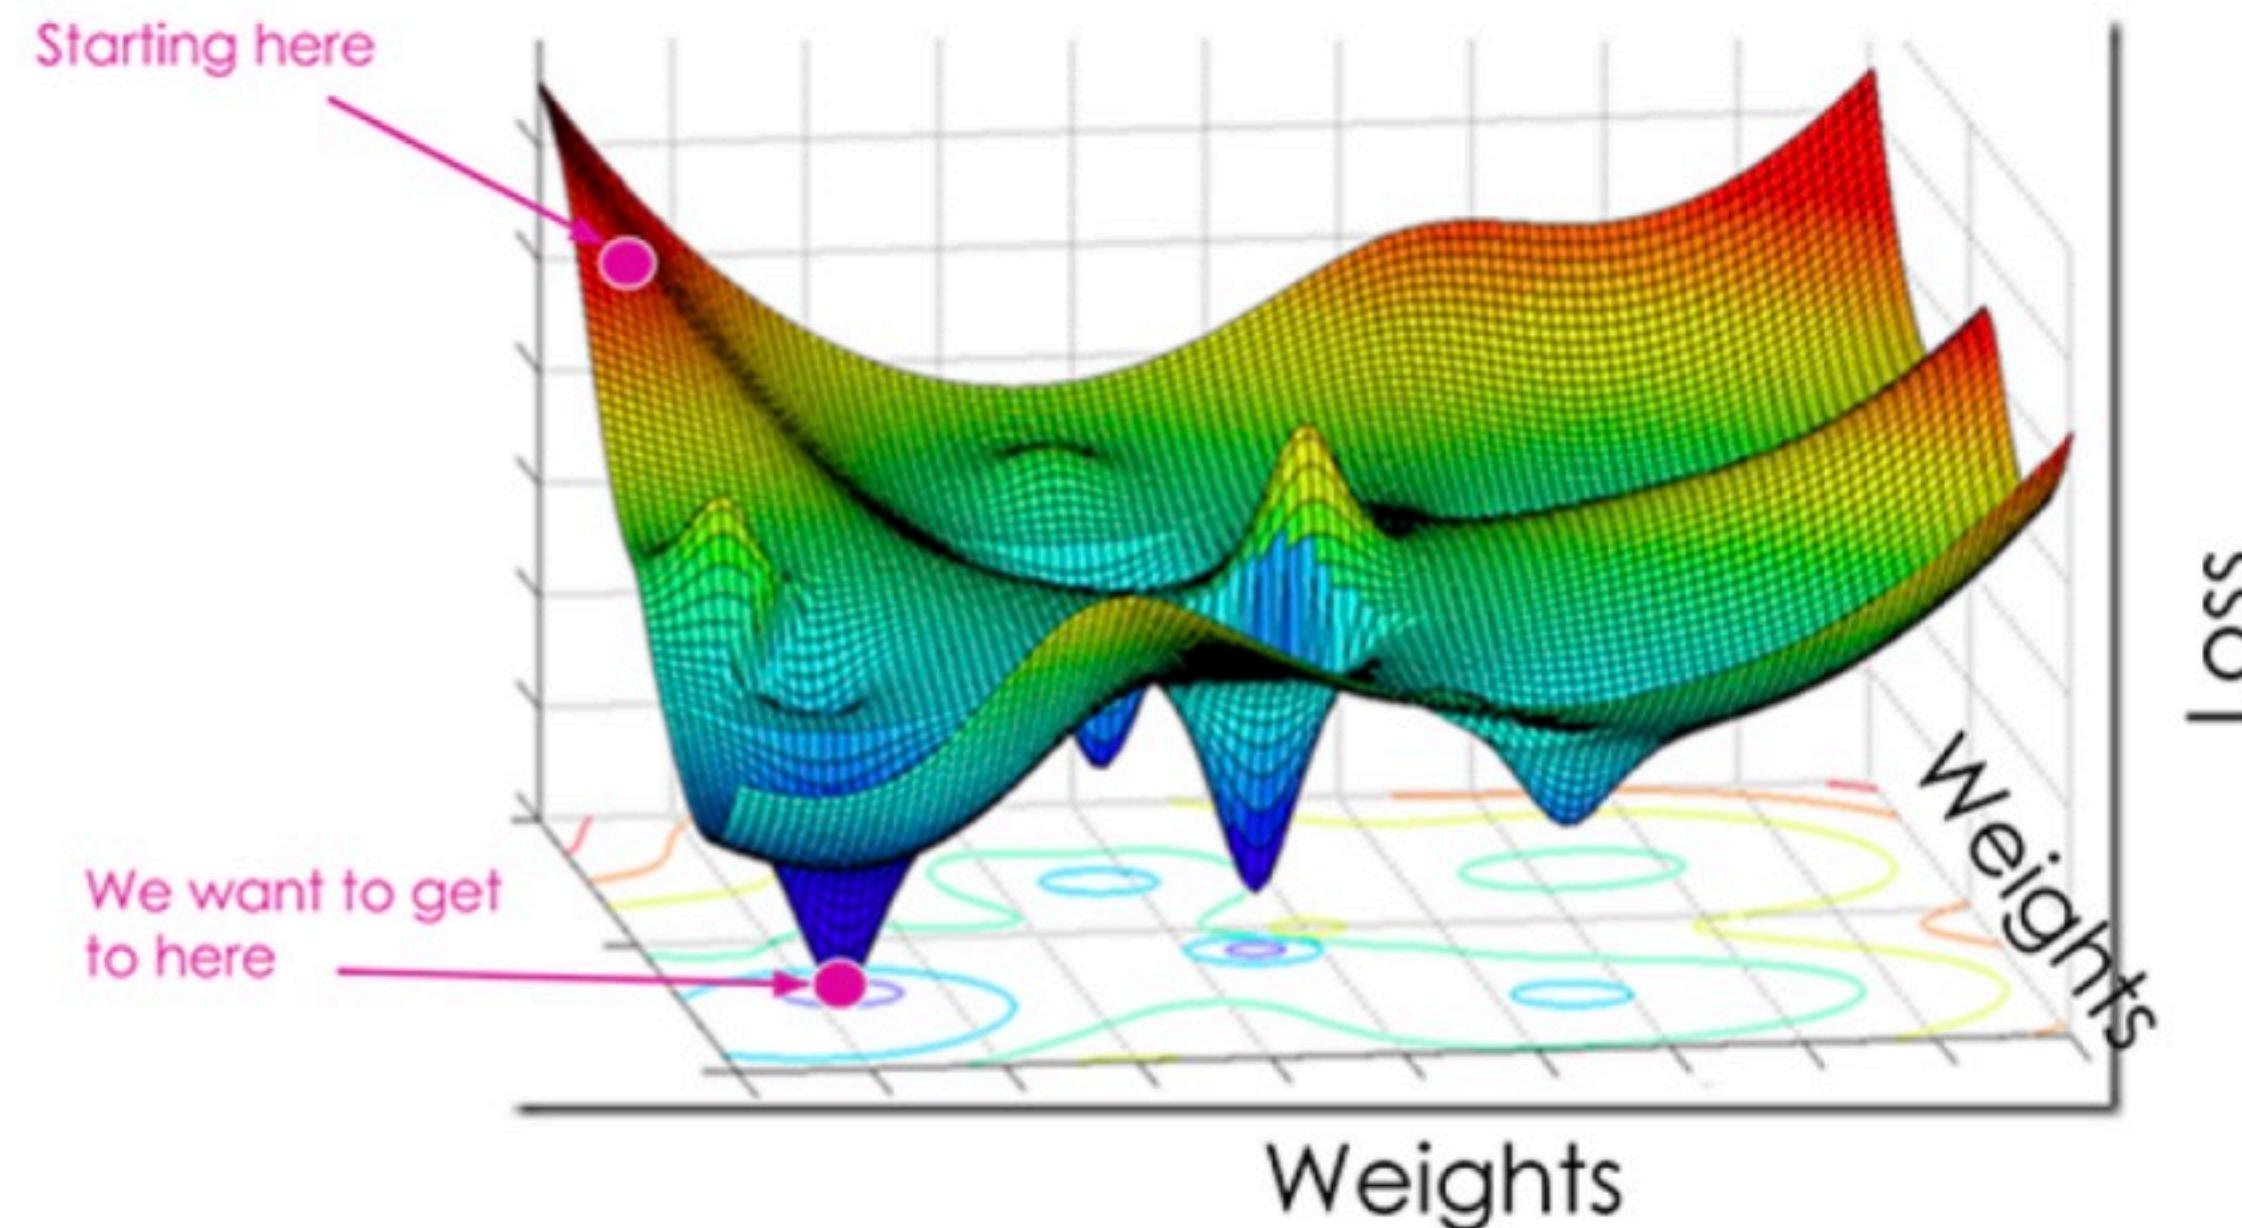


Cost Function (Loss Function)



*The objective of a ML model, is to find parameters, weights or a structure that **minimises** the cost function.*

Cost Function (Loss Function)



*The objective of a ML model, is to find parameters, weights or a structure that **minimises** the cost function.*

Types of Cost Function

Regression Loss Functions

1. Mean Squared Error Loss
2. Mean Squared Logarithmic Error Loss
3. Mean Absolute Error Loss

```
model.compile(optimizer='adam', loss='mse')
```



Binary Classification Loss Functions

1. Binary Cross-Entropy
2. Hinge Loss
3. Squared Hinge Loss

```
model.compile(loss='binary_crossentropy', optimizer='adam')
```



Multi-Class Classification Loss Functions

1. Multi-Class Cross-Entropy Loss
2. Sparse Multiclass Cross-Entropy Loss
3. Kullback Leibler Divergence Loss

```
1 model.compile(loss='categorical_crossentropy',
```

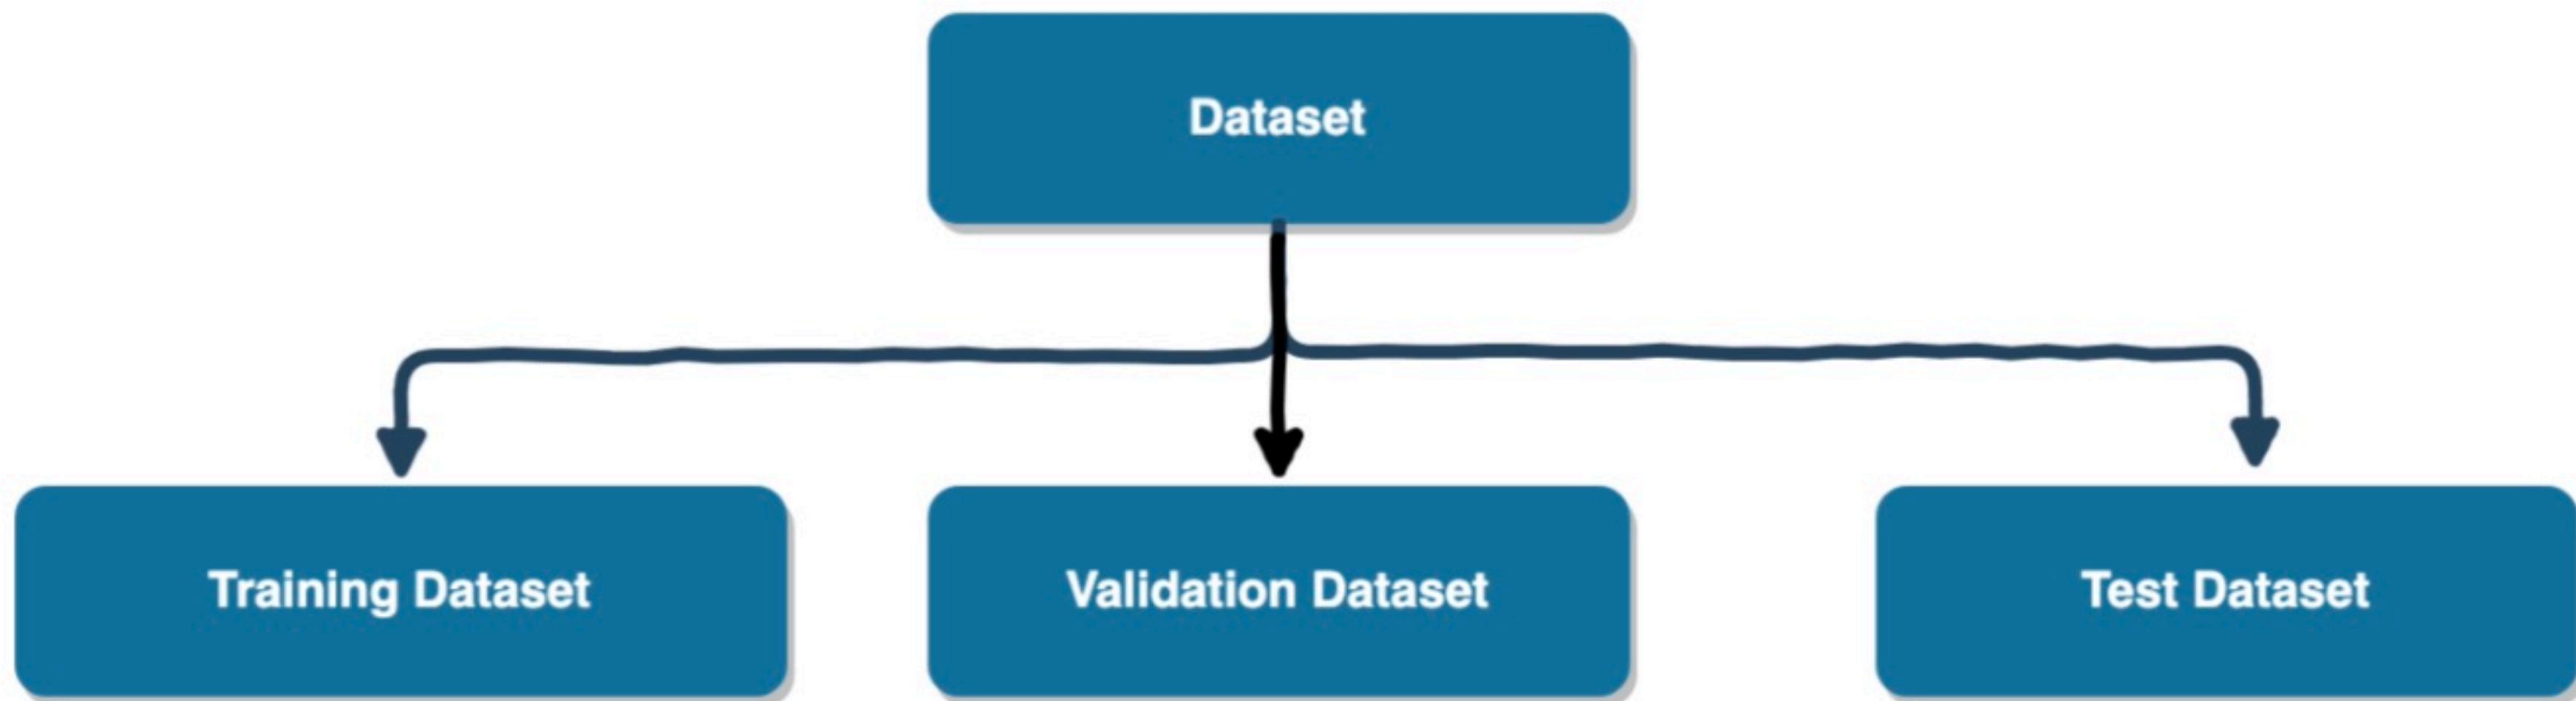


Cost Function (Loss Function)

```
model = Sequential()  
  
model.add(Dense(19,activation='relu'))  
model.add(Dense(19,activation='relu'))  
model.add(Dense(19,activation='relu'))  
model.add(Dense(19,activation='relu'))  
model.add(Dense(1))  
  
model.compile(optimizer='adam' loss='mse')
```

```
model = Sequential()  
model.add(Dense(units=30,activation='relu'))  
model.add(Dense(units=15,activation='relu'))  
model.add(Dense(units=1,activation='sigmoid'))  
model.compile loss='binary_crossentropy' , optimizer='adam'
```

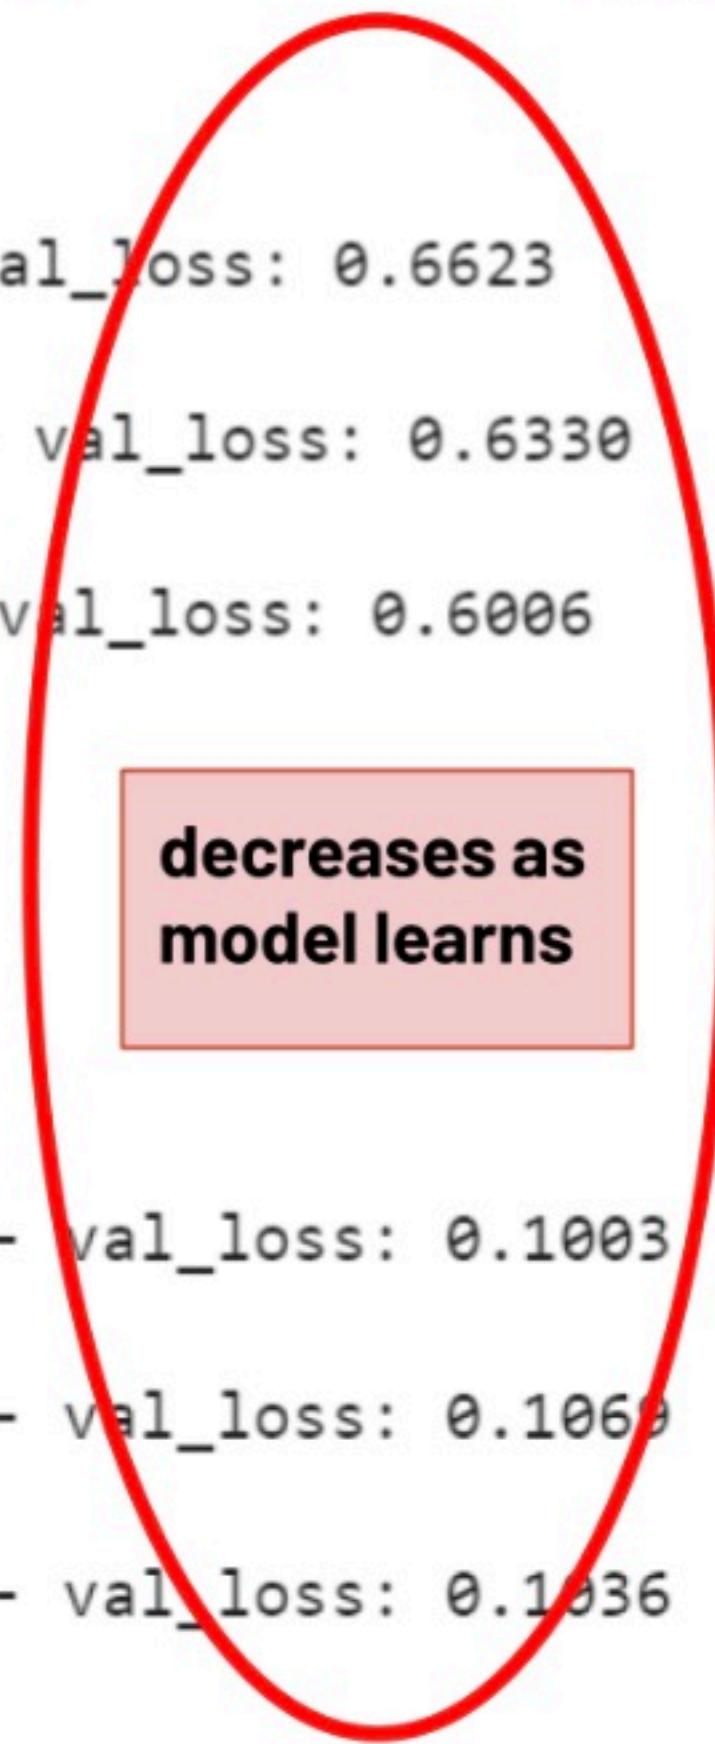
Cost Function (Loss Function)



```
model.fit(x = X_train, y = y_train,  
          validation_split = 0.1,  
          batch_size = 32,  
          epochs = 1000, verbose=1)
```

Cost Function (Loss Function)

```
Epoch 1/600  
426/426 [=====] - 0s 1ms/sample loss: 0.6807 - val_loss: 0.6623  
Epoch 2/600  
426/426 [=====] - 0s 108us/sample - loss: 0.6498 - val_loss: 0.6330  
Epoch 3/600  
426/426 [=====] - 0s 80us/sample - loss: 0.6188 - val_loss: 0.6006  
...  
Epoch 99/600  
426/426 [=====] - 0s 68us/sample - loss: 0.0485 - val_loss: 0.1003  
Epoch 100/600  
426/426 [=====] - 0s 73us/sample - loss: 0.0483 - val_loss: 0.1069  
Epoch 101/600  
426/426 [=====] - 0s 68us/sample - loss: 0.0500 - val_loss: 0.1036  
Epoch 00101: early stopping
```



**decreases as
model learns**

Cost Function (Loss Function)



UNDERFITTING VS OVERFITTING

	Underfitting	Just right	Overfitting
Deep learning illustration			
Possible remedies	<ul style="list-style-type: none">• Complexify model• Add more features• Train longer		<ul style="list-style-type: none">• Perform regularization• Get more data

Cost Function (Loss Function)

Problem Type	Output Type	Final Activation Function	Loss Function
Regression	Numerical value	Linear	Mean Squared Error (MSE)
Classification	Binary outcome	Sigmoid	Binary Cross Entropy
Classification	Single label, multiple classes	Softmax	Cross Entropy
Classification	Multiple labels, multiple classes	Sigmoid	Binary Cross Entropy